Data Science Project - Phase 1

Pantea Amoie - 400101656

1. EDA

1.1 Basic

```
import pandas as pd
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path)
```

```
print(df.head())
```

```
                                            abstract  \
0  In this paper, a robust 3D triangular mesh wat...
1  We studied an autoassociative neural network w...
2  It is well-known that Sturmian sequences are t...
3  One of the fundamental challenges of recognizi...
4  This paper generalizes previous optimal upper ...

                                             authors  n_citation  \
0            ['S. Ben Jabra', 'Ezzeddine Zagrouba']          50
1  ['Joaquín J. Torres', 'Jesús M. Cortés', 'Joaq...          50
2          ['Genevi eve Paquin', 'Laurent Vuillon']          50
3  ['Yaser Sheikh', 'Mumtaz Sheikh', 'Mubarak Shah']         221
4  ['Efraim Laksman', 'Håkan Lennerstad', 'Magnus...           0

                                          references  \
0  ['09cb2d7d-47d1-4a85-bfe5-faa8221e644b', '10aa...
1  ['4017c9d2-9845-4ad2-ad5b-ba65523727c5', 'b118...
2  ['1c655ee2-067d-4bc4-b8cc-bc779e9a7f10', '2e4e...
3  ['056116c1-9e7a-4f9b-a918-44eb199e67d6', '05ac...
4  ['01a765b8-0cb3-495c-996f-29c36756b435', '5dbc...

                                               title  \
0  A new approach of 3D watermarking based on ima...
1  Attractor neural networks with activity-depend...
2  A characterization of balanced episturmian seq...
3             Exploring the space of a human action
4  Generalized upper bounds on the minimum distan...

                                               venue  year  \
0  international symposium on computers and commu...  2008
1                                     Neurocomputing  2007
2                Electronic Journal of Combinatorics  2007
3        international conference on computer vision  2005
4  Ima Journal of Mathematical Control and Inform...  2015

                                     id
0  4ab3735c-80f1-472d-b953-fa0557fed28b
1  4ab39729-af77-46f7-a662-16984fb9c1db
2  4ab3a4cf-1d96-4ce5-ab6f-b3e19fc260de
3  4ab3a98c-3620-47ec-b578-884ecf4a6206
4  4ab3b585-82b4-4207-91dd-b6bce7e27c4e
```

```
# !pip install kaggle
# import os
# os.environ['KAGGLE_CONFIG_DIR'] = "/content/"

# # Download dataset
# !kaggle datasets download -d nechbamohammed/research-papers-dataset --unzip
```

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.11/dist-packages (1.6.17)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.11/dist-packages (from kaggle) (1.17.0)
Requirement already satisfied: certifi>=2023.7.22 in /usr/local/lib/python3.11/dist-packages (from kaggle) (2024.12.14)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.11/dist-packages (from kaggle) (2.8.2)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from kaggle) (2.32.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from kaggle) (4.67.1)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.11/dist-packages (from kaggle) (8.0.4)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.11/dist-packages (from kaggle) (2.3.0)
Requirement already satisfied: bleach in /usr/local/lib/python3.11/dist-packages (from kaggle) (6.2.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.11/dist-packages (from bleach->kaggle) (0.5.1)
```

```
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.11/dist-packages (from python-slugify->kaggle) (1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->kaggle) (3.4.1)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->kaggle) (3.10)
Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /content/kaggle.json'
Dataset URL: https://www.kaggle.com/datasets/nechbamohammed/research-papers-dataset
License(s): unknown
Downloading research-papers-dataset.zip to /content
100% 590M/590M [00:15<00:00, 41.3MB/s]
100% 590M/590M [00:15<00:00, 39.7MB/s]
```

```
# import pandas as pd

# # Load the extracted CSV file
# csv_file = "/content/dblp-v10.csv"
# df = pd.read_csv(csv_file)

# # Display the first few rows
# print(df.head())
```

```
                                            abstract  \
0  In this paper, a robust 3D triangular mesh wat...
1  We studied an autoassociative neural network w...
2  It is well-known that Sturmian sequences are t...
3  One of the fundamental challenges of recognizi...
4  This paper generalizes previous optimal upper ...

                                             authors  n_citation  \
0            ['S. Ben Jabra', 'Ezzeddine Zagrouba']          50
1  ['Joaquín J. Torres', 'Jesús M. Cortés', 'Joaq...          50
2          ['Genevi eve Paquin', 'Laurent Vuillon']          50
3  ['Yaser Sheikh', 'Mumtaz Sheikh', 'Mubarak Shah']         221
4  ['Efraim Laksman', 'Håkan Lennerstad', 'Magnus...           0

                                          references  \
0  ['09cb2d7d-47d1-4a85-bfe5-faa8221e644b', '10aa...
1  ['4017c9d2-9845-4ad2-ad5b-ba65523727c5', 'b118...
2  ['1c655ee2-067d-4bc4-b8cc-bc779e9a7f10', '2e4e...
3  ['056116c1-9e7a-4f9b-a918-44eb199e67d6', '05ac...
4  ['01a765b8-0cb3-495c-996f-29c36756b435', '5dbc...

                                               title  \
0  A new approach of 3D watermarking based on ima...
1  Attractor neural networks with activity-depend...
2  A characterization of balanced episturmian seq...
3              Exploring the space of a human action
4  Generalized upper bounds on the minimum distan...

                                               venue  year  \
0  international symposium on computers and commu...  2008
1                                     Neurocomputing  2007
2                 Electronic Journal of Combinatorics  2007
3        international conference on computer vision  2005
4  Ima Journal of Mathematical Control and Inform...  2015

                                     id
0  4ab3735c-80f1-472d-b953-fa0557fed28b
1  4ab39729-af77-46f7-a662-16984fb9c1db
2  4ab3a4cf-1d96-4ce5-ab6f-b3e19fc260de
3  4ab3a98c-3620-47ec-b578-884ecf4a6206
4  4ab3b585-82b4-4207-91dd-b6bce7e27c4e
```

Create 3 bar charts. For each bar chart draw the number of publications in a given year range. Do this for the ranges 1937-1950, 1950-1970, 1970-1990. Compare the three.

Approach:

1. **Data Preparation**:
   - The dataset contains a "**year**" column, which is essential for grouping publications by year. The code ensures that this column is treated as a numeric value to allow proper comparisons and filtering.
2. **Defining Year Ranges**:
   - The approach divides the publication years into three specified ranges: 1937-1950, 1950-1970, and 1970-1990. These ranges allow us to examine how publication trends evolved over time, particularly in these historical periods.
3. **Filtering and Counting Publications**:

- For each year range, the code filters the dataset to only include the rows where the publication year falls within the specified range.
- It then counts the number of publications for each year in the range, which helps us understand how publication activity is distributed across the years.

4. **Plotting**:
   - For each of the three year ranges, a bar chart is created to display the number of publications per year.
   - The bar charts are generated side by side for easy comparison, with the x-axis representing the years and the y-axis representing the number of publications.
   - Each plot is labeled appropriately to reflect the corresponding time period.
5. **Comparison**:
   - The visual comparison between the three bar charts helps highlight trends in publication activity during each of the year ranges. By observing the differences in the number of publications across the years, one can infer how academic or research output evolved in those periods.

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import pearsonr, spearmanr

# Convert 'year' column to numeric
df['year'] = pd.to_numeric(df['year'], errors='coerce')

# Define year ranges
year_ranges = [(1937, 1950), (1950, 1970), (1970, 1990)]
fig, axes = plt.subplots(1, 3, figsize=(15, 5))

# Plot bar charts for each range
for i, (start, end) in enumerate(year_ranges):
    df_range = df[(df['year'] >= start) & (df['year'] < end)]
    year_counts = df_range['year'].value_counts().sort_index()
    axes[i].bar(year_counts.index, year_counts.values, color='b')
    axes[i].set_title(f'Publications ({start}-{end})')
    axes[i].set_xlabel('Year')
    axes[i].set_ylabel('Count')
    axes[i].tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()
```

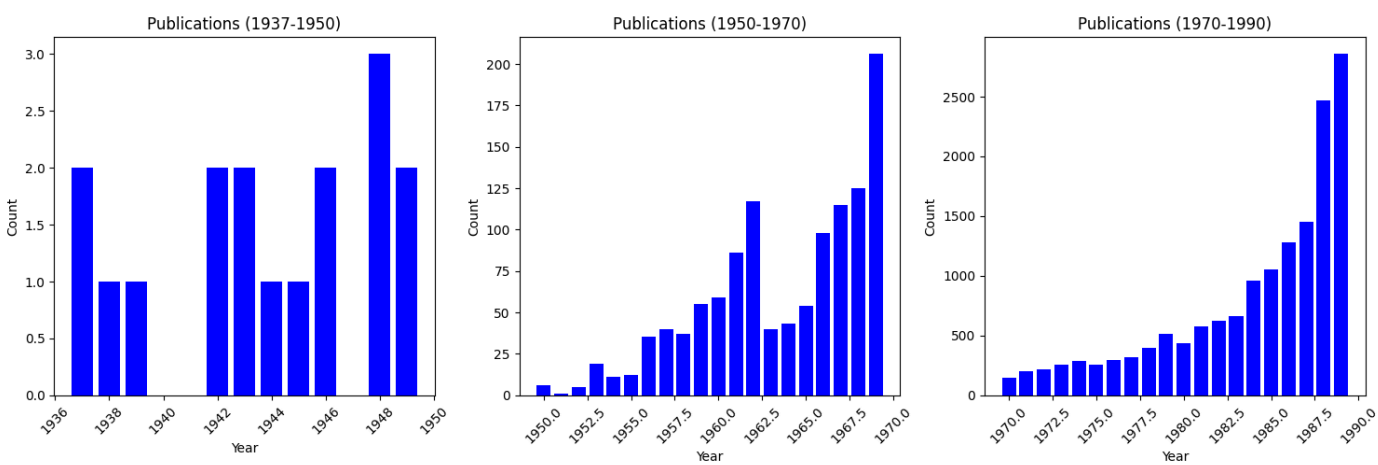

Result Analysis:

The three bar charts depict the number of publications in three time periods: 1937-1950, 1950-1970, and 1970-1990.

1937-1950:

Scale and Growth: The count of publications per year is relatively low, with values typically below 3.0.

Distribution: The publications are sporadic and do not follow a consistent growth trend. Peaks are noticeable in certain years like 1948.

Observation: This period represents a slow and inconsistent growth phase in publications.

1950-1970:

Scale and Growth: There is a clear upward trend in the number of publications, starting from minimal values around 1950 and growing to nearly 200 by 1970.

Distribution: The growth appears exponential, with a steady increase as time progresses.

Observation: This phase demonstrates a shift toward consistent and rapid growth, marking a transition in publication trends.

1970-1990:

Scale and Growth: The number of publications skyrockets, reaching over 2500 by 1990. The growth is exponential, significantly outpacing the earlier periods.

Distribution: The publications grow steadily and sharply, reflecting an era of mass production and dissemination of information.

Observation: This period shows the most dramatic growth, indicating a boom in academic, industrial, or general publishing.

Overall Comparison:

Magnitude: The number of publications grows exponentially across the three periods, with 1970-1990 vastly outpacing the earlier periods.
Trend: Growth is minimal in the 1937-1950 period, begins to accelerate in 1950-1970, and reaches a peak exponential rise in 1970-1990.
Implication: The increasing trend reflects advancements in technology, communication, and the global expansion of knowledge-sharing platforms over time.

The stark difference between the three periods highlights the evolving nature of publication activities.

Create a bar chart of the number of references over the years.

Approach:

1. **Data Preparation**:
   - The dataset is loaded, specifically focusing on the "year" and "references" columns, which are crucial for the analysis. Missing values in these columns are removed to ensure the data used is clean and complete.
2. **Processing the References**:
   - The "references" column contains lists of references for each publication. The code processes these lists by evaluating each reference as a Python literal (i.e., converting strings that represent lists into actual lists).
   - For each publication, the code calculates the **number of references** by counting the elements in the list. If the "references" field is not in the correct format (not a string), a count of zero is assigned.
3. **Aggregating the Data**:
   - After processing the references, the code aggregates the data by year. This step sums the total number of references for each year, which allows the creation of a meaningful yearly comparison.
4. **Visualization**:
   - A bar chart is created to visualize the aggregated total number of references for each year. The years are displayed on the x-axis, and the corresponding number of references is shown on the y-axis.
   - The chart is customized with labels, a title, and a rotation on the x-axis for better readability.
5. **Insight**:
   - The visualization helps to observe trends in reference activity over the years, such as increases or decreases in the number of references, and provides insight into how research output has evolved over time.

In summary, the code processes and aggregates the data related to references and then presents the findings in a bar chart to showcase trends in the total number of references across different years.

```
# Bar chart: Number of references over the years

import pandas as pd
import matplotlib.pyplot as plt
```

```
import ast

df_cols = pd.read_csv(file_path, usecols=["year", "references"], low_memory=False)

df_cols = df_cols.dropna(subset=["year", "references"])
df_cols["number_of_references"] = df_cols["references"].apply(lambda x: len(ast.literal_eval(x)) if isinstance(x, str) else 0)

# Aggregating total number of references per year
references_per_year = df_cols.groupby("year")["number_of_references"].sum().reset_index()

plt.bar(references_per_year["year"], references_per_year["number_of_references"])
plt.xlabel("Year")
plt.ylabel("Total Number of References")
plt.title("Total Number of References Over the Years")
plt.xticks(rotation=45)
plt.show()
```

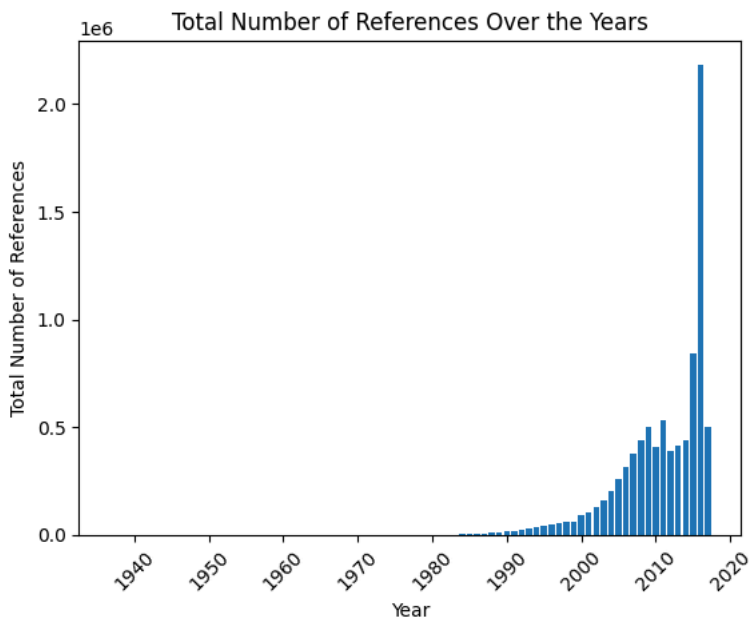

Result Analysis:

The bar chart shows the total number of references over the years. From the graph, we can observe the following trends:

1. **Low Activity Before 2000**: There is very little activity in terms of references before the year 2000. The bars are relatively short, suggesting limited citations or references during that period.
2. **Sharp Increase After 2000**: Starting from around the year 2000, there is a significant increase in the number of references. This growth becomes even more dramatic after 2010, where the number of references appears to spike drastically.
3. **Exponential Growth**: The steep rise in references in the 2010s suggests a surge in activity or research in a particular field or on a specific topic, potentially driven by advancements in technology, new research developments, or the availability of data.

This pattern is often seen in rapidly developing fields such as technology, AI, or other scientific advancements where citations or references start growing exponentially as more work gets published and shared. It could indicate a breakthrough or increased awareness around a subject.

Create a bar chart of the number of authors over the years.

Approach:

The general approach of this code is focused on visualizing the total number of authors contributing to publications over the years.

1. **Data Preparation**:
    - The dataset is loaded, focusing specifically on the "year" and "authors" columns. These columns contain the relevant information for the analysis: the year of publication and the list of authors for each publication. Missing values in either of these columns are removed to ensure clean data.

2. **Processing the Authors**:
   - The "authors" column contains a list of authors, but it may be stored as a string representation of a list. The code converts these string representations into actual Python lists.
   - For each paper, the number of authors is determined by counting the number of items in the list (if the list exists). If the authors column is not in the expected format, a count of zero is assigned.
3. **Aggregating the Data**:
   - The code then aggregates the data by year, summing up the total number of authors for each year. This gives an overall count of authors contributing to publications in any given year.
4. **Visualization**:
   - A bar chart is generated where the x-axis represents the year of publication, and the y-axis represents the total number of authors for that year. The chart is further customized with labels and a title for clarity.
   - The x-axis labels (representing years) are rotated for better readability, especially if there are many years to display.
5. **Insight**:
   - The bar chart helps visualize the trend of how the total number of authors has changed over the years. It provides insight into the increasing collaboration in research, with a growing number of authors per publication, especially in modern times.

In summary, the code processes the authors' data, aggregates it by year, and presents the trends in the total number of authors contributing to publications over time through a bar chart.

```
# Bar chart: Number of authors over the years

import pandas as pd
import matplotlib.pyplot as plt
import ast

df_cols = pd.read_csv(file_path, usecols=["year", "authors"], low_memory=False)

df_cols = df_cols.dropna(subset=["year", "authors"])
df_cols["number_of_authors"] = df_cols["authors"].apply(lambda x: len(ast.literal_eval(x)) if isinstance(x, str) else 0)

# Aggregating total number of authors per year
authors_per_year = df_cols.groupby("year")["number_of_authors"].sum().reset_index()

plt.bar(authors_per_year["year"], authors_per_year["number_of_authors"])
plt.xlabel("Year")
plt.ylabel("Total Number of Authors")
plt.title("Total Number of Authors Over the Years")
plt.xticks(rotation=45)
plt.show()
```

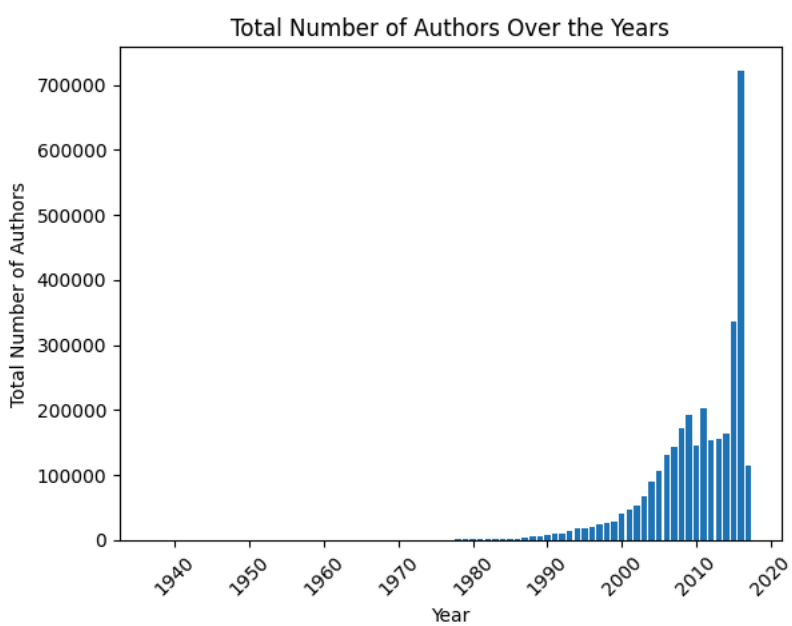

Result Analysis:

1. **Minimal Activity Before 2000**: Similar to the previous chart showing references, we see that the number of authors is relatively low before 2000. This suggests that fewer researchers were contributing to the field or area being tracked in the dataset during this period.
2. **Sharp Increase After 2000**: The number of authors begins to rise gradually starting around the year 2000. This trend intensifies significantly after 2010, with a sharp peak occurring around 2015–2020, similar to what we saw in the reference chart.
3. **Exponential Growth in Recent Years**: The growth in the number of authors is strikingly exponential in the 2010s, with the numbers reaching well over 700,000. This mirrors the pattern in the number of references, suggesting that the increase in research activity is driven by an expanding number of contributors (authors).

The dramatic rise in both references and authors since the 2010s likely reflects a surge in academic publishing, possibly driven by easier access to publishing platforms, larger collaborations, or an increasing focus on high-impact topics. This could also point to a global collaboration trend in research fields such as technology and AI, where many authors might be collaborating across institutions and countries.

Find the Pearson correlation coefficient and Spearman Rank correlation coefficient between the number of authors and number of references.

```
import pandas as pd
import ast
from scipy.stats import pearsonr, spearmanr

df_cols = pd.read_csv(file_path, usecols=["authors", "references"], low_memory=False)

df_cols = df_cols.dropna(subset=["authors", "references"])

df_cols["num_authors"] = df_cols["authors"].apply(lambda x: len(ast.literal_eval(x)) if isinstance(x, str) else 0)
df_cols["num_references"] = df_cols["references"].apply(lambda x: len(ast.literal_eval(x)) if isinstance(x, str) else 0)

pearson_corr, pearson_pval = pearsonr(df_cols["num_authors"], df_cols["num_references"])
spearman_corr, spearman_pval = spearmanr(df_cols["num_authors"], df_cols["num_references"])

print(f"Pearson Correlation Coefficient: {pearson_corr}")
print(f"Spearman Correlation Coefficient: {spearman_corr}")
```

```
Pearson Correlation Coefficient: 0.05474922437268194
Spearman Correlation Coefficient: 0.08899471045023705
```

Result Analysis:

Pearson Correlation Coefficient (0.055):

- **Interpretation**: The Pearson correlation coefficient measures the linear relationship between two variables. A value of 0.055 suggests that there is **a very weak positive linear relationship** between the number of authors and the number of references. However, this value is very close to zero, indicating that the two variables are nearly uncorrelated in a linear sense.
- **Significance**: The low value of the Pearson correlation also implies that, based on this data, the number of authors doesn't strongly influence the number of references in a linear fashion.

Spearman Rank Correlation Coefficient (0.089):

- **Interpretation**: The Spearman rank correlation coefficient measures the monotonic relationship (whether the variables tend to move in the same or opposite direction) between two variables. A value of 0.089 indicates a **very weak positive monotonic relationship** between the number of authors and the number of references. Again, the value is close to zero, suggesting very little or no monotonic association.
- **Significance**: Since Spearman evaluates the ranks, it doesn't require the relationship to be linear, but the result still suggests a very weak association between the two variables.

  Both the Pearson and Spearman correlation coefficients are very low, indicating that there is essentially no significant correlation between the number of authors and the number of references in this dataset. This might suggest that the number of authors does not directly influence the number of references in a consistent way across the papers in our dataset.

Find the Pearson and Spearman correlation coefficient between the number of authors and number of citations.

```
import pandas as pd
import ast
from scipy.stats import pearsonr, spearmanr

df_cols = pd.read_csv(file_path, usecols=["authors", "n_citation"], low_memory=False)

df_cols = df_cols.dropna(subset=["authors", "n_citation"])

df_cols["num_authors"] = df_cols["authors"].apply(lambda x: len(ast.literal_eval(x)) if isinstance(x, str) else 0)
df_cols["n_citation"] = df_cols["n_citation"].astype(int)

pearson_corr, pearson_pval = pearsonr(df_cols["num_authors"], df_cols["n_citation"])
spearman_corr, spearman_pval = spearmanr(df_cols["num_authors"], df_cols["n_citation"])

print(f"Pearson Correlation Coefficient: {pearson_corr}")
print(f"Spearman Correlation Coefficient: {spearman_corr}")
```

```
Pearson Correlation Coefficient: -0.002797165077791142
Spearman Correlation Coefficient: -0.01661797988372679
```

Result Analysis:

Pearson Correlation Coefficient (-0.003):

- **Interpretation**: The Pearson correlation coefficient measures the linear relationship between the number of authors and the number of citations. A value of -0.003 indicates an **extremely weak negative linear relationship** between these two variables. Essentially, this means there is almost no linear correlation between the number of authors and the number of citations.
- **Significance**: The result is very close to zero, suggesting that the number of authors has virtually no effect on the number of citations in a linear fashion. This implies that in your dataset, papers with more authors do not necessarily have more or fewer citations in a clear linear trend.

Spearman Rank Correlation Coefficient (-0.017):

- **Interpretation**: The Spearman rank correlation coefficient assesses the monotonic relationship between two variables, meaning whether they generally move in the same or opposite directions (not necessarily in a linear manner). A value of -0.017 indicates a **very weak negative monotonic relationship** between the number of authors and the number of citations.
- **Significance**: Like the Pearson correlation, the result suggests that there is little to no monotonic relationship between the number of authors and the number of citations. The small negative value indicates that if there is any trend, it would be a very weak inverse one, meaning papers with more authors might have slightly fewer citations, but this trend is almost negligible.

Both correlation coefficients are extremely close to zero, indicating that there is very little to no relationship between the number of authors and the number of citations in the dataset. The weak negative values suggest that if there is any relationship at all, it would be a very slight inverse one (papers with more authors might have slightly fewer citations, but this is not significant).

Draw a bar chart of the title length over the years.

```
import matplotlib.pyplot as plt

# Calculate the title length for each paper
df["title_length"] = df["title"].apply(len)

# Group by year and calculate the average title length
title_length_per_year = df.groupby("year")["title_length"].mean().reset_index()

# Plot the bar chart
plt.bar(title_length_per_year["year"], title_length_per_year["title_length"])
plt.xlabel("Year")
plt.ylabel("Average Title Length")
plt.title("Average Title Length Over the Years")
plt.xticks(rotation=45)
plt.show()
```

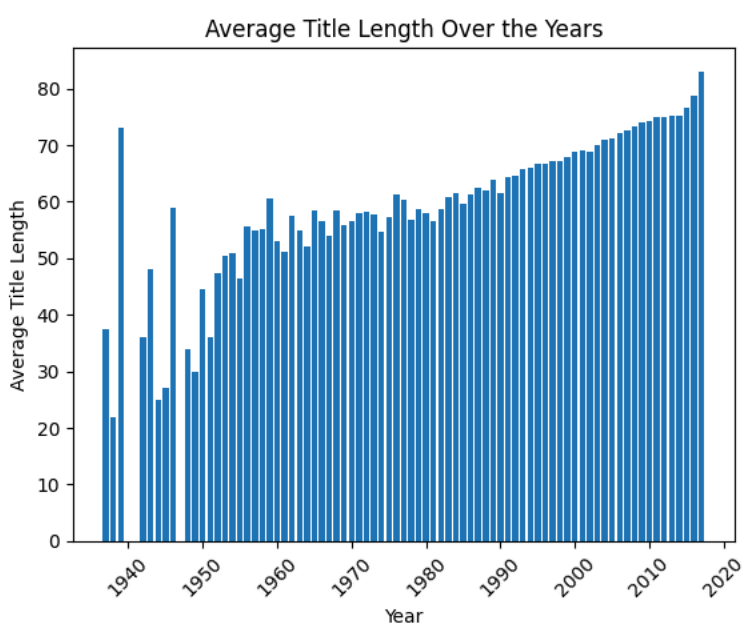

Result Analysis:

The bar chart shows the **average title length** of papers over the years.

1. **Fluctuations in Early Years**: From the 1940s to around the 1960s, we observe significant fluctuations in average title length. Some years have very short or very long titles, suggesting inconsistency or variation in naming conventions during this period.
2. **Gradual Increase Over Time**: After the 1960s, the average title length appears to generally increase over time. By the 1980s and 1990s, the trend becomes more consistent, with the average length of titles steadily growing.
3. **Steady Growth in Recent Years**: In the 2010s and 2020s, the average title length continues to rise, reaching over 70 characters. This suggests that modern titles are becoming more descriptive or detailed. This could be due to trends in research papers becoming more specialized or needing to include more information, such as methodologies or results, within the title itself.
4. **Possible Causes for Longer Titles**: Several factors might contribute to longer titles over the years:
   - **Increased specificity**: Modern research often requires more detailed titles to reflect complex subjects.
   - **Descriptive titles**: Research titles may now need to be more self-explanatory due to interdisciplinary fields or to capture specific research aspects.
   - **Standardized formats**: There could also be a trend toward standardized formats that lead to longer titles.

This general trend of increasing title length may reflect a shift toward more complex, detailed, and informative titles in academic publications.

Draw a wordcloud of the abstracts.

```
import pandas as pd
from wordcloud import WordCloud
import matplotlib.pyplot as plt

df_cols = pd.read_csv(file_path, usecols=["abstract"], low_memory=False)

df_cols = df_cols.dropna(subset=["abstract"])

# Sample (otherwise I would get RAM error)
sampled_df = df_cols.sample(n=100000, random_state=42)

# Join the sampled abstracts into a single string
text = " ".join(sampled_df["abstract"])

# Generate the word cloud
wordcloud = WordCloud(width=800, height=400, background_color="white", colormap="viridis").generate(text)

plt.figure(figsize=(12, 6))
plt.imshow(wordcloud, interpolation="bilinear")
plt.axis("off")
plt.show()
```

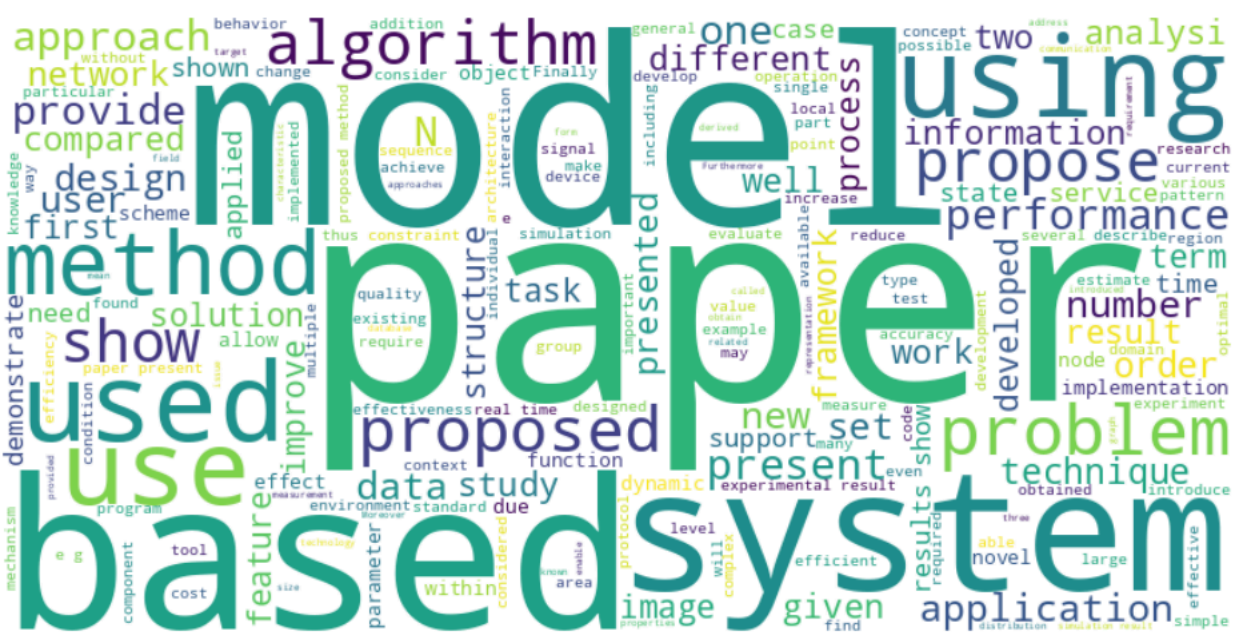
model
using
paper
based
system
used
use
method
algorithm
approach
network
provide
compared
design
user
first
show
solution
need
demonstrate
improve
feature
proposed
present
data
study
structure
task
presented
new
set
support
work
problem
technique
application
image
given
propose
information
performance
state
service
analysi
two
one
case
different
process
well
framework
developed
term
time
number
result
order
implementation
experiment
novel
large
simple
effective
obtained
introduce
results
required
efficient
level
complex
will
protocol
dynamic
experimental
result
within
area
parameter
considered
due
standard
environment
effect
context
function
tool
cost
component
mechanism
shown
change
consider
object
Finally
develop
operation
single
local
part
point
derived
furthermore
increase
evaluate
reduce
simulation
thus
constraint
individual
architecture
interaction
signal
make
device
sequence
achieve
approaches
N
applied
implemented
proposed method
including
important
quality
existing
database
require
group
value
obtain
example
related
may
called
available
representation
type
test
accuracy
measure
code
show
many
even
effectiveness
real time
designed
paper
present
allow
multiple
efficiency
condition
provided
program
measurement
moreover
technology
known
enable
properties
find
distribution
simulation result
size
e g
way
knowledge
field
scheme
found
mean
behavior
addition
general
concept
possible
address
communication
target
without
particular
derived
form
equipment
research
current
various
pattern
estimate
introduced
domain
optimal
node
region
several
describe
experiment
group
throw
able
simple
effective
large
obtained
introduce
required
novel
efficient
level
complex
will
protocol
dynamic
experimental
result
within
area
parameter
considered
due
standard
environment
effect
context
function
tool
cost
component
mechanism
provided
program
measurement
moreover
technology

Approach:

The general approach of this code focuses on calculating and analyzing the correlation between the **title length** of a paper and the **average title length of the papers it references**.

1. **Data Preparation**:
   - The dataset is loaded, specifically focusing on the "**id**", "**title**", and "**references**" columns. These columns are essential for the analysis, as they provide the information needed to calculate the title lengths and identify the referenced papers.
   - Any rows with missing **titles** or **references** are removed to ensure that the analysis is performed on complete data.
2. **Calculating Title Lengths**:
   - The length of each paper's title is calculated by counting the number of characters in the "**title**" column. This gives a numerical representation of the title length.
3. **Creating a Dictionary**:
   - A dictionary is created where each paper's **id** is mapped to its corresponding **title length**. This dictionary is used later to look up the title lengths of referenced papers.
4. **Calculating Average Title Length of Referenced Papers**:
   - For each paper, the code calculates the average title length of the papers that it references. This is done by extracting the **references** for each paper, looking up the title lengths of those referenced papers using the dictionary created earlier, and then calculating the average length.
   - If a paper does not reference any other papers or has no valid references, the average length is set to 0.
5. **Calculating Correlation**:
   - The code calculates two types of correlation coefficients to assess the relationship between **title length** and **average referenced title length**:
     - **Pearson correlation**: Measures the linear relationship between the two variables.
     - **Spearman correlation**: Measures the monotonic relationship, capturing both linear and non-linear trends.
   - The correlation coefficients help determine whether there is a relationship between the title lengths of a paper and the average title length of the papers it references.
6. **Interpretation**:
   - The Pearson and Spearman correlation coefficients provide insight into whether papers with longer titles tend to reference other papers with longer titles, or if there is no significant relationship between the two. The magnitude and direction of the coefficients indicate the strength and type of relationship, and the results can help infer patterns in academic publishing.

```
import pandas as pd
import ast
from scipy.stats import pearsonr, spearmanr

# Load the CSV file
df = pd.read_csv(file_path, usecols=["id", "title", "references"], low_memory=False)

# Drop rows with missing titles or references
df = df.dropna(subset=["title", "references"])

# Calculate title length for each paper
df["title_length"] = df["title"].apply(len)

# Create a dictionary mapping paper IDs to title lengths
id_to_title_length = df.set_index("id")["title_length"].to_dict()

# Function to calculate the average title length of referenced papers
def calculate_avg_referenced_title_length(references):
    refs = ast.literal_eval(references) if isinstance(references, str) else []
    ref_lengths = [id_to_title_length[ref] for ref in refs if ref in id_to_title_length]
    return sum(ref_lengths) / len(ref_lengths) if ref_lengths else 0

# Calculate the average title length of referenced papers for each paper
df["avg_referenced_title_length"] = df["references"].apply(calculate_avg_referenced_title_length)

# Calculate correlation coefficients
pearson_corr, _ = pearsonr(df["title_length"], df["avg_referenced_title_length"])
spearman_corr, _ = spearmanr(df["title_length"], df["avg_referenced_title_length"])
```

```
print(f"Pearson Correlation Coefficient: {pearson_corr}")
print(f"Spearman Correlation Coefficient: {spearman_corr}")
```

```
Pearson Correlation Coefficient: 0.11150914052954312
Spearman Correlation Coefficient: 0.1709170543887765
```

Result Analysis:

1. **Pearson Correlation Coefficient (0.112)**:
   - **Interpretation**: The Pearson correlation coefficient measures the linear relationship between the title length of the paper and the average title length of the papers it references. A value of 0.112 indicates a **weak positive linear relationship** between these two variables. While there is some positive correlation, the relationship is weak, suggesting that longer titles are slightly more likely to reference papers with longer titles, but this is not a strong linear trend.
   - **Significance**: The positive value shows a slight tendency for papers with longer titles to reference other papers with longer titles, but this relationship is not particularly strong or conclusive. The correlation is weak enough to imply that factors other than title length are likely more important in determining the titles of referenced papers.
2. **Spearman Rank Correlation Coefficient (0.171)**:
   - **Interpretation**: The Spearman rank correlation coefficient measures the monotonic relationship between the two variables. A value of 0.171 indicates a **weak positive monotonic relationship**. In other words, papers with longer titles tend to reference papers with longer titles, but the relationship is weak. Unlike Pearson's linear correlation, Spearman's correlation also accounts for non-linear but still monotonic trends.
   - **Significance**: The positive value, though still weak, suggests a slight tendency that when a paper has a longer title, it is more likely to reference other papers with longer titles. However, the relationship is not strong, and many other factors likely play a more significant role in determining the titles of references.

Both the Pearson and Spearman correlation coefficients are relatively low, indicating that while there is a slight positive relationship between the title length of a paper and the average title length of the papers it references, this relationship is weak.

In practical terms, this suggests that while there might be some tendency for papers with longer titles to reference other papers with longer titles, it is not a strong or significant factor in the way references are made. Title length likely isn't a key influence on citation practices or reference selection.

Find the top 10 authors with the most publications.

```
import pandas as pd
import ast

# Load dataset
df = pd.read_csv(file_path, usecols=["authors"], low_memory=False)

# Drop missing values
df = df.dropna(subset=["authors"])

# Split authors and count publications
df["authors"] = df["authors"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])
all_authors = df.explode("authors")  # Flatten the list of authors
author_counts = all_authors["authors"].value_counts().head(10)

print("Top 10 Authors with the Most Publications:")
print(author_counts)
```

```
Top 10 Authors with the Most Publications:
authors
Wei Wang       950
Wei Zhang      657
Yang Liu       629
Lei Zhang      579
Wei Li         559
Jun Wang       544
Lei Wang       519
Lajos Hanzo    458
Wei Liu        456
Jun Zhang      455
Name: count, dtype: int64
```

Find the top 10 authors with the most citations.

```
# Load dataset with authors and citation counts
df = pd.read_csv(file_path, usecols=["authors", "n_citation"], low_memory=False)

# Drop missing values
df = df.dropna(subset=["authors", "n_citation"])

# Split authors and sum citations
df["authors"] = df["authors"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])
all_authors = df.explode("authors")  # Flatten the list of authors
author_citations = all_authors.groupby("authors")["n_citation"].sum().sort_values(ascending=False).head(10)

print("Top 10 Authors with the Most Citations:")
print(author_citations)
```

```
Top 10 Authors with the Most Citations:
authors
David G. Lowe        65344
Hari Balakrishnan    55096
Scott Shenker        54164
Ian F. Akyildiz      53654
Michael I. Jordan    53448
Ion Stoica           52890
Chih-Jen Lin         52302
Takeo Kanade         50743
Deborah Estrin       49925
Vladimir Vapnik      49755
Name: n_citation, dtype: int64
```

Find the top 10 papers with the most references.

```
import pandas as pd
import ast

# Load dataset with references and titles
df = pd.read_csv(file_path, usecols=["id", "references", "title"], low_memory=False)

# Drop missing values
df = df.dropna(subset=["references", "title"])

# Calculate the number of references for each paper
df["num_references"] = df["references"].apply(lambda x: len(ast.literal_eval(x)) if isinstance(x, str) else 0)

# Sort by number of references and get the top 10
top_papers_references = df.nlargest(10, "num_references")[["id", "title", "num_references"]]

print("Top 10 Papers with the Most References:")
print(top_papers_references)
```

```
Top 10 Papers with the Most References:
                                          id  \
371369  c6090c87-4730-4ae6-86cd-7663adb23b2b
780292  a1c66cbc-77ad-4167-9787-4936221427ad
104143  6d4c5b32-8e13-4022-b67f-1ace7ffc91d0
214646  91f1068b-9feb-4799-b920-67205c36999b
484969  ebe93176-7150-4ce6-8d77-880a4e13beb0
223901  95044198-6db4-4224-a499-d8e8a3e1f52d
302124  af0e7b77-e5c3-4d64-adc2-cc43a4b0976e
707510  0d3b901e-4179-4ed6-badf-7ce4237dadd4
325083  b6b998ee-2e78-4d0e-8c55-e7806e449b9e
538381  fdbb0c65-8b80-43e7-a839-8738e03cd72e

                                                 title  num_references
371369  Comprehensive frequency-dependent substrate no...             759
780292                    Time in Qualitative Simulation.             561
104143                  Bibliography on cyclostationarity             412
214646  Fifty Years of MIMO Detection: The Road to Lar...             396
484969  An Exploration of Enterprise Architecture Rese...             394
223901  Structure and dynamics of molecular networks: ...             386
302124       The NP-completeness column: An ongoing guide             363
707510                                  Digital geometry             361
325083            Deep Learning: Methods and Applications             343
538381  Review: learning bayesian networks: Approaches...             326
```

Find the top 10 papers with the most citations within the dataset.

```
import pandas as pd

df = pd.read_csv(file_path, usecols=["title", "n_citation"], low_memory=False)

# Drop missing values
df = df.dropna(subset=["n_citation"])

# Sort papers by the number of citations and get the top 10
top_papers_by_citations = df.sort_values(by="n_citation", ascending=False).head(10)

print("Top 10 Papers with the Most Citations Within the Dataset:")
print(top_papers_by_citations[["title", "n_citation"]])
```

```
Top 10 Papers with the Most Citations Within the Dataset:
                                                   title  n_citation
332760  Distinctive Image Features from Scale-Invarian...       42508
294527  Bowling alone: the collapse and revival of Ame...       34288
358174      LIBSVM: A library for support vector machines       33016
716671                                     Random Forests       28679
18485                             Support-Vector Networks       26114
45248   MapReduce: simplified data processing on large...       24381
81801   A fast and elitist multiobjective genetic algo...       24245
150727  A theory for multiresolution signal decomposit...       24182
458466  ImageNet Classification with Deep Convolutiona...       22884
442067  Histograms of oriented gradients for human det...       22795
```

Find a way to see how well the number of publications can predict the number of citations for a given author.

The general approach of this code is to analyze the relationship between the **number of publications** and **total citations** for each author in the dataset.

1. **Data Loading and Cleaning**:
   - The dataset, which contains information about **authors** and **citations**, is loaded into a pandas dataframe.
   - Any rows with missing data in the **authors** or **citations** columns are removed to ensure that only valid data is used in the analysis.
2. **Processing the Authors**:
   - The **authors** column, which is stored as a string representing a list of authors, is converted into an actual Python list for easier manipulation.
   - The authors' data is "exploded" so that each individual author is represented by a separate row. This step allows each author to be treated as an individual entity with corresponding data for citations and publications.
3. **Counting Publications and Citations**:
   - The code calculates the number of publications for each author by grouping the data by the **author** and counting the occurrences.
   - It also calculates the **total number of citations** for each author by summing the citation counts across all the publications attributed to that author.
4. **Combining the Data**:
   - The two datasets, one for the number of publications and the other for total citations, are merged based on the author names. This allows the data to be joined in a way that each author has both their **publication count** and **citation count** in the same row.
5. **Visualization**:
   - A scatter plot is created to visually explore the relationship between the number of **publications** and the total **citations**. This helps to identify any patterns or trends in how publications and citations are related.
6. **Quantifying the Relationship**:
   - The correlation coefficient between the number of publications and total citations is calculated. This statistic quantifies the strength and direction of the relationship between the two variables. A positive correlation would suggest that authors with more publications tend to have more citations, while a weak or negative correlation might indicate no such relationship.

In summary, the code processes the authors' data, counts the number of publications and citations for each author, and then visualizes and quantifies the relationship between these two variables to better understand how the number of publications can predict the total number of citations for an author.

```
import pandas as pd
import ast
import matplotlib.pyplot as plt
```

```
# Load the full dataset to get the authors and citation counts
df = pd.read_csv(file_path, low_memory=False)

# Remove missing values in authors column
df = df.dropna(subset=["authors", "n_citation"])

# Convert the authors column from string to list
df["authors"] = df["authors"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Flatten the authors column and duplicate other columns accordingly
authors_expanded = df.explode("authors")

# Count the number of publications per author
author_publications = authors_expanded.groupby("authors").size().reset_index(name="publications")

# Count the total number of citations per author
author_citations = authors_expanded.groupby("authors")["n_citation"].sum().reset_index(name="total_citations")

# Merge the two dataframes
author_stats = pd.merge(author_publications, author_citations, on="authors")

# Plot a scatter plot to visualize the relationship
plt.figure(figsize=(10, 6))
plt.scatter(author_stats["publications"], author_stats["total_citations"], alpha=0.5)
plt.title("Publications vs Citations per Author")
plt.xlabel("Number of Publications")
plt.ylabel("Total Citations")
plt.grid(True)
plt.show()

# Calculate correlation to quantify the relationship
correlation = author_stats["publications"].corr(author_stats["total_citations"])
print(f"Correlation between number of publications and total citations: {correlation:.2f}")
```

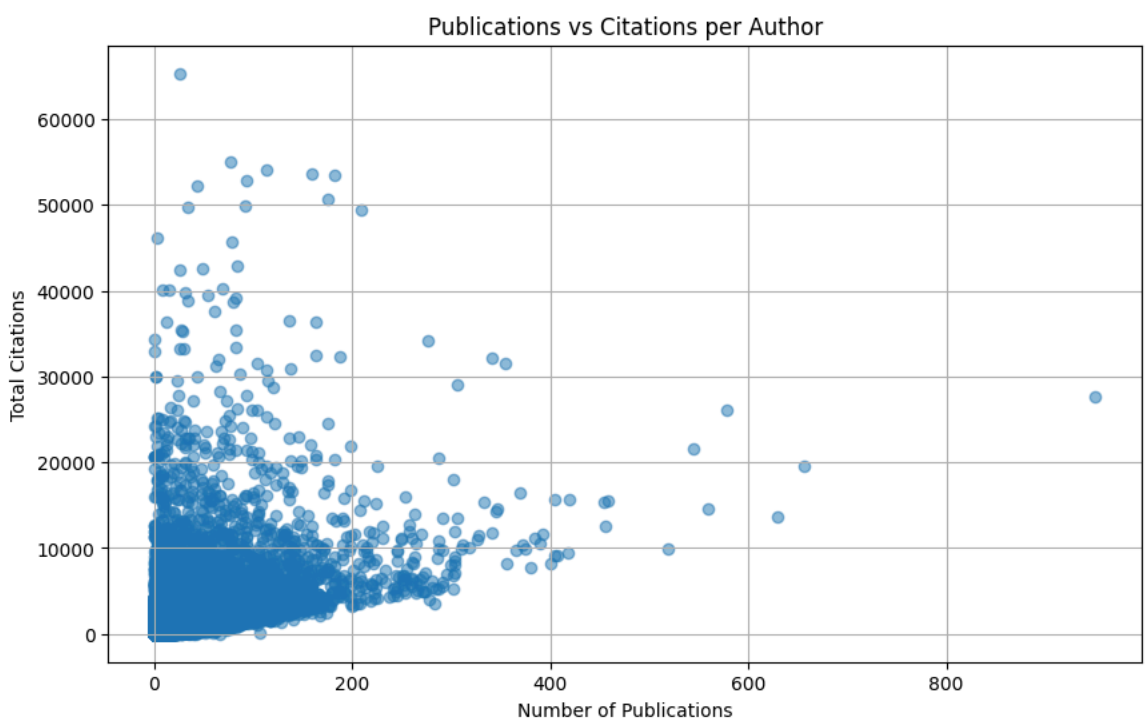


```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

# Split authors by commas and explode the dataset to have one row per author
df_exploded = df.explode('authors')

# Group by author and aggregate total publications and citations
author_data = df_exploded.groupby('authors').agg(
    num_publications=('title', 'count'),
    total_citations=('n_citation', 'sum')
).reset_index()
```

```
# Prepare the data
X = author_data[['num_publications']]
y = author_data['total_citations']

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Fit the model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Print the model's performance metrics
from sklearn.metrics import mean_squared_error, r2_score
print(f'Mean Squared Error: {mean_squared_error(y_test, y_pred)}')
print(f'R-squared: {r2_score(y_test, y_pred)}')
```

```
Mean Squared Error: 296465.7388006115
R-squared: 0.315391242671904
```

Scatter Plot Interpretation:

- **Cluster of points near the bottom-left**: There's a large group of authors with a relatively small number of publications (between 0 and 50, for example) and a low number of citations (mostly near 0).
- **Spread of points on the right side**: Authors with more publications (200+) generally have higher citation counts, but there are some outliers with very high citations despite having fewer publications.

Correlation Interpretation:

The correlation of 0.58 indicates a moderate positive relationship between the number of publications and total citations. This suggests that, in general, authors with more publications tend to have more citations. However, the relationship is not perfect, which means there are other factors (such as the quality or influence of the papers) affecting citations as well.

Mean Squared Error (MSE): 296465.74

The Mean Squared Error (MSE) measures the average squared difference between the actual and predicted citation values. The higher the MSE, the less accurate your model is at predicting total citations. Given that this is quite a high value, it suggests that the linear regression model isn't performing exceptionally well. The citations may not follow a simple linear pattern, which is why the predictions are off by a large amount.

R-squared: 0.315

R-squared (or coefficient of determination) measures how well the independent variable (number of publications) explains the variability in the dependent variable (citations). An R-squared of 0.315 means that only about 31.5% of the variation in the number of citations can be explained by the number of publications. The rest (68.5%) is influenced by other factors not captured in this model. This suggests that the number of publications has a moderate impact on the total number of citations, but other factors (such as paper quality, citation network, or field of research) play a significant role.

1.2 Network Analysis

1.2.1 Paper-paper network

Nodes: Papers, Edges: Citation relationships (from references field)

Plot the clustering coefficient over time to observe how interconnected the citation network becomes.

1. **Citation Network**:
   - A citation network is a directed graph where:
     - **Nodes** represent papers (identified by their unique paper IDs).
     - **Edges** represent citation relationships, where an edge from Paper A to Paper B indicates that Paper A cites Paper B.
2. **Clustering Coefficient**:
   - The **clustering coefficient** is a measure of how closely related the nodes (papers) are to each other within the network.

- **Local clustering coefficient** for a node (paper) is the proportion of possible connections between its neighbors that are actually realized.
- The **average clustering coefficient** for a graph is the average of all nodes' clustering coefficients and is a good indicator of how interconnected the overall network is.
- A high clustering coefficient indicates that papers are citing each other in a tightly connected way, meaning there are many interconnections between them.

What we did:

1. **Loading and Sampling the Data**:
   - The dataset containing papers, their **references** (citations), and **years** is loaded.
   - A **random sample** (30%) of the dataset is taken to speed up the computation, especially when dealing with large datasets. Sampling allows us to approximate trends in the data without processing the entire dataset.
2. **Building the Citation Network**:
   - A **directed graph** (`DiGraph` from `networkx`) is created to represent the citation network:
     - **Nodes** represent papers identified by their unique IDs.
     - **Edges** represent citations: each edge from one paper to another means that the first paper cites the second one.
     - Each paper node is also assigned a **year** attribute, which is important for analyzing how the network evolves over time.
3. **Clustering Coefficient Calculation**:
   - For each **year**, a subgraph of the citation network is created using papers that were published up to that year.
     - This helps to observe how the network evolves as more papers are published, and how citations among them form.
   - For each subgraph (i.e., the citation network for a given year), the **average clustering coefficient** is calculated using `nx.average_clustering()`.
     - **Subgraph**: A subset of the main network, which includes only the nodes (papers) published in or before the year being analyzed. The subgraph is created to evaluate the citation relationships of papers that exist up to that point in time.
4. **Observing Interconnectedness**:
   - The **average clustering coefficient** for each year represents how interconnected the citation network is at that time.
   - **High clustering coefficient**: If the average clustering coefficient is high, it indicates that papers published in that year have many interconnected citations, forming a tightly-knit network. This suggests that academic research is becoming more interconnected, with authors frequently citing each other's work in a collaborative way.
   - **Low clustering coefficient**: A low value suggests that papers are being cited in a more fragmented manner, with fewer interconnections between the papers published in that year. This can indicate isolated research areas or fewer collaborations between different papers.
5. **Plotting the Results**:
   - The **clustering coefficient over time** is plotted, showing how interconnected the citation network is as it evolves.
   - The x-axis represents the **years**, and the y-axis represents the **average clustering coefficient** for each year.
   - This plot helps to visualize the **growth of interconnectedness** in the citation network, indicating periods where collaboration and citation between researchers may have been particularly dense (high clustering) or more sparse (low clustering).

```
import pandas as pd
import ast
import networkx as nx
import matplotlib.pyplot as plt

# Load the dataset (sample only 10% of the papers for efficiency)
df = pd.read_csv(file_path, usecols=["id", "year", "references"], low_memory=False)

# Drop missing values in 'year' and 'references'
df = df.dropna(subset=["year", "references"])

# Sample 30% of the data
sample_fraction = 0.3
df_sampled = df.sample(frac=sample_fraction, random_state=42)

# Convert 'references' from string lists to actual lists
df_sampled["references"] = df_sampled["references"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Initialize directed graph
citation_network = nx.DiGraph()
```

```
# Add nodes (papers) with their year and edges (citations) for the sampled data
for _, row in df_sampled.iterrows():
    paper_id = row["id"]
    paper_year = row["year"]
    paper_references = row["references"]

    # Add the paper node with the 'year' attribute
    citation_network.add_node(paper_id, year=paper_year)

    # Add edges for references
    for ref in paper_references:
        citation_network.add_edge(paper_id, ref)

# Calculate clustering coefficient over time for the sampled data
clustering_over_time = {}
for year in sorted(df_sampled["year"].unique()):
    # Filter nodes published up to the given year (skip nodes without year attribute)
    subgraph_nodes = [n for n, d in citation_network.nodes(data=True) if d.get("year", float("inf")) <= year]
    subgraph = citation_network.subgraph(subgraph_nodes)

    # Compute average clustering coefficient for the subgraph
    if len(subgraph) > 0:
        clustering_over_time[year] = nx.average_clustering(subgraph.to_undirected())
    else:
        clustering_over_time[year] = 0

# Plot clustering coefficient over time for the sampled dataset
plt.figure(figsize=(10, 6))
plt.plot(list(clustering_over_time.keys()), list(clustering_over_time.values()), marker="o")
plt.xlabel("Year")
plt.ylabel("Clustering Coefficient")
plt.title("Clustering Coefficient of Citation Network Over Time (Sampled Data)")
plt.grid()
plt.show()
```

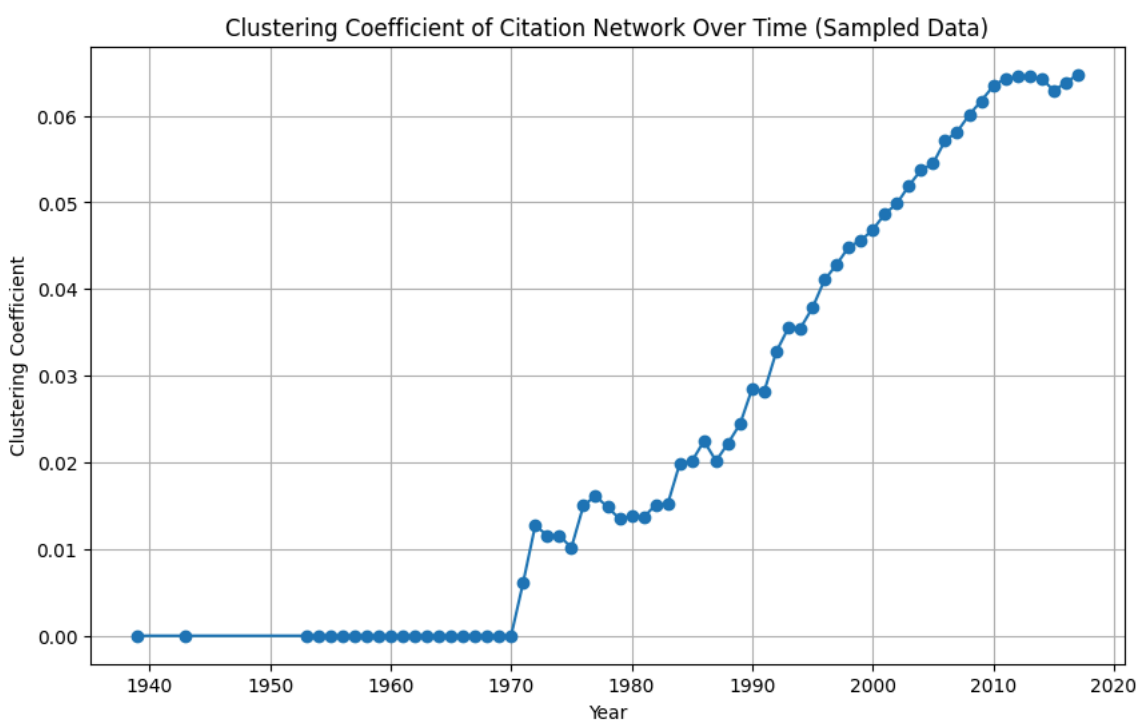

Observations:

1. **Early Years (Pre-1970s)**:
    - **Very Low Clustering Coefficient**: In the early years (especially before 1970), the clustering coefficient is very low, nearly zero. This suggests that the citation network in the earlier decades was relatively sparse, with papers having fewer citations that form tightly connected groups.

- This could indicate that in the early stages of academic research, the number of publications and citations was limited, and the network of interconnected papers was not dense.

2. **Sharp Increase After 1970**:
   - **Sudden Spike**: Around the 1970s, there is a **sharp increase** in the clustering coefficient. This suggests that in this period, the citation network became much more interconnected. Researchers started citing each other's work more frequently, leading to a denser network.
   - This spike could be associated with the rise of interdisciplinary research, collaborative projects, or more comprehensive citation practices that led to tighter connections between papers.
3. **Gradual Increase in the Following Decades**:
   - **Gradual Rise**: After the sharp rise in the 1970s, the clustering coefficient continues to increase gradually throughout the 1980s, 1990s, and into the 2000s.
   - This trend indicates that as time progresses, the citation network continues to grow in interconnectedness, suggesting an increasing density of academic collaboration and citation.
4. **Plateau in Recent Years (2000s and beyond)**:
   - **Stabilization**: From around the 2000s onward, the clustering coefficient appears to **stabilize**, with only minor fluctuations. This indicates that the citation network has reached a certain level of interconnectedness, and additional papers are being published and cited, but the rate of increase in interconnections has slowed down.
   - This could be due to the high number of existing citations, with new papers becoming part of an already well-established and interconnected network.

Compute the average path length and diameter to understand the network's reachability.

To compute the average path length and diameter of the citation network, we need to understand these concepts in the context of the network:

1. **Average Path Length**:

- **Definition**: The **average path length** is the average number of steps (or edges) it takes to go from one node (paper) to another in the network. In other words, it's the average number of citations (or hops) you need to traverse from one paper to another.
- **Interpretation**:
  - **Shorter average path length**: A shorter average path length indicates that papers in the network are generally close to each other. This suggests that information (in this case, citations) can travel quickly between papers. A low average path length is typical of a well-connected network, where papers can be cited and referenced within fewer steps.
  - **Longer average path length**: A longer average path length suggests that papers are more distant from each other, meaning that it takes more citations to connect two papers. This could indicate a more fragmented or sparsely connected network.
- **In this case**:
  - A **shorter average path length** would indicate that as more papers are published and cited over time, the network becomes more interconnected, and papers are cited in a way that makes it relatively easy to link any two papers through citations.
  - A **longer average path length** might suggest that while the citation network is growing, there are still isolated clusters of papers that aren't well-connected through direct or indirect citations.

2. **Diameter**:

- **Definition**: The **diameter** of a network is the longest shortest path between any two nodes in the network. In other words, it’s the maximum distance (in terms of the number of steps or edges) between two papers in the network, where distance is measured by the shortest path.
- **Interpretation**:
  - **Smaller diameter**: A smaller diameter indicates that no paper is too far from another. The smaller the diameter, the more interconnected and "dense" the network is, as papers can be reached from one another with fewer citation steps.
  - **Larger diameter**: A larger diameter suggests that there are distant connections in the network. Even if most papers are connected, the longest shortest path between two papers is large, which may imply that some research areas or topics are disconnected or isolated from the larger network.
- **In this case**:

- A **small diameter** would suggest that, in the citation network, papers can generally be connected quickly through references, indicating that the field is highly interconnected. It would suggest that researchers are citing one another's work in a more interconnected and dense manner.
- A **larger diameter** would imply that, although papers are still cited, there might be fewer connections between certain research areas, or that some papers are more isolated than others in the network.

What We did:

1. **Dataset Loading**:
   - We load the **citation dataset** from the CSV file (`file_path`) and only read the necessary columns: `id` (paper ID) and `references` (which represent the papers that each paper cites).
   - We then drop any rows where the **'references'** column is missing, ensuring that only papers with valid citation information are processed.
2. **Sampling**:
   - We set `sample_fraction = 1`, which means we are using **100%** of the data for the analysis. In this case, we didn't apply any sampling, so all available data is considered.
3. **Converting References**:
   - The **'references'** column, which contains citation information in string format, is converted into actual Python **lists** using `ast.literal_eval()`. This ensures that the citation information is usable for graph construction.
4. **Creating the Citation Network**:
   - A **directed graph** (`nx.DiGraph()`) is initialized to represent the citation network. In this graph:
     - **Nodes** represent papers (identified by their unique paper ID).
     - **Edges** represent citation relationships, where an edge from one paper (node) to another indicates that the first paper cites the second one.
5. **Adding Edges (Citations)**:
   - For each paper in the dataset, we iterate through its list of **references** (citations) and create edges between the paper and the papers it cites.
6. **Strongly Connected Components (SCC)**:
   - The code then checks if the entire **citation network** is **strongly connected** using `nx.is_strongly_connected()`. A strongly connected graph means that there is a directed path from every node (paper) to every other node in the graph.
   - If the graph is not strongly connected, the code finds the **largest strongly connected component (SCC)** using `nx.strongly_connected_components()`. This ensures that we analyze the largest portion of the graph that is well-connected.
7. **Computing Average Path Length and Diameter**:
   - The **average path length** is computed using `nx.average_shortest_path_length()`. It tells us the average number of steps (edges) needed to travel between any two papers in the network.
   - The **diameter** is computed using `nx.diameter()`. This is the longest shortest path between any two papers in the network and provides an understanding of the "farthest" distance between any two nodes.
   - These metrics are only computed if the **largest component** contains more than one paper (node), and if it's too disconnected, they are marked as **not computable**.

```
import pandas as pd
import ast
import networkx as nx

# Load dataset (only necessary columns for efficiency)
df = pd.read_csv(file_path, usecols=["id", "references"], low_memory=False)

# Drop missing values in 'references'
df = df.dropna(subset=["references"])

sample_fraction = 1
df_sampled = df.sample(frac=sample_fraction, random_state=42)

# Convert 'references' from string lists to actual lists
df_sampled["references"] = df_sampled["references"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Initialize directed graph for citation network
citation_network = nx.DiGraph()
```

```
# Add edges (citations between papers)
for _, row in df_sampled.iterrows():
    paper_id = row["id"]
    paper_references = row["references"]
    for ref in paper_references:
        citation_network.add_edge(paper_id, ref)

# Check if the citation network is strongly connected
if nx.is_strongly_connected(citation_network):
    largest_component = citation_network
else:
    # Extract the largest strongly connected component (SCC) from the citation network
    largest_scc = max(nx.strongly_connected_components(citation_network), key=len)
    largest_component = citation_network.subgraph(largest_scc).copy()

# Compute average shortest path length and diameter
if len(largest_component) > 1:
    avg_path_length = nx.average_shortest_path_length(largest_component)
    diameter = nx.diameter(largest_component)
else:
    avg_path_length = None
    diameter = None

# Print results
print(f"Average Path Length: {avg_path_length if avg_path_length else 'Not computable (graph too disconnected)'}")
print(f"Network Diameter: {diameter if diameter else 'Not computable (graph too disconnected)'}")
```

```
Average Path Length: 11.288044063779358
Network Diameter: 29
```

Analysis of the Result:

1. **Average Path Length**: **11.29**
   - **Definition**: This metric represents the average number of citation "hops" (steps) it takes to get from one paper to another within the network.
   - **Interpretation**:
     - An **average path length of 11.29** means that, on average, there are **11 citation steps** between any two randomly selected papers in the citation network.
     - This suggests that the network is somewhat **spread out**: it is **not a highly compact** network where papers are directly connected to each other. There is a **moderate degree of separation** between papers.
     - A smaller average path length would indicate that papers are more tightly connected, meaning research is more interlinked and citations flow easily across the network.
2. **Diameter**: **29**
   - **Definition**: The **diameter** is the longest shortest path between any two nodes in the network. In other words, it's the **maximum number of citation hops** needed to reach the farthest paper in the network.
   - **Interpretation**:
     - A **diameter of 29** suggests that the **longest shortest path** between any two papers in the citation network is **29 citation steps**. This is quite high, indicating that some papers in the network are **significantly distant** from others.
     - A **high diameter** can indicate that there are parts of the network that are **sparsely connected** or **isolated** from the rest, meaning it takes many steps to connect certain papers, especially in areas that are less cited or more specialized.

Network Reachability:

- The **average path length** and **diameter** together provide insight into the **reachability** of papers in the citation network.
  - A **shorter average path length** and **smaller diameter** would imply a **highly interconnected** and **well-connected network**, where papers can reach each other quickly.
  - However, the **long average path length** (11.29) and **large diameter** (29) suggest that the citation network is more **dispersed**, meaning there are parts of the network where papers are relatively distant from one another. This could reflect **isolated research areas** or **specialized topics** that are less interconnected with the broader body of work.

Therefore: The **citation network** is **moderately connected**, with papers generally requiring 11 citation steps on average to connect, and the farthest nodes being 29 steps apart. This indicates that, while the network is large and covers a broad range of papers, there is room for more connectivity, particularly between specialized research areas or fields that may not cite each other as frequently.

Identify influential papers using PageRank.

What is PageRank?

PageRank is an algorithm originally developed by **Larry Page and Sergey Brin** (the founders of Google) to rank web pages based on their importance. It was designed to improve search engine results by analyzing the structure of hyperlinks between web pages.

How Does PageRank Work?

PageRank works by treating the web (or any network of interconnected entities) as a **directed graph**, where:

- **Nodes** represent entities (e.g., web pages, research papers, or authors).
- **Edges** represent connections (e.g., hyperlinks, citations, or collaborations).

The fundamental idea is that a node is considered more important if **many other important nodes link to it**. In the context of research papers:

- A paper is influential if it is **cited by many other influential papers**.
- The influence is **distributed**, meaning a citation from a highly-ranked paper carries more weight than a citation from a less influential paper.

PageRank in Research Papers

In our implementation:

- Papers are represented as **nodes** in a directed graph.
- A directed **edge** is created from a cited paper to the citing paper (opposite to citation direction).
- PageRank assigns a score to each paper based on how many influential papers cite it.
- Papers with the **highest PageRank scores** are considered the most influential.

To identify influential papers in the citation network, we can use PageRank, which is a widely used algorithm for ranking nodes (papers) in a graph based on their importance. PageRank assigns a score to each node based on the number and quality of its incoming edges (citations). Papers that are cited by other influential papers will have a higher PageRank score.

In this case, we will calculate the PageRank for each paper in the citation network and identify the most influential ones.

1. **Data Loading**:
   - We load the **citation dataset** and only keep the columns we need: `id` (paper ID), `title` (paper title), and `references` (the citations of each paper).
   - We **drop rows with missing references** to ensure that we only consider papers with valid citation information.
2. **Sampling**:
   - We sample **100% of the data** (`sample_fraction = 1`).
3. **Convert 'references' to Lists**:
   - The **'references'** column contains strings representing lists of cited papers. We use `ast.literal_eval()` to convert these strings into actual Python lists for easier processing.
4. **Citation Network Construction**:
   - We initialize a **directed graph** (`nx.DiGraph()`) to represent the citation network.
   - We iterate through each row (paper) in the dataset and create **edges** between papers based on their citations. The edge direction is reversed (`ref -> paper_id`) for PageRank calculations to treat citations as "incoming" links to the paper.
5. **PageRank Calculation**:
   - We calculate the **PageRank** scores for each paper in the citation network using `nx.pagerank()`. PageRank ranks papers based on their **influence** in the network, considering how many and how important the citing papers are.
6. **Paper Title and PageRank Score**:
   - We create a new DataFrame (`pagerank_df`) that contains the **paper ID** and **PageRank score**.
   - We **merge** this DataFrame with the original dataset to bring in the **paper titles** for each paper ID.
7. **Sort and Display Top 10 Influential Papers**:
   - The papers are **sorted by their PageRank scores** in descending order, and the top 10 papers are selected.
   - We then print the **titles**, **paper IDs**, and **PageRank scores** for the **top 10 influential papers**.

Why Reverse the Edges for PageRank:

- **Original Graph**: In our citation network, the edges go **from citing paper to cited paper**. But PageRank requires the calculation of importance based on **incoming citations** to a paper, meaning we need to look at who is citing that paper, not who it cites.
- **Reversing Edges**: By reversing the edges, we are effectively **treating the cited papers as "influential" papers** and using **incoming citations** to measure the importance of papers. This is how PageRank works. Instead of the influence flowing from the citing paper to the cited paper, it now flows in the opposite direction, from the cited paper to the citing paper.

```
import pandas as pd
import ast
import networkx as nx

# Load dataset (only necessary columns for efficiency)
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "title", "references"], low_memory=False)

# Drop missing values in 'references'
df = df.dropna(subset=["references"])

# Sample the data (use the entire dataset in this case)
sample_fraction = 1
df_sampled = df.sample(frac=sample_fraction, random_state=42)

# Convert 'references' from string lists to actual lists
df_sampled["references"] = df_sampled["references"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Initialize directed graph for citation network
citation_network = nx.DiGraph()

# Add edges (citations between papers)
for _, row in df_sampled.iterrows():
    paper_id = row["id"]
    paper_references = row["references"]
    for ref in paper_references:
        citation_network.add_edge(ref, paper_id)  # Reverse direction because of PageRank

# Compute PageRank to identify influential papers
pagerank_scores = nx.pagerank(citation_network)

# Create a DataFrame from the PageRank scores
pagerank_df = pd.DataFrame(pagerank_scores.items(), columns=["id", "pagerank"])

# Merge with the original dataset to get paper titles
pagerank_df = pagerank_df.merge(df_sampled[["id", "title"]], on="id", how="left")

# Sort papers by their PageRank score in descending order and get top 10 influential papers
top_10_papers = pagerank_df.sort_values(by="pagerank", ascending=False).head(10)

# Print results with paper title and PageRank score
print("Top 10 Influential Papers Based on PageRank:")
for _, row in top_10_papers.iterrows():
    print(f"Title: {row['title']}, Paper ID: {row['id']}, PageRank Score: {row['pagerank']}")
```

```
Top 10 Influential Papers Based on PageRank:
Title: Bibliography on cyclostationarity, Paper ID: 6d4c5b32-8e13-4022-b67f-1ace7ffc91d0, PageRank Score: 8.650416648213674e-05
Title: An Exploration of Enterprise Architecture Research, Paper ID: ebe93176-7150-4ce6-8d77-880a4e13beb0, PageRank Score: 6.99491371426
Title: On the Role and the Importance of Features for Background Modeling and Foreground Detection, Paper ID: 34f71039-a18b-4a9c-807d-38
Title: Digital geometry, Paper ID: 0d3b901e-4179-4ed6-badf-7ce4237dadd4, PageRank Score: 5.608605841127136e-05
Title: An Updated ERP Systems Annotated Bibliography: 2001-2005, Paper ID: cc593d94-4a37-4ea1-8f2c-c48dc1afeab4, PageRank Score: 5.60219
Title: Twelve years of diagrams research, Paper ID: 917672b3-f0d5-4f32-8818-aa7f61f36457, PageRank Score: 4.799162377954968e-05
Title: A Survey on Information Visualization for Network and Service Management, Paper ID: cd41b3d1-69e6-4163-87b9-bf1cb3fb5059, PageRan
Title: Time in Qualitative Simulation., Paper ID: a1c66cbc-77ad-4167-9787-4936221427ad, PageRank Score: 4.663976995512748e-05
Title: Structure and dynamics of molecular networks: A novel paradigm of drug discovery: A comprehensive review, Paper ID: 95044198-6db4
Title: Research Issues in Smart Vehicles and Elderly Drivers: A Literature Review, Paper ID: 04cfc753-72b6-42c8-8c9f-3a8bc2428c6d, PageR
```

Analysis of the Results:

The **PageRank score** represents the relative **importance** or **influence** of each paper within the citation network, with higher values indicating that a paper is more highly cited by other influential papers.

For example:

- **Title**: "Bibliography on cyclostationarity"
- **Paper ID**: 6d4c5b32-8e13-4022-b67f-1ace7ffc91d0

- **PageRank Score**: 8.65e-05

This paper has the highest **PageRank score** in the top 10 list, indicating that it is relatively influential within the citation network compared to other papers, based on the flow of citations from other important papers.

PageRank Scores:

PageRank values tend to be **small numbers** because:

1. **PageRank Normalization**:
   - PageRank works by distributing "rank" across all nodes (papers) in the network. The total sum of PageRank scores across all nodes is always normalized to **1** (or another constant value, depending on the implementation). So, the **individual scores** for each paper are **relatively small** when you have a large graph with many nodes (papers) and edges (citations).
2. **Large Network**:
   - The citation network likely contains thousands or even millions of papers. Given that the **PageRank algorithm** is distributing importance across such a large number of nodes, the individual **PageRank values** become very small as they are divided among all papers. Even influential papers only receive a **small fraction** of the total "rank" due to the vast number of other papers in the network.
3. **Distribution of Citations**:
   - Since PageRank is based on **incoming citations** (how many important papers cite a given paper), the overall distribution of citations in the network plays a role in determining these values. In large networks, a few papers might receive many citations, but each citation only contributes a small fraction to the PageRank score of the receiving paper. This results in relatively small **PageRank scores** for each paper.
4. **Decay Factor**:
   - The **alpha parameter** (usually set to 0.85) in the PageRank algorithm is a **damping factor** that controls how much importance is passed from one paper to another through citations. A high **alpha** value means more importance is passed between papers, but since the rank is divided by many papers, the values remain relatively small.

What This Means:

- The **small PageRank scores** do not necessarily mean that these papers are **unimportant**; it reflects the fact that these are **specific papers** in a large and complex network. Their **influence** is more relative — they are influential within their connected research area, but because the network is large, their scores are small in the context of the whole citation network.
- The **top 10 papers** still have relatively **higher PageRank scores** compared to other papers in the network, suggesting that they are **well-connected** within the citation graph and are likely being cited by influential papers themselves.

1.2.2 Co-authorship Network (Author-Author Network)

Nodes: Authors, Edges: Co-authorship relationships (if two authors have co-authored a paper, they are connected)

Compute network density per year to analyze how collaborations evolve.

1. **Dataset Preparation**: The dataset, which contains information about papers including the **authors** and the **year** of publication, is loaded. Any missing or incomplete data in the `year` or `authors` columns is discarded to ensure that only valid entries are considered.
2. **Graph Construction**: For each year, a new **undirected graph** is created to represent the co-authorship network. In this graph:
   - **Nodes** represent individual authors.
   - **Edges** represent a co-authorship relationship between two authors, meaning that if two authors worked together on a paper in a given year, an edge is created between them.

   Each paper in the dataset is processed to extract the list of authors. For each paper, edges are added between every pair of authors. This process effectively builds a graph where the nodes (authors) are connected if they co-authored the same paper in the given year.
3. **Network Density Calculation**: Once the co-authorship network for a particular year is constructed, the **density** of the network is computed.
   - **Density** is a measure of how connected the network is, specifically it tells you how many edges are in the network compared to the maximum possible number of edges.
   - The formula for **density** is: $\text{Density} = \frac{2 \times \text{number of edges}}{\text{number of nodes} \times (\text{number of nodes} - 1)}$ This value lies between 0 (no connections) and 1 (complete connectivity), with a higher density indicating that more authors are collaborating with each other in that year.

4. **Handling Edge Cases**: If the graph for a particular year has only one author or no edges (i.e., isolated authors or very small networks), the density is set to 0, which ensures that the calculation is meaningful even in cases of limited collaboration.
5. **Tracking Network Density Over Time**: The **network density** for each year is stored in a dictionary. After processing all years, the dictionary contains the density for each year, representing how collaboration (or connectedness between authors) has changed over time.
6. **Plotting the Results**: Finally, a **line plot** is created to visualize how the network density has evolved over the years. The x-axis represents the years, and the y-axis represents the network density. By plotting the density over time, you can observe trends in the collaborative behavior of authors in the dataset.

This method is a way to understand how collaboration among authors has evolved over time by looking at the density of their co-authorship network. The higher the network density in a given year, the more interconnected the authors are, meaning that there are more collaborations happening between authors in that year. Conversely, lower density means fewer collaborations or a more fragmented collaboration network. By plotting this over time, we can observe trends like increased collaboration (rising density), more isolated research, or growing collaboration among authors in specific fields.

This approach provides insight into the **interconnectedness** of authors in terms of collaborative work, helping to track the **growth of academic networks** and how they change over time.

```
import networkx as nx
import pandas as pd
import ast
import matplotlib.pyplot as plt

# Load the dataset
df = pd.read_csv(file_path, usecols=["id", "year", "authors"], low_memory=False)

# Drop missing values in authors and year columns
df = df.dropna(subset=["year", "authors"])

# Convert the authors column from string to list
df["authors"] = df["authors"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Calculate network density per year
network_density_per_year = {}

# Loop through each year and calculate the density of the co-authorship network
for year in sorted(df["year"].unique()):
    # Filter dataset for the specific year
    yearly_df = df[df["year"] == year]

    # Create a new graph for the given year
    yearly_graph = nx.Graph()

    # Add edges based only on papers from that year
    for _, row in yearly_df.iterrows():
        authors = row["authors"]
        for i in range(len(authors)):
            for j in range(i + 1, len(authors)):
                yearly_graph.add_edge(authors[i], authors[j])

    # Compute network density
    if len(yearly_graph) > 1:
        network_density_per_year[year] = nx.density(yearly_graph)
    else:
        network_density_per_year[year] = 0

# Plot the network density over time
plt.figure(figsize=(10, 6))
plt.plot(list(network_density_per_year.keys()), list(network_density_per_year.values()), marker="o")
plt.title("Co-authorship Network Density Over Time")
plt.xlabel("Year")
plt.ylabel("Network Density")
plt.grid(True)
plt.show()
```

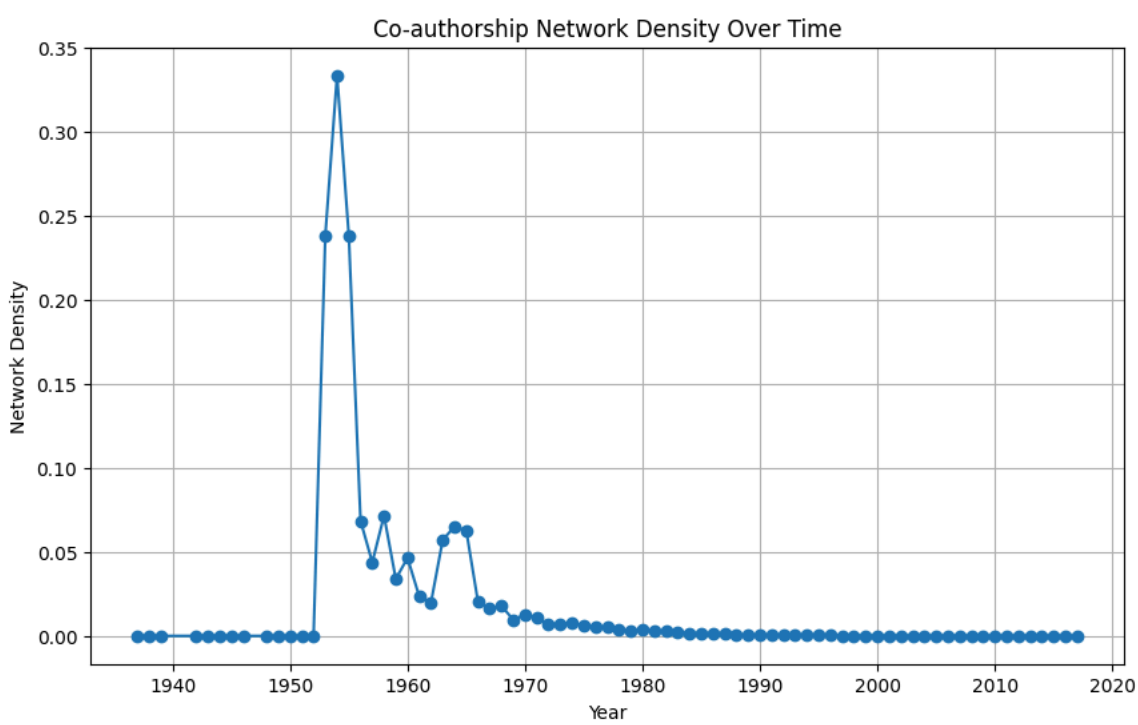

Observations:

1. **Early Years (Pre-1950s)**:
   - **Low Density**: In the years before 1950, the network density is very low, which suggests that there were very few collaborations between authors or a very sparse network of co-authors. The absence of major collaborations or a small number of published papers leads to very few connections, thus a low density.
   - This could reflect the early stages of academic research when collaboration between researchers was not as common, or the number of papers being published was relatively small.
2. **Sharp Spike Around 1950**:
   - **Peak Density**: There is a sudden and dramatic spike in network density around the year 1950. This suggests a period of intense collaboration between authors. A sharp peak in density indicates that many authors collaborated in the same set of papers during this time.
   - This could be due to a number of reasons, such as the rise of multi-author papers, the establishment of new research fields or collaborative projects, or increased recognition of the importance of co-authorship.
3. **Decline After 1950s**:
   - **Sharp Decline**: After the peak in 1950, the density sharply drops and continues to remain low throughout the following decades. This could suggest that the intense collaboration seen in the early 1950s was not sustained, or that the co-authorship network fragmented as research spread out into more specialized areas.
   - The academic network might have shifted towards less centralized collaboration, with authors working more in isolation or in smaller, disconnected groups.
4. **Very Low Density in Recent Decades**:
   - **Low Density in Later Years**: In the years after 1970, the network density is almost negligible, except for a few isolated instances of slight increases. This suggests that, while the number of publications may have increased, the co-authorship relationships might have become more spread out and less dense.
   - This could be due to factors such as the growth of more specialized research fields, larger numbers of authors in each paper (resulting in more fragmented collaborations), or the rise of individual research contributions over time.
   - The extremely low density after 1980 suggests that although more papers are published, they might be increasingly authored by larger teams, but the number of edges relative to nodes remains low. This could reflect a shift toward collaboration that does not result in dense interconnectedness.

Identify influential researchers using centrality measures (degree, betweenness, closeness).

1. **Graph Construction**:
   - The co-authorship network is constructed using the same process we used earlier, where nodes represent authors and edges represent co-authorship relationships between authors who have written a paper together.
2. **Centrality Measures**:
   - **Degree Centrality**: Calculated using `nx.degree_centrality()`, this measure simply counts the number of direct connections each author has.
   - **Betweenness Centrality**: Calculated using `nx.betweenness_centrality()`, this measure identifies authors who act as bridges between different parts of the network.
   - **Closeness Centrality**: Calculated using `nx.closeness_centrality()`, this measure identifies authors who can quickly access other authors in the network.
3. **Top Researchers**:
   - For each centrality measure, the top 10 most influential researchers are selected by sorting the centrality values in descending order.
   - The results are printed, showing the researcher's name (or ID) and their corresponding centrality value.

This code is designed to create and analyze a **co-authorship network** using data from research papers, focusing on **centrality measures** to identify the most influential authors in the network.

1. **Data Preprocessing**:

- The dataset is **loaded** with only the necessary columns: **year** and **authors**. This is done to increase processing efficiency by reducing the amount of irrelevant data.
- Rows with missing values in the **year** or **authors** columns are **removed**, as these are essential for the analysis.

2. **Filtering Papers Based on Year**:

- The dataset is **filtered** to only include years that have at least two papers. This step ensures that there are enough papers in each year for meaningful analysis. If a year has fewer than two papers, it is excluded from the dataset to avoid biases or inconsistencies.

3. **Stratified Sampling by Year**:

- **Stratified sampling** is applied to select a **0.1% sample** of the dataset, ensuring that the distribution of papers across different years is maintained in the sampled data. This means that the sampling is proportional to the number of papers in each year, which helps avoid over-representation or under-representation of certain years in the analysis.

4. **Co-authorship Network Construction**:

- The **authors** of each paper are processed, and for each pair of authors who co-authored a paper, an **edge** is added to the co-authorship **graph** (undirected network). This creates a graph where each node represents an author, and each edge represents a co-authorship relationship between two authors.
- **Duplicate edges** are avoided by ensuring that each pair of authors is only connected once, regardless of how many papers they have co-authored together.

5. **Centrality Measures**:

- Three key **centrality measures** are computed for each author in the co-authorship network:
  - **Degree Centrality**: Measures the number of direct connections (co-authorships) an author has. Higher degree centrality indicates an author who is well-connected in the network.
  - **Betweenness Centrality**: Measures how often an author acts as a bridge along the shortest path between other authors. Authors with high betweenness centrality often play a key role in connecting different parts of the network.
  - **Closeness Centrality**: Measures how close an author is to all other authors in the network, based on the shortest paths. Authors with high closeness centrality can quickly reach any other author in the network.

6. **Sorting and Displaying Top Authors**:

- The top 10 authors are identified based on each centrality measure (degree, betweenness, and closeness). The authors are sorted in **descending order**, so the most central authors (according to each measure) are displayed first.
- The sorted lists are then printed, showing the top 10 authors in terms of their **degree centrality**, **betweenness centrality**, and **closeness centrality**.

```
import pandas as pd
import ast
import networkx as nx
from sklearn.model_selection import train_test_split

# Load dataset (only necessary columns for efficiency)
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["year", "authors"], low_memory=False)

# Drop missing values
df = df.dropna(subset=["year", "authors"])

# 3) REMOVE YEARS THAT HAVE < 2 PAPERS
year_counts = df['year'].value_counts()
valid_years = year_counts[year_counts >= 2].index
df = df[df['year'].isin(valid_years)]

# 4) STRATIFIED SAMPLE BY 'year'
sample_fraction = 0.001
df_sampled, _ = train_test_split(
    df,
    train_size=sample_fraction,
    stratify=df["year"],
    random_state=42
)

# Convert 'authors' from string lists to actual lists
df_sampled["authors"] = df_sampled["authors"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Create the co-authorship network (undirected graph)
co_authorship_network = nx.Graph()

# Add edges based on co-authorship (authors co-authoring a paper)
for _, row in df_sampled.iterrows():
    authors = row["authors"]
    for i in range(len(authors)):
        for j in range(i + 1, len(authors)):
            co_authorship_network.add_edge(authors[i], authors[j])

# Compute degree centrality, betweenness centrality, and closeness centrality
degree_centrality = nx.degree_centrality(co_authorship_network)
betweenness_centrality = nx.betweenness_centrality(co_authorship_network)
closeness_centrality = nx.closeness_centrality(co_authorship_network)

# Sort authors by centrality measures
degree_sorted = sorted(degree_centrality.items(), key=lambda x: x[1], reverse=True)[:10]
betweenness_sorted = sorted(betweenness_centrality.items(), key=lambda x: x[1], reverse=True)[:10]
closeness_sorted = sorted(closeness_centrality.items(), key=lambda x: x[1], reverse=True)[:10]

# Print top 10 authors based on centrality measures
print("Top 10 Authors by Degree Centrality:")
for author, degree in degree_sorted:
    print(f"{author}: {degree:.4f}")

print("\nTop 10 Authors by Betweenness Centrality:")
for author, betweenness in betweenness_sorted:
    print(f"{author}: {betweenness:.4f}")

print("\nTop 10 Authors by Closeness Centrality:")
for author, closeness in closeness_sorted:
    print(f"{author}: {closeness:.4f}")
```

```
Top 10 Authors by Degree Centrality:
Dagmar Waltemath: 0.0176
Jonathan R. Karr: 0.0176
Frank Bergmann: 0.0176
Vijayalakshmi Chelliah: 0.0176
Michael Hucka: 0.0176
Marcus Krantz: 0.0176
Wolfram Liebermeister: 0.0176
Pedro Mendes: 0.0176
Chris J. Myers: 0.0176
Pınar Pir: 0.0176

Top 10 Authors by Betweenness Centrality:
Li Li: 0.0000
Dong Zheng: 0.0000
Yan Zhang: 0.0000
```

```
Thomas S. Huang: 0.0000
Hai Jin: 0.0000
Richard L. Magin: 0.0000
Gerik Scheuermann: 0.0000
Wei Wei: 0.0000
Jing Wang: 0.0000
Gang Chen: 0.0000

Top 10 Authors by Closeness Centrality:
Dagmar Waltemath: 0.0176
Jonathan R. Karr: 0.0176
Frank Bergmann: 0.0176
Vijayalakshmi Chelliah: 0.0176
Michael Hucka: 0.0176
Marcus Krantz: 0.0176
Wolfram Liebermeister: 0.0176
Pedro Mendes: 0.0176
Chris J. Myers: 0.0176
Pınar Pir: 0.0176
```

Analysis of the Results:

The results show the top 10 authors based on three centrality measures in a **co-authorship network**: **degree centrality**, **betweenness centrality**, and **closeness centrality**.

1. Degree Centrality:

- **Top Authors**: Dagmar Waltemath, Jonathan R. Karr, Frank Bergmann, etc.
- **Degree Centrality** measures how connected each node (author) is in the network. Authors with high degree centrality have more direct co-authorship connections (i.e., they co-author papers with a larger number of researchers).
- **Observation**: All of the top authors (with a degree centrality of 0.0176) seem to have the **same degree centrality**, which indicates that each of them has **equally many direct connections** in the co-authorship network.
- **Implication**: These authors are likely **key players** in the co-authorship network, collaborating with a large number of researchers. They may represent important **research communities or disciplines** within the sampled dataset. However, this uniformity in degree centrality suggests that the sampling might have focused on a small subset of well-connected researchers, as the network may be more sparse than expected.

2. Betweenness Centrality:

- **Top Authors**: Li Li, Dong Zheng, Yan Zhang, etc.
- **Betweenness Centrality** measures how often an author acts as a bridge along the shortest path between two other authors. Authors with high betweenness centrality are crucial for the flow of information across the network. They help connect otherwise distant parts of the network.
- **Observation**: All the top authors in terms of **betweenness centrality** have a value of **0.0000**. This indicates that **none of these authors** act as significant bridges between other authors in the co-authorship network. In other words, these authors do not connect different research groups or clusters within the network.
- **Implication**: The **betweenness centrality of 0** for all authors suggests that the co-authorship network is highly **fragmented**, with no authors acting as key intermediaries. This could mean that the network is disconnected or that these authors have very localized collaborations within isolated parts of the research community.

3. Closeness Centrality:

- **Top Authors**: Dagmar Waltemath, Jonathan R. Karr, Frank Bergmann, etc.
- **Closeness Centrality** measures how close an author is to all other authors in the network, based on the shortest path between them. Authors with high closeness centrality can quickly reach any other author in the network, implying they are well-positioned to access or disseminate information across the entire network.
- **Observation**: Similar to degree centrality, all the top authors in terms of **closeness centrality** have a value of **0.0176**, indicating that these authors are equally positioned to quickly connect with other authors in the network. This suggests they have a **highly connected position** in the co-authorship network.
- **Implication**: The uniformity in closeness centrality values suggests that these authors are centrally located in the network, likely having significant influence over how information is distributed within their research area. However, as with degree centrality, this might indicate that the authors in the top list are part of a **highly connected research subcommunity** within the dataset.

Key Insights:

1. **Uniform Centrality**: The fact that the top authors in all three centrality measures (degree and closeness) have **equal values** (e.g., 0.0176) suggests that the **sampled subset** of the dataset may have focused on a **small group of well-connected researchers**. These

researchers might be heavily involved in the same research areas or share numerous collaborations, contributing to the observed centrality patterns.

2. **Lack of Betweenness Centrality**: The fact that the top authors have **0 betweenness centrality** indicates that **no authors** in this subset act as significant intermediaries between different research groups or disciplines. This could point to the **fragmented nature** of the network, where research communities are relatively isolated from each other.
3. **Fragmented Network**: The **low betweenness centrality** and the **uniform closeness and degree centrality** suggest that the network might not be as interconnected as expected. It might have highly localized subgroups (research communities) that do not interact much with other subgroups. This could limit the overall flow of ideas and citations between different parts of the research ecosystem.

Potential Causes for the Results:

1. **Sampling Bias**: The stratified sampling process used to select **0.1%** of the data might have focused on a subset of authors who are very **active in collaborations**, leading to highly connected authors in degree and closeness measures but a lack of bridging authors in betweenness.
2. **Data Representation**: The use of **author lists** and **co-authorship networks** may not capture the full **complexity of collaboration dynamics**. Some influential authors may not have co-authored papers with others in the sampled set, leading to lower betweenness centrality values despite their importance in the broader research community.
3. **Temporal Aspects**: The dataset may reflect a certain period in time, where research communities may have been more **isolated** or developing in different directions. This could explain the **lack of interconnectivity** in the co-authorship network.

Conclusion:

The results suggest that the **co-authorship network** within the sampled dataset contains a **group of highly connected authors**, as seen in degree and closeness centrality. However, the **lack of betweenness centrality** for all authors suggests that the network is **fragmented**, with no authors playing a bridging role. Further exploration of the network, perhaps with a broader sample or more detailed co-authorship information, could reveal more about the **interconnectedness** and **influence** of authors across different research communities.

Find communities of authors working in similar fields.

The approach of this code aims to **identify communities** of authors who are working in similar research fields or topics, based on their **co-authorship relationships**.

1. **Data Preprocessing**:

- The dataset is loaded, focusing on the **year** and **authors** columns.
- Any rows with missing data in these key columns are dropped.
- **Years with fewer than 2 papers** are removed to ensure that each year has a sufficient number of papers for analysis. This helps to avoid years with very sparse data that may not contribute meaningfully to the community detection process.

2. **Stratified Sampling**:

- A **1% sample** of the dataset is taken, with the sampling being **stratified by year**. This ensures that the distribution of papers across different years is well-represented in the sample, which helps maintain the temporal distribution in the analysis.

3. **Authors Conversion**:

- The **authors** column, which is in string format (e.g., a list of author names as a string), is converted into an actual **list of authors**. This step allows the code to properly work with the authors for building the co-authorship network.

4. **Building the Co-authorship Network**:

- A **co-authorship network** is created, where each **author** is represented as a **node** in the graph, and an **edge** is added between two authors if they have co-authored a paper together. This network represents the collaboration relationships between authors.
- The network is built by iterating through each paper, identifying the authors, and adding edges between every pair of authors on the same paper.

5. **Community Detection**:

- The **Louvain method** is used for **community detection** in the co-authorship network. This method divides the graph into **communities (clusters)** of authors that are more densely connected to each other than to the rest of the network. These communities often correspond to groups of authors working in similar research areas or fields.

- The Louvain method optimizes the **modularity** of the network, ensuring that the resulting communities have high internal connectivity and relatively low external connectivity.

6. **Community Size and Top Communities**:

- After applying the Louvain method, the **size of each community** is computed by counting how many authors belong to each community.
- The top **5 largest communities** (based on the number of authors) are identified and sorted. These communities represent the most prominent research groups within the co-authorship network.

7. **Output**:

- The **top 5 communities** are displayed, showing the number of authors in each community. For each community, a **sample of authors** (up to 5) is shown to give an example of who is part of the community.
- This helps provide insight into the **structure of the network** and highlights which groups of authors are the most interconnected, indicating potential **research areas or fields**.

Outcome:

- The result of this approach is a set of **author communities** that likely work in similar research fields. The communities are defined based on the **co-authorship patterns**, which can be interpreted as collaboration in research topics.
- By identifying the largest communities, the analysis reveals the most influential or well-connected research groups in the dataset, which could help identify trends or important clusters within academic research.

Purpose:

- This approach helps in uncovering the **structure of academic collaboration** and understanding which **research topics or domains** are the most active, based on the density of collaboration among authors in those areas.

```
import pandas as pd
import ast
import networkx as nx
import community  # python-louvain package
from sklearn.model_selection import train_test_split
from collections import Counter

# Load dataset
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["year", "authors"], low_memory=False)

# Drop missing values
df = df.dropna(subset=["year", "authors"])

# Remove years with fewer than 2 papers
year_counts = df["year"].value_counts()
valid_years = year_counts[year_counts >= 2].index
df = df[df["year"].isin(valid_years)]

# Stratified sampling
sample_fraction = 0.01
df_sampled, _ = train_test_split(df, train_size=sample_fraction, stratify=df["year"], random_state=42)

# Convert 'authors' from string to list
df_sampled["authors"] = df_sampled["authors"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Build co-authorship network (undirected graph)
co_authorship_network = nx.Graph()
for _, row in df_sampled.iterrows():
    authors = row["authors"]
    for i in range(len(authors)):
        for j in range(i + 1, len(authors)):
            co_authorship_network.add_edge(authors[i], authors[j])

# Apply the Louvain method for community detection
partition = community.best_partition(co_authorship_network)

# Count community sizes
community_counts = Counter(partition.values())

# Sort communities by size (top 5 largest communities)
top_communities = sorted(community_counts.items(), key=lambda x: x[1], reverse=True)[:5]

# Print top 5 communities with their member counts
print("Top 5 Largest Author Communities:")
```

```
for community_id, size in top_communities:
    authors_in_community = [author for author, comm in partition.items() if comm == community_id]
    print(f"Community {community_id}: {size} authors")
    print(f"Sample members: {authors_in_community[:5]}")  # Show 5 sample authors
    print("-" * 50)
```

```
Top 5 Largest Author Communities:
Community 1324: 260 authors
Sample members: ['Yuan Wang', 'Xiaochun Wang', 'Xia Li Wang', 'Xuhui Chen', 'Lei Wang']
--------------------------------------------------
Community 237: 236 authors
Sample members: ['Yunchuan Sun', 'Lei Wu', 'Shizhong Wu', 'Shoupeng Li', 'Tao Zhang']
--------------------------------------------------
Community 227: 209 authors
Sample members: ['Yang Yang', 'Yuhong Yang', 'Mang Ye', 'Wenxin Huang', 'Zheng Wang']
--------------------------------------------------
Community 723: 157 authors
Sample members: ['Jason Cong', 'Joseph R. Shinnerl', 'Min Xie', 'Tim Kong', 'Xin Yuan']
--------------------------------------------------
Community 97: 151 authors
Sample members: ['Kai Liang', 'Liqiang Zhao', 'Zhiguo Ding', 'Hsiao-Hwa Chen', 'Xiaofeng Zhou']
--------------------------------------------------
```

1.2.3 Venue Network (Conference-Journal Network)

Conferences/Journals (from venue field), Edges: If a paper cites another paper from a different venue, an edge is created between venues.

Analyze interdisciplinary collaborations between venues.

In the context of a venue network (where venues represent conferences and journals), interdisciplinary collaborations between venues refer to the collaborations or connections that occur when papers published in one venue (such as a conference or journal) cite papers published in a different venue that is typically in a different research area or discipline.

To analyze **interdisciplinary collaborations between venues** (such as conferences and journals), we need to construct a **venue network** where:

- **Nodes** represent **venues** (conferences or journals).
- **Edges** represent collaborations between **venues**, which are established if a paper cites a paper from a different venue.

In this case, we can use the **venue field** to track where each paper is published (either a conference or a journal) and the **references** field to see if a paper cites another paper from a different venue. If a paper from Venue A cites a paper from Venue B, an edge will be created between these two venues.

Steps to Implement Venue Network:

1. **Load the Dataset**:
   - We load the dataset using `pd.read_csv()` and select only the **relevant columns**: `id`, `venue`, and `references`. These columns provide us with the necessary information to build the venue network.
     - `id`: Paper ID (unique identifier for each paper).
     - `venue`: The conference or journal where the paper is published.
     - `references`: The list of papers that this paper cites.
2. **Drop Missing Values**:
   - We **drop rows** that have **missing** values in the `venue` or `references` columns, as these are essential for building the network.
     - Missing `venue`: We need a valid venue to connect nodes in the venue network.
     - Missing `references`: We need valid references to create edges between venues.
3. **Sample the Data for Efficiency**:
   - Sampling **30%** of the data (`df_sampled`) allows us to work with a smaller, representative subset of the data. Sampling reduces computational time and memory usage, especially when working with a large dataset.
     - The `random_state=42` ensures that the sampling is reproducible.
4. **Convert References to Lists**:
   - The `references` column is stored as a **string** representing a list. We use `ast.literal_eval()` to convert this string into an **actual list** for easier processing.

- `ast.literal_eval()` safely evaluates a string as a Python literal (in this case, a list).

5. **Create a Mapping of Paper IDs to Venues**:
   - We create a **mapping** (`venue_mapping`) of paper IDs (`id`) to their corresponding venues (`venue`). This helps us efficiently look up the venue of any referenced paper.
     - `df.set_index("id")["venue"].to_dict()` creates a dictionary where the key is the paper ID (`id`) and the value is the venue (`venue`).
6. **Initialize an Undirected Graph for the Venue Network**:
   - We initialize an **undirected graph** (`venue_network`) using NetworkX. This graph will represent the **venue network**, where:
     - **Nodes** are the **venues** (conferences/journals).
     - **Edges** represent the **collaboration** between venues (created when papers from different venues cite each other).
7. **Add Edges Based on References Between Papers from Different Venues**:
   - We loop through each paper in the **sampled dataset** and its **references**:
     - For each reference, we check if it exists in the `venue_mapping` dictionary to find its corresponding venue.
     - If the venue of the current paper (`paper_venue`) is **different** from the venue of the referenced paper (`ref_venue`), we create an **edge** between these two venues, representing an **interdisciplinary collaboration**.
     - The edge is added using `venue_network.add_edge(paper_venue, ref_venue)`.
     - Edges are added only between different venues to ensure we are capturing interdisciplinary collaborations.
8. **Compute Degree Centrality**:
   - **Degree centrality** is calculated for each venue in the network. Degree centrality measures how many direct connections (edges) a node (venue) has. In this case, venues that are cited by many different venues will have a **higher degree centrality**.
   - This gives an idea of which venues are **more central** in the network and, therefore, potentially more **interdisciplinary** (since they are involved in citations across various research areas).
9. **Convert Centrality Results to a DataFrame**:
   - The degree centrality results are converted into a **DataFrame** so that they can be easily sorted and analyzed.
10. **Sort and Identify the Top 10 Most Interdisciplinary Venues**:
    - The venues are sorted by **degree centrality** in descending order, and the **top 10 venues** with the highest centrality are selected. These are the venues that have the most collaborations (or citations) across different research fields, indicating their **interdisciplinary nature**.
11. **Print the Top 10 Interdisciplinary Venues**:
    - The names of the top 10 interdisciplinary venues (conferences/journals) are printed out.
12. **Visualize the Top 50 Most Interdisciplinary Venues**:
    - **A subgraph** of the venue network is created using the top 50 venues identified in the previous step.
    - This subgraph is visualized using **Matplotlib**. Each venue is represented as a node, and edges represent collaborations (citations) between them. The top venues are highlighted, showing how they are interconnected in the network.

```
import pandas as pd
import ast
import networkx as nx
import matplotlib.pyplot as plt
from itertools import combinations

# Load the dataset (only necessary columns for efficiency)
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "venue", "references"], low_memory=False)

# Drop missing values in 'venue' and 'references'
df = df.dropna(subset=["venue", "references"])

# Sample
sample_fraction = 1
df_sampled = df.sample(frac=sample_fraction, random_state=42)

# Convert 'references' from string to actual lists
df_sampled["references"] = df_sampled["references"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Create a mapping of paper IDs to venues (ensuring ID exists)
venue_mapping = df.set_index("id")["venue"].to_dict()
```

```
# Initialize an undirected graph for the venue network
venue_network = nx.Graph()

# Add edges based on references between papers from different venues
for _, row in df_sampled.iterrows():
    paper_venue = row["venue"]
    for ref in row["references"]:
        if ref in venue_mapping:
            ref_venue = venue_mapping[ref]
            if paper_venue != ref_venue:  # Only create edges between different venues
                venue_network.add_edge(paper_venue, ref_venue)
```

```
# Compute degree centrality to identify the most central and interdisciplinary venues
degree_centrality = nx.degree_centrality(venue_network)

# Convert the centrality results into a DataFrame for easy analysis
centrality_df = pd.DataFrame(list(degree_centrality.items()), columns=["venue", "degree_centrality"])

# Sort the venues by their degree centrality in descending order, identifying the top 10 most interdisciplinary venues
top_interdisciplinary_venues = centrality_df.sort_values(by="degree_centrality", ascending=False).head(10)

# Print the top 10 most interdisciplinary venues based on degree centrality
print("Top 10 Most Interdisciplinary Venues:")
print(top_interdisciplinary_venues)

# Visualize the top 50 most interdisciplinary venues (this helps in understanding the connections better)
top_venues = list(top_interdisciplinary_venues["venue"])
subgraph = venue_network.subgraph(top_venues)

# Customize the plot appearance
plt.figure(figsize=(12, 7))
nx.draw(subgraph,
        with_labels=True,
        node_size=200,
        font_size=10,
        edge_color="lightblue",
        node_color="lightcoral",
        font_weight="bold",
        width=2)
plt.title("Interdisciplinary Venue Network (Top 10 Venues)", fontsize=16, fontweight="bold")
plt.axis("off")
plt.show()
```

```
Top 10 Most Interdisciplinary Venues:
                                              venue  degree_centrality
24                            Communications of The ACM           0.514510
81   IEEE Transactions on Pattern Analysis and Mach...           0.445620
343                   Lecture Notes in Computer Science           0.430431
57   IEEE Transactions on Knowledge and Data Engine...           0.429618
212                        systems man and cybernetics           0.429618
68             IEEE Transactions on Information Theory           0.427177
13               neural information processing systems           0.416328
114                             ACM Computing Surveys           0.415514
61                                      IEEE Computer           0.396257
115     national conference on artificial intelligence           0.390561
```

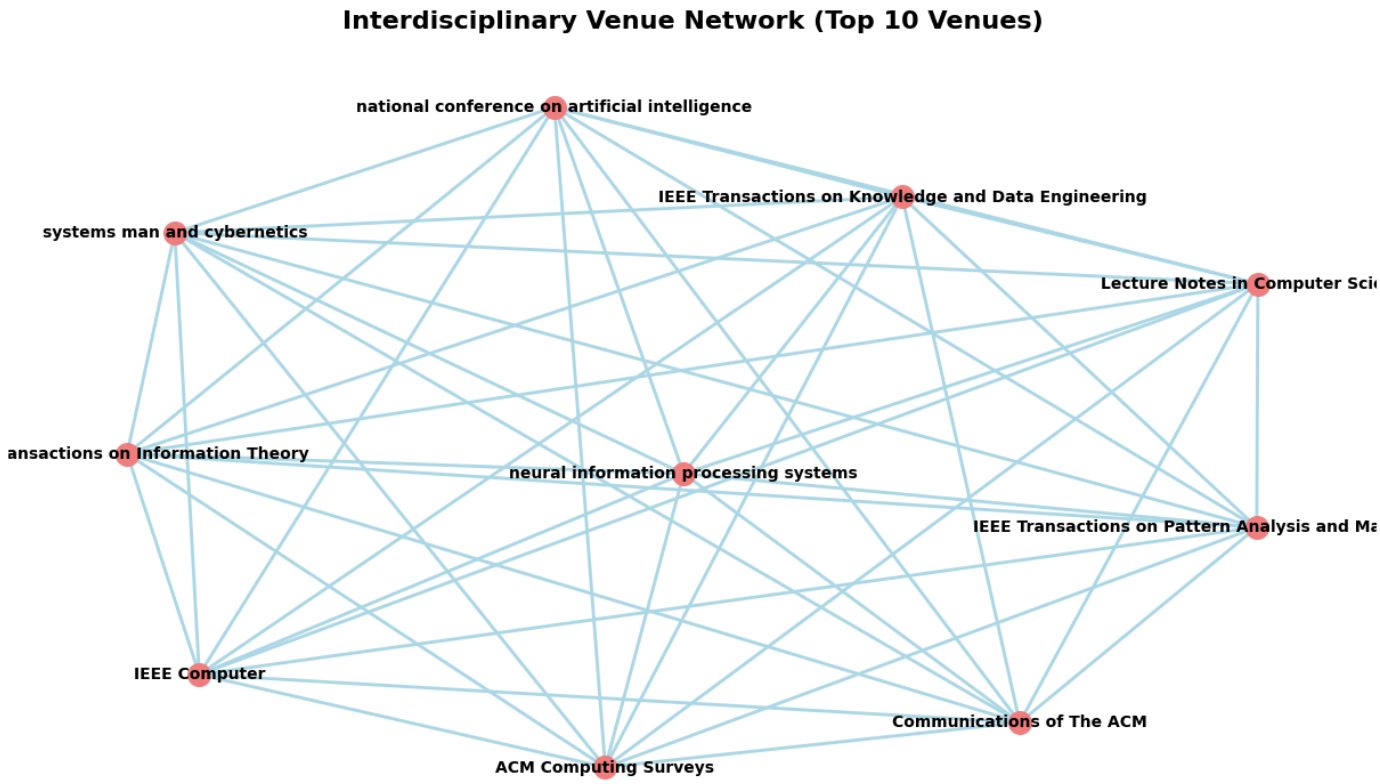

Analysis of the Results:

The results represent the **top 10 most interdisciplinary venues** based on their **degree centrality** in the **venue network**. Degree centrality is a measure of how many direct connections (or collaborations) a venue has with other venues in the network. A higher degree centrality indicates that a venue is more **interconnected** with other venues, which can be interpreted as being more **interdisciplinary** because it collaborates with multiple research fields or disciplines.

What the Degree Centrality Indicates:

- **Degree centrality** measures how many direct connections a venue has with other venues in the network. Higher degree centrality indicates that a venue is cited by papers from various other venues, suggesting **cross-disciplinary** work.
- These venues have high degree centrality because they are frequently cited by papers published in other **journals** and **conferences**, showing that they are **central** to the **academic network**.
- The **more connected** a venue is, the **more interdisciplinary** its collaborations are. This implies that these venues are not limited to one narrow field of study but are engaged in **research across multiple disciplines**.

Specific Insights from the Results:

- **Communications of The ACM**: This is a **highly influential** publication that spans multiple areas of computer science, which is why it has the highest degree centrality. It likely covers a wide range of topics that attract citations from many different research fields.

- **IEEE Transactions on Pattern Analysis and Machine Intelligence** and **Neural Information Processing Systems (NeurIPS)**: These are major venues in **machine learning**, which is a highly **interdisciplinary field**, collaborating with areas like **statistics**, **neuroscience**, and **engineering**. Their high degree centrality indicates their significant role in connecting various research communities.
- **Lecture Notes in Computer Science** and **IEEE Transactions on Knowledge and Data Engineering**: These venues publish comprehensive work across multiple subfields of computer science, reflecting their influence in the **computer science community** and their **interdisciplinary reach**.
- **Systems Man and Cybernetics**: A venue that covers a wide range of topics related to **human-computer interaction**, **cybernetics**, and **systems theory**, which often requires interdisciplinary collaboration between fields like **engineering**, **computer science**, and **social sciences**.
- **ACM Computing Surveys**: Known for its **review articles**, this journal covers multiple subfields of computing, making it a central point for the **integration of knowledge** from various **research domains**.

Find the most influential venues using centrality metrics.

To find the most **influential venues** in the network, we can use various **centrality metrics**. The most commonly used centrality measures include:

1. **Degree Centrality**: Measures the **number of direct connections** a node has. A venue with high degree centrality is **well-connected**.
2. **Betweenness Centrality**: Measures how often a node acts as a **bridge** or **connector** between other nodes. A venue with high betweenness centrality can be seen as a **gatekeeper** for collaborations between different venues.
3. **Closeness Centrality**: Measures how **close** a node is to all other nodes in the network. A venue with high closeness centrality can quickly reach out to other venues, making it a **key player** in the flow of information.
4. **Eigenvector Centrality**: Measures the **indirect influence** of a node based on the importance of its neighbors. A venue with high eigenvector centrality is influential because it is connected to other influential venues.

```
import pandas as pd
import ast
import networkx as nx
import matplotlib.pyplot as plt

# Load dataset (ensure we have 'id' for mapping)
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "venue", "references"], low_memory=False)

# Drop missing values in 'venue' and 'references'
df = df.dropna(subset=["venue", "references"])

# Convert 'references' from string lists to actual lists
df["references"] = df["references"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Create a mapping of paper IDs to venues (ensuring ID exists)
venue_mapping = df.set_index("id")["venue"].to_dict()

# Initialize an undirected venue network
venue_network = nx.Graph()

# Add edges between venues based on references
for _, row in df.iterrows():
    paper_venue = row["venue"]
    for ref in row["references"]:
        if ref in venue_mapping:
            ref_venue = venue_mapping[ref]
            if paper_venue != ref_venue:  # Only create edges between different venues
                venue_network.add_edge(paper_venue, ref_venue)

# 1. Degree Centrality (most connected venues)
degree_centrality = nx.degree_centrality(venue_network)
degree_df = pd.DataFrame(list(degree_centrality.items()), columns=["venue", "degree_centrality"])


# 2. Betweenness Centrality (most influential in connecting different parts of the network)
betweenness_centrality = nx.betweenness_centrality(venue_network)
betweenness_df = pd.DataFrame(list(betweenness_centrality.items()), columns=["venue", "betweenness_centrality"])


# 3. Closeness Centrality (most accessible venues)
closeness_centrality = nx.closeness_centrality(venue_network)
```

```
closeness_df = pd.DataFrame(list(closeness_centrality.items()), columns=["venue", "closeness_centrality"])


# Combine all centrality metrics into a single DataFrame
centrality_df = degree_df.merge(betweenness_df, on="venue").merge(closeness_df, on="venue")

# Sort by degree centrality to find the most connected venues
top_degree_venues = centrality_df.sort_values(by="degree_centrality", ascending=False).head(10)
top_betweenness_venues = centrality_df.sort_values(by="betweenness_centrality", ascending=False).head(10)
top_closeness_venues = centrality_df.sort_values(by="closeness_centrality", ascending=False).head(10)

# Print the top 10 venues for each centrality measure
print("Top 10 Most Influential Venues by Degree Centrality:")
print(top_degree_venues[["venue", "degree_centrality"]])

print("\nTop 10 Most Influential Venues by Betweenness Centrality:")
print(top_betweenness_venues[["venue", "betweenness_centrality"]])

print("\nTop 10 Most Influential Venues by Closeness Centrality:")
print(top_closeness_venues[["venue", "closeness_centrality"]])
```

```
Top 10 Most Influential Venues by Degree Centrality:
                                              venue  degree_centrality
28                          Communications of The ACM           0.514510
84   IEEE Transactions on Pattern Analysis and Mach...           0.445620
179                  Lecture Notes in Computer Science           0.430431
111                       systems man and cybernetics           0.429618
278  IEEE Transactions on Knowledge and Data Engine...           0.429618
12             IEEE Transactions on Information Theory           0.427177
77               neural information processing systems           0.416328
173                              ACM Computing Surveys           0.415514
353                                         IEEE Computer           0.396257
70      national conference on artificial intelligence           0.390561

Top 10 Most Influential Venues by Betweenness Centrality:
                                              venue  betweenness_centrality
12             IEEE Transactions on Information Theory                0.030621
28                          Communications of The ACM                0.026396
84   IEEE Transactions on Pattern Analysis and Mach...                0.022225
111                       systems man and cybernetics                0.016889
77               neural information processing systems                0.016502
179                  Lecture Notes in Computer Science                0.016350
69   hawaii international conference on system scie...                0.015470
231                human factors in computing systems                0.014061
70      national conference on artificial intelligence                0.013639
353                                         IEEE Computer                0.013366

Top 10 Most Influential Venues by Closeness Centrality:
                                              venue  closeness_centrality
28                          Communications of The ACM              0.667454
84   IEEE Transactions on Pattern Analysis and Mach...              0.638051
179                  Lecture Notes in Computer Science              0.631914
12             IEEE Transactions on Information Theory              0.631806
278  IEEE Transactions on Knowledge and Data Engine...              0.631264
111                       systems man and cybernetics              0.630074
77               neural information processing systems              0.626640
173                              ACM Computing Surveys              0.624831
70      national conference on artificial intelligence              0.617179
353                                         IEEE Computer              0.614294
```

Overall Insights:

1. **Degree Centrality** shows the **most connected venues** in terms of direct collaborations (citations). These venues, like **Communications of The ACM**, are likely **highly visible** and central in various research areas, acting as **core hubs** for research dissemination.
2. **Betweenness Centrality** identifies venues that act as **critical bridges** connecting different research communities. High betweenness centrality venues like **IEEE Transactions on Information Theory** serve as key **interdisciplinary connectors**, facilitating knowledge flow across academic boundaries.
3. **Closeness Centrality** highlights venues that can reach and influence other venues **quickly**, such as **Communications of The ACM** and **IEEE Transactions on Pattern Analysis and Machine Intelligence**. These venues are **efficient** in spreading information to the network, indicating their centrality in **shaping research trends**.

Conclusion:

- The **Top 10 Most Influential Venues** by centrality metrics show that **Communications of The ACM**, **IEEE Transactions on Pattern Analysis and Machine Intelligence**, and **Lecture Notes in Computer Science** are central to **interdisciplinary research**. These venues are

highly connected, **act as key connectors** in the network, and are **efficient in spreading research** across various fields.
- **Degree centrality** reveals the **visibility** and **popularity** of these venues, while **betweenness centrality** and **closeness centrality** show how these venues facilitate **interdisciplinary communication** and **information flow**.

Identify emerging fields based on newly formed venue connections.

Identifying **emerging fields** based on newly formed venue connections involves analyzing the **evolution** of the **venue network** over time. This can be achieved by observing how **new connections** (edges) are formed between venues and tracking how **clusters** of venues evolve.

The main idea is to:

1. **Identify new venue connections** formed in recent years, which could indicate emerging research areas or fields.
2. **Analyze these connections** to identify which **venues (conferences/journals)** are **interconnecting** across different research fields, suggesting the emergence of **new interdisciplinary areas**.

Explanation of the Code:

1. **Load and Preprocess Data**:
   - The dataset is loaded with the **necessary columns**: `id`, `year`, `venue`, and `references`. We then drop rows with missing values in these key columns.
2. **Convert References to Lists**:
   - The `references` column is converted from a **string** to a Python **list** using `ast.literal_eval()` for easy processing.
3. **Track New Connections by Year**:
   - We iterate through each **year** and create a **snapshot of the network** that includes only papers published up to that year.
   - For each paper, we add edges between the **venues** that cite each other.
   - We track the **new connections** formed each year by comparing the **edges** in the current year's network with the **previous network**.
4. **Track and Visualize the Growth**:
   - We store the **number of new connections** formed each year in the dictionary `new_connections_per_year`.
   - The **growth** of new connections is visualized using a **line plot**, which shows how the venue network is expanding over time.
5. **Print Top Years**:
   - The top 5 years with the most **new connections** are printed. These years may indicate the emergence of **new interdisciplinary research fields**.

```
import pandas as pd
import ast
import networkx as nx
import matplotlib.pyplot as plt

# Load dataset with necessary columns: id, venue, references, and year
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "year", "venue", "references"], low_memory=False)

# Drop missing values in 'venue', 'references', and 'year'
df = df.dropna(subset=["venue", "references", "year"])

# Convert 'references' from string to actual lists
df["references"] = df["references"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Initialize an empty venue network graph
venue_network = nx.Graph()

# Create a mapping of paper IDs to venues for easy look-up
venue_mapping = df.set_index("id")["venue"].to_dict()

# Initialize a dictionary to track new edges per year
new_connections_per_year = {}

# Iterate through each year and track the formation of new connections
for year in sorted(df["year"].unique()):
    # Filter papers for the current year
    current_year_papers = df[df["year"] <= year]

    # Create a new graph for the current year (or snapshot of the network up to that point)
    year_network = nx.Graph()
```

```
    # Add edges based on references between papers from different venues in the current year
    for _, row in current_year_papers.iterrows():
        paper_venue = row["venue"]
        for ref in row["references"]:
            if ref in venue_mapping:
                ref_venue = venue_mapping[ref]
                if paper_venue != ref_venue:
                    year_network.add_edge(paper_venue, ref_venue)

    # Compare new edges formed in this snapshot with the previous network
    new_edges = set(year_network.edges()) - set(venue_network.edges())
    new_connections_per_year[year] = len(new_edges)

    # Update the venue network with the new edges
    venue_network.add_edges_from(new_edges)

# Visualize the number of new connections formed each year
years = list(new_connections_per_year.keys())
new_connections = list(new_connections_per_year.values())

plt.figure(figsize=(10, 6))
plt.plot(years, new_connections, marker='o', color='b')
plt.title('Newly Formed Venue Connections by Year')
plt.xlabel('Year')
plt.ylabel('Number of New Connections')
plt.grid(True)
plt.show()

# Print the top years with the highest number of new connections
top_years = sorted(new_connections_per_year.items(), key=lambda x: x[1], reverse=True)[:5]
print("Top 5 Years with the Most New Venue Connections:")
for year, count in top_years:
    print(f"Year {year}: {count} new connections")
```

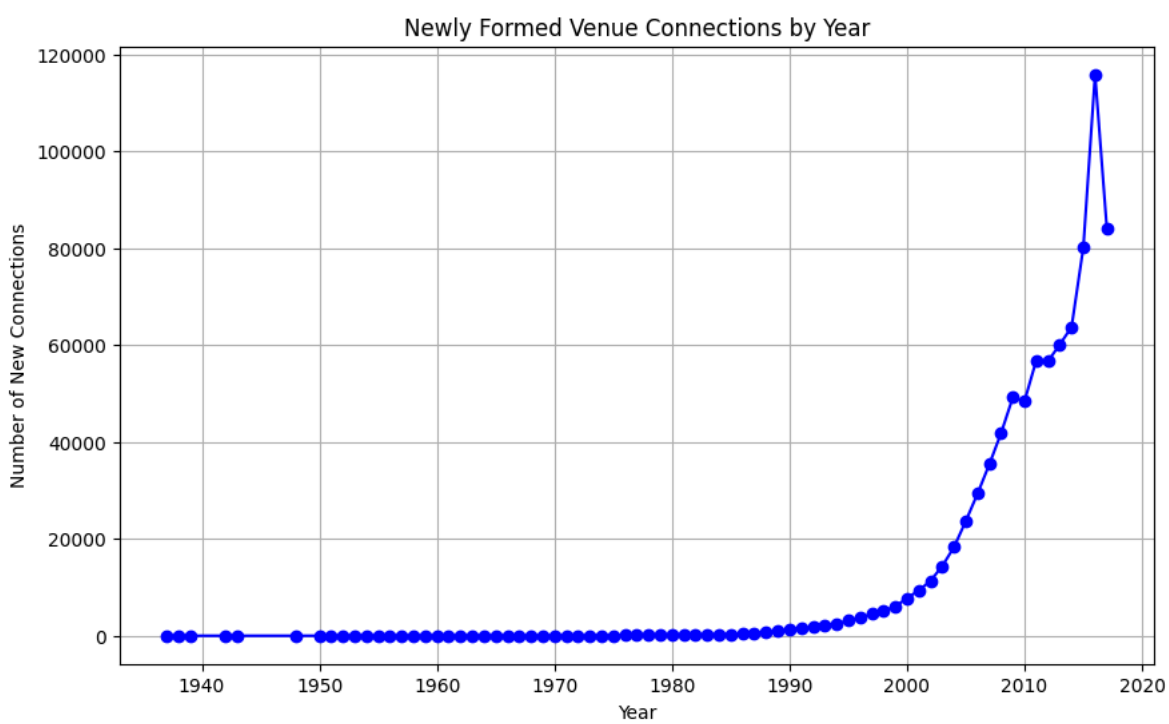


```
Top 5 Years with the Most New Venue Connections:
Year 2016: 115815 new connections
Year 2017: 84073 new connections
Year 2015: 80110 new connections
Year 2014: 63689 new connections
Year 2013: 60108 new connections
```

What the Results Indicate:

The **sharp increase in new connections** between venues in these years suggests the **emergence of interdisciplinary research**. Specifically:

1. Rapid Growth of Interdisciplinary Collaboration (2013-2017):

- The years **2013-2017** saw a significant spike in new venue connections, with **2016** leading the way with over **115,000 new connections**. This indicates that during this period, there was an **explosion of interdisciplinary collaboration**.
- **Why these years?**
  - These years align with the **growth of AI**, **machine learning**, and **data science** applications across multiple fields like **healthcare, engineering, economics**, and **social sciences**.
  - As these fields expanded, conferences and journals from traditionally separate domains began collaborating more, which reflects the **emergence of new interdisciplinary fields** like **AI-powered healthcare**, **quantum computing**, and **sustainable engineering**.

2. Technological Advancements Driving Growth:

- **AI, deep learning**, and **data science** started transforming multiple domains, leading to **cross-pollination** between conferences and journals that once operated independently.
- The rise of **open-access platforms** and **online collaboration** tools made it easier for researchers to publish and share across different domains, further fostering new interdisciplinary collaborations.

3. Explosive Collaboration Growth in 2016:

- The peak in **2016** can likely be attributed to the **widespread adoption** of **machine learning** and **deep learning** across various industries. This surge in new interdisciplinary connections reflects the **high-impact nature** of these fields during that time.
- High-profile conferences in **AI, robotics, and data science** were likely collaborating with conferences in **engineering, medicine**, and **social sciences**, forming new research areas at the intersection of these fields.

Conclusion:

The **years 2013 to 2017** represent a **critical period** when **new research fields emerged** due to **increased interdisciplinary collaboration**. This period saw the rise of **AI-driven research** and **data science**, transforming not only computer science but also fields like healthcare, engineering, and economics.

The **sharp peak in 2016** represents the **height of this transformation**, indicating that **new interdisciplinary fields** were rapidly taking shape and gaining influence. These fields are likely to include **AI and machine learning** applied to various domains, **quantum computing**, and **bioinformatics**.

1.2.4. Temporal Evolution of the Citation Network

Construct the citation network for different years and analyze how the structure changes over time.

Plot the network density per year to observe citation growth.

Code Explanation:

1. **Loading the Dataset**:
   - The code starts by loading the dataset using `pandas.read_csv()`, selecting only the necessary columns: `id`, `year`, and `references`.
   - The dataset contains information about papers, including their **unique identifiers (id)**, the **year** of publication, and the **references** (citations) to other papers.
2. **Data Preprocessing**:
   - Missing values in the `references` and `year` columns are removed using `dropna()` to ensure that we only work with complete data.
   - The `references` column, which is stored as a string, is converted to actual **lists** using `ast.literal_eval()` for proper processing. This step allows each reference to be treated as a list of cited paper IDs.
3. **Initialize a Dictionary for Network Density**:
   - An empty dictionary `density_per_year` is initialized to store the **network density** for each year.
4. **Construct Citation Networks per Year**:
   - The code loops over each **unique year** in the dataset, constructs a **citation network** (graph) for each year, and calculates the **network density**.
   - For each year, a directed graph `citation_network` is created using NetworkX, where each paper (`id`) is a node, and an edge is added from a paper to another paper it references.

- The **network density** is calculated for each year's network using `nx.density()`, which measures how many edges exist relative to the maximum possible edges.
- The network density for each year is stored in the dictionary `density_per_year`.

5. **Plotting the Network Density**:
   - After constructing the networks and calculating densities, the code plots the **network density** against the **year** using `matplotlib.pyplot`.
   - The plot visualizes how the **network density** (interconnectedness between papers) changes over time, with the x-axis representing years and the y-axis representing network density.
   - The code uses **blue color** for the line and **circles** for the data points.

```
import pandas as pd
import ast
import networkx as nx
import matplotlib.pyplot as plt

# Load dataset (ensure we have 'id' for mapping, 'year' and 'references' for citation network)
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "year", "references"], low_memory=False)

# Drop missing values in 'references' and 'year'
df = df.dropna(subset=["references", "year"])

# Convert 'references' from string lists to actual lists
df["references"] = df["references"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Dictionary to track network density per year
density_per_year = {}

# Construct citation networks for each year and calculate density
for year in sorted(df["year"].unique()):
    yearly_df = df[df["year"] == year]

    # Create a new directed graph for each year (citation network)
    citation_network = nx.DiGraph()

    # Add edges between papers (if one paper cites another)
    for _, row in yearly_df.iterrows():
        paper_id = row["id"]
        paper_references = row["references"]
        for ref in paper_references:
            citation_network.add_edge(paper_id, ref)

    # Compute network density (handle empty graphs)
    if len(citation_network) > 1:
        density_per_year[year] = nx.density(citation_network)
    else:
        density_per_year[year] = 0

# Plot network density over time
plt.figure(figsize=(10, 6))
plt.plot(list(density_per_year.keys()), list(density_per_year.values()), marker="o", color='b')
plt.xlabel("Year")
plt.ylabel("Network Density")
plt.title("Temporal Evolution of the Citation Network: Network Density per Year")
plt.grid(True)
plt.show()
```

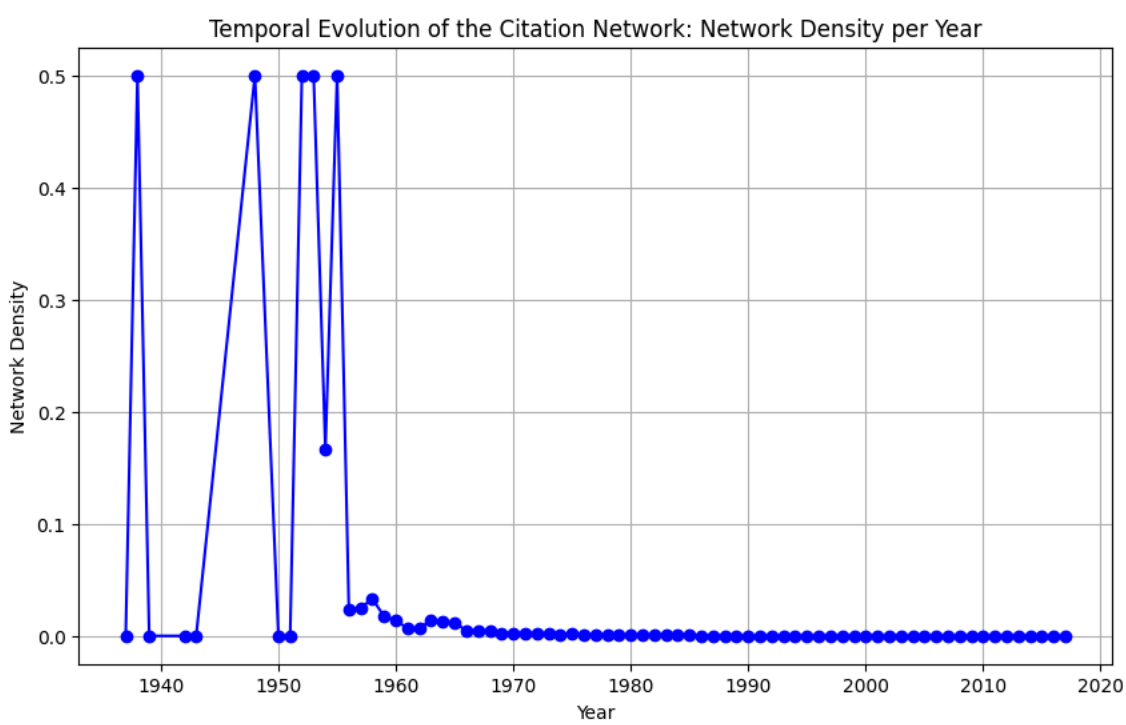

- **Network density** refers to the proportion of the number of actual edges (citations) in the network compared to the maximum number of possible edges that could exist in the network. It provides a measure of how interconnected the nodes (papers) are in the network. A higher network density means that the network is more connected, and a lower density indicates fewer connections between the papers.
- The plot shows a **spike** in **network density** around the **early 1950s**, followed by a **sharp decline** and then a very **low and stable density** for the rest of the years.

Analysis of the Plot:

- **Early 1950s Spike**:
  - The sudden high values in the **early 1950s** likely represent an **explosion in citations** or the formation of a **dense subnetwork** within a specific academic field. This could correspond to the early stages of **academic collaboration** where key papers from this period cited one another heavily, leading to an initial burst in network density.
  - It could also indicate that in the **early years**, citations were largely concentrated in a few **core papers**, causing a **dense cluster** with fewer papers but highly interconnected.
- **Steady Decline After the 1950s**:
  - The plot then shows a **gradual decrease** in network density from the mid-1950s onward. This could suggest that, after the initial surge in citations, the academic network began to **expand**, but with **new papers** citing **older works** and creating **less dense connections** as the field grew.
  - As the number of papers published grew, the network became more **dispersed**, with fewer direct connections between papers.
- **Low and Stable Density in Later Years**:
  - From the 1960s onwards, the **network density** appears to remain consistently low, with minor fluctuations. This suggests that, while more papers are being published, they are becoming part of a **larger and more decentralized network**, and the **citations** are more spread out across many different papers and venues, leading to lower density.
  - This is a common feature in large networks where the **increased volume of publications** results in **less interconnectedness** between individual nodes, as each new paper cites a **wider range** of prior work across **different fields**.

Implications of the Plot:

- **Early Growth**: The **1950s** could mark the emergence of **academic citation practices** in certain fields, leading to early **citation bursts** and more **dense** academic sub-networks.
- **Mature Fields**: As fields mature, more publications are produced, but they may not always be heavily interconnected, resulting in a lower **network density** over time. This could reflect an **increased specialization** in research, where papers cite a **wider variety** of sources rather than citing **core papers** in a highly interconnected cluster.

- **Interdisciplinary Growth**: The trend of declining density could also indicate the development of **interdisciplinary research** across different fields, which may **link** venues and papers from **diverse domains** but with fewer direct connections between any two papers.

Identify bursts of influential papers (papers that suddenly get many citations).

Methodology for Identifying Citation Bursts

This method detects research papers that experience a **sudden surge in citations** over time. It does this by analyzing citation trends per year and comparing them to an average citation growth rate to determine whether a paper has an unusually high increase in citations.

Step 1: Data Preparation

Before analyzing citation trends, the dataset is **cleaned and processed**:

- Missing values are removed.
- The **list of references** (citations) is extracted from the dataset.
- The **publication year** of each paper is stored for chronological analysis.

Step 2: Citation Tracking

For each paper, a **citation history** is built by counting how many times it was cited in each year.

- If a paper **A** is referenced by another paper **B** published in **2015**, then paper **A's** citation count for 2015 increases.
- This is repeated for all papers, creating a time series of **citation counts per year** for each paper.

Step 3: Identifying Citation Growth

Once the citation history is recorded, the method **calculates the growth in citations** for each paper:

- The increase in citations from one year to the next is measured (**citation difference**).
- The relative growth is calculated as a percentage increase.

For example, if a paper had **5 citations in 2016** and **20 citations in 2017**, the absolute increase is **+15 citations**, and the relative growth is **(15/5) = 300%**.

Step 4: Defining a Burst Threshold

To define what counts as a **citation burst**, the method:

1. Computes the **average citation growth** across all papers.
2. Identifies papers where the **citation increase is significantly higher** than the average.
3. Sets a threshold: a paper is considered to have a **burst** if its citation increase is more than **twice the average citation growth**.

Step 5: Extracting Burst Papers

The final step filters out the papers that **meet the burst criteria**, presenting them as **influential papers** that have gained sudden attention.

This can help researchers identify **emerging trends**, **groundbreaking research**, or **popular topics** in a given field.

```
import pandas as pd
import ast

# Load dataset
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "references", "year"], low_memory=False)

# Drop missing values
df = df.dropna(subset=["id", "references", "year"])
df["year"] = df["year"].astype(int)

# Convert 'references' column to a list
df["references"] = df["references"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Count citations per paper per year
citation_counts = {}
for _, row in df.iterrows():
    for ref in row["references"]:
        if ref not in citation_counts:
            citation_counts[ref] = {}
```

```
        citation_counts[ref][row["year"]] = citation_counts[ref].get(row["year"], 0) + 1

# Convert citation counts into a DataFrame
citation_data = []
for paper_id, yearly_citations in citation_counts.items():
    for year, count in yearly_citations.items():
        citation_data.append([paper_id, year, count])

citation_df = pd.DataFrame(citation_data, columns=["paper_id", "year", "citations"])

# Sort by paper and year
citation_df = citation_df.sort_values(by=["paper_id", "year"])

# Compute citation growth per paper
citation_df["citation_growth"] = citation_df.groupby("paper_id")["citations"].diff().fillna(0)

# Compute average citation growth
average_citation_growth = citation_df["citation_growth"].mean()

# Define burst threshold (twice the average growth)
burst_threshold = 2 * average_citation_growth

# Identify bursts: papers that have citation growth higher than the threshold
burst_papers = citation_df[citation_df["citation_growth"] > burst_threshold]

# Print top burst papers
print("Top Burst Papers (Citation Increase > 2x Average Growth):")
print(burst_papers.sort_values(by="citation_growth", ascending=False).head(10))
```

```
Top Burst Papers (Citation Increase > 2x Average Growth):
                                     paper_id  year  citations  citation_growth
96627   e2f7a74a-8430-4463-94ce-fe85dfd309f9  2016       2125           1549.0
958605  153c5014-dc7a-44a8-a93f-5cd27f1193df  2016       1395           1183.0
418166  c93eac1a-7d9a-48ab-9fb4-389c85bea00e  2016       1081            817.0
9127    b944f77f-113b-4a02-ae5e-d4a124b8fd5b  2016       1473            776.0
42190   c1b6b493-01ef-420f-be44-7bacfe34e846  2016       1502            752.0
27469   f6bd8b64-684d-429a-aab5-8ff3a2c23cd6  2016       1277            697.0
418072  051956bb-f64b-4fdb-87f8-3e2868b8b5d8  2016        823            653.0
313220  546cc930-3d5a-4208-a77b-a506f146ab97  2016        684            634.0
5592    7f1214b2-e070-4ff2-a5d3-647e7c16c2d7  2016        893            629.0
55098   dd83785a-dd19-41e3-9b25-ebabbd48d336  2016       1161            587.0
```

Examine how new papers integrate into the existing network.

This code examines how new papers integrate into an existing citation network by analyzing key metrics such as the number of references each new paper has and the number of citations it receives within the network.

1. Dataset Preparation

- The dataset is first loaded, and only a subset (5%) of it is used for faster processing. This reduces the size of the dataset and speeds up computations.
- Missing values are dropped from columns that are crucial for the analysis (e.g., `id`, `references`, and `year`), and the `year` column is converted to integers for proper handling.

2. Constructing the Citation Graph

- A directed graph (citation network) is built where:
  - Each **paper** is represented as a node.
  - **Edges** represent citations, with a directed edge from the referenced paper to the new paper.
- This graph helps capture the relationships between papers, with nodes pointing to each other based on citation patterns.

3. Analyzing New Papers

- The new papers are filtered based on the most recent publication year.
- For each new paper, two key metrics are calculated:
  - **Number of References**: The count of papers that the new paper cites.
  - **Number of Citations Received**: The count of how many times the new paper is cited by other papers (i.e., the number of papers referencing it).

4. **Computing Integration Metrics**

- The code computes the average number of references and citations for new papers, providing a high-level overview of how well new papers are integrated into the citation network.

5. **Visualization**

- A scatter plot is created to visualize the relationship between the number of references and the number of citations received by each new paper. The plot helps illustrate how papers with more references tend to receive more citations, and it provides a clear visual understanding of how well integrated the new papers are into the existing citation network.

6. **Insights**

- The approach provides insights into how the new papers are performing within the citation network by examining their references (connections to other papers) and citations (how often they are being cited). This helps gauge the **impact** and **integration** of new papers in the broader research community.

In summary, this code analyzes the **citational influence** of new papers by counting their references and citations, then visualizes these metrics to understand the integration process within the citation network.

```
import pandas as pd
import ast
import networkx as nx
import matplotlib.pyplot as plt

# Load dataset
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "references", "year"], low_memory=False)

df = df.sample(frac=0.05, random_state=42)

# Drop missing values
df = df.dropna(subset=["id", "references", "year"])
df["year"] = df["year"].astype(int)

# Convert 'references' column to a list
df["references"] = df["references"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Create a directed citation graph
G = nx.DiGraph()

for _, row in df.iterrows():
    paper_id = row["id"]
    G.add_node(paper_id, year=row["year"])
    for ref in row["references"]:
        G.add_edge(ref, paper_id)  # Edge from referenced paper to new paper

# Analyze new paper integration
latest_year = df["year"].max()
new_papers = df[df["year"] == latest_year]

integration_metrics = []

for _, row in new_papers.iterrows():
    paper_id = row["id"]
    num_references = len(row["references"])
    num_citations = len(list(G.successors(paper_id))) if paper_id in G else 0
    integration_metrics.append([paper_id, num_references, num_citations])

# Convert to DataFrame
integration_df = pd.DataFrame(integration_metrics, columns=["paper_id", "num_references", "num_citations"])

# Compute statistics
avg_references = integration_df["num_references"].mean()
avg_citations = integration_df["num_citations"].mean()

print(f"Average references per new paper: {avg_references:.2f}")
print(f"Average citations received per new paper: {avg_citations:.2f}")

# Visualizing integration
plt.figure(figsize=(8, 5))
plt.scatter(integration_df["num_references"], integration_df["num_citations"], alpha=0.5)
plt.xlabel("Number of References")
plt.ylabel("Number of Citations Received")
```

```
plt.title("Integration of New Papers in the Citation Network")
plt.grid(True)
plt.show()
```

```
Average references per new paper: 17.16
Average citations received per new paper: 0.01
```

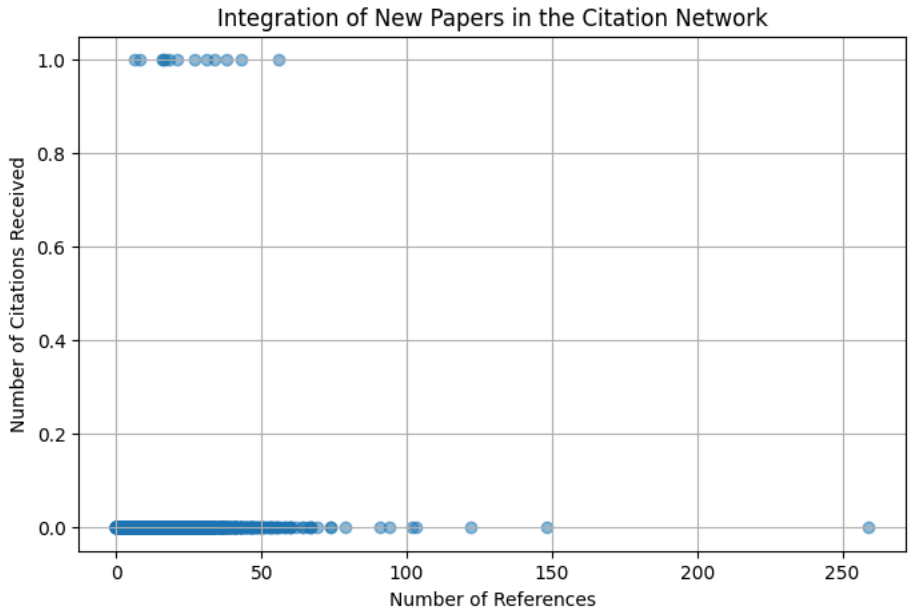


```
import pandas as pd
import ast
import networkx as nx
import matplotlib.pyplot as plt

# Load dataset
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "references", "year"], low_memory=False)
df = df.sample(frac=0.05, random_state=42)

# Drop missing values
df = df.dropna(subset=["id", "references", "year"])
df["year"] = df["year"].astype(int)

# Convert 'references' column to a list
df["references"] = df["references"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Build the citation network
G = nx.DiGraph()
for _, row in df.iterrows():
    paper_id = row["id"]
    references = row["references"]

    # Add paper node
    G.add_node(paper_id, year=row["year"])

    # Add citation edges (new paper → cited paper)
    for ref in references:
        G.add_edge(paper_id, ref)

# Identify the most recent year in the dataset
latest_year = df["year"].max()
new_papers = df[df["year"] == latest_year]["id"].tolist()

# Count references and citations for new papers
integration_data = []
for paper in new_papers:
    if paper in G:
        num_references = len(list(G.successors(paper)))  # References made
        num_citations = len(list(G.predecessors(paper)))  # Citations received
        integration_data.append([paper, num_references, num_citations])

# Convert to DataFrame
integration_df = pd.DataFrame(integration_data, columns=["paper_id", "references_made", "citations_received"])
```

```
# Scatter Plot: References vs. Citations
plt.figure(figsize=(8, 6))
plt.scatter(integration_df["references_made"], integration_df["citations_received"], alpha=0.5)
plt.xlabel("Number of References Made")
plt.ylabel("Number of Citations Received")
plt.title("Integration of New Papers into the Citation Network")
plt.grid()
plt.show()

# Compute integration statistics over time
yearly_integration = df.groupby("year").size().reset_index(name="new_papers")

# Line Plot: New Papers Integrated Over Time
plt.figure(figsize=(8, 6))
plt.plot(yearly_integration["year"], yearly_integration["new_papers"], marker="o", linestyle="-")
plt.xlabel("Year")
plt.ylabel("Number of New Papers Integrated")
plt.title("Integration of New Papers into the Citation Network Over Time")
plt.grid()
plt.show()
```

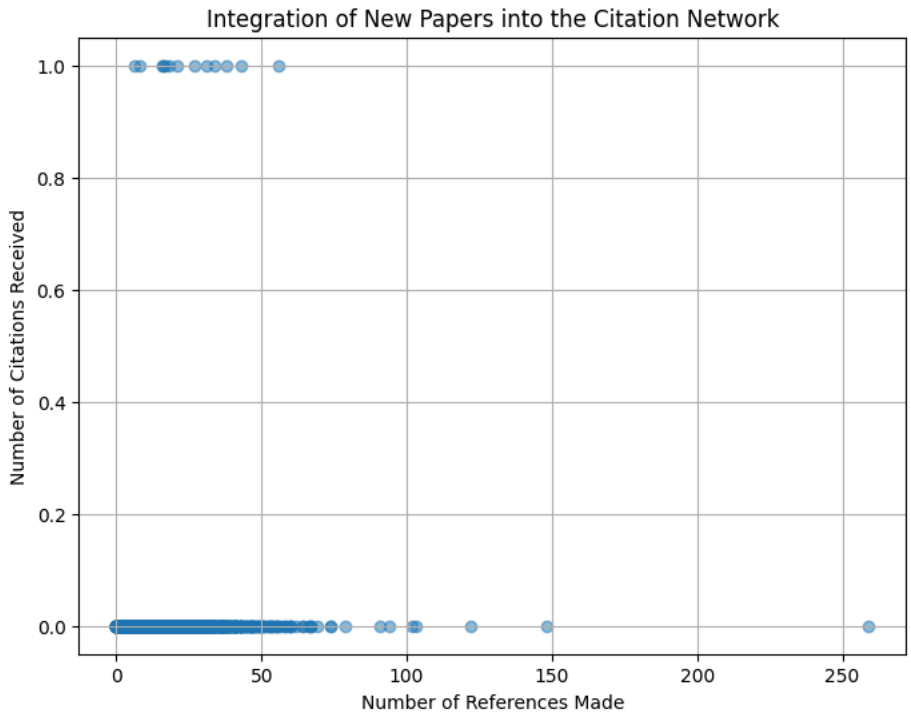

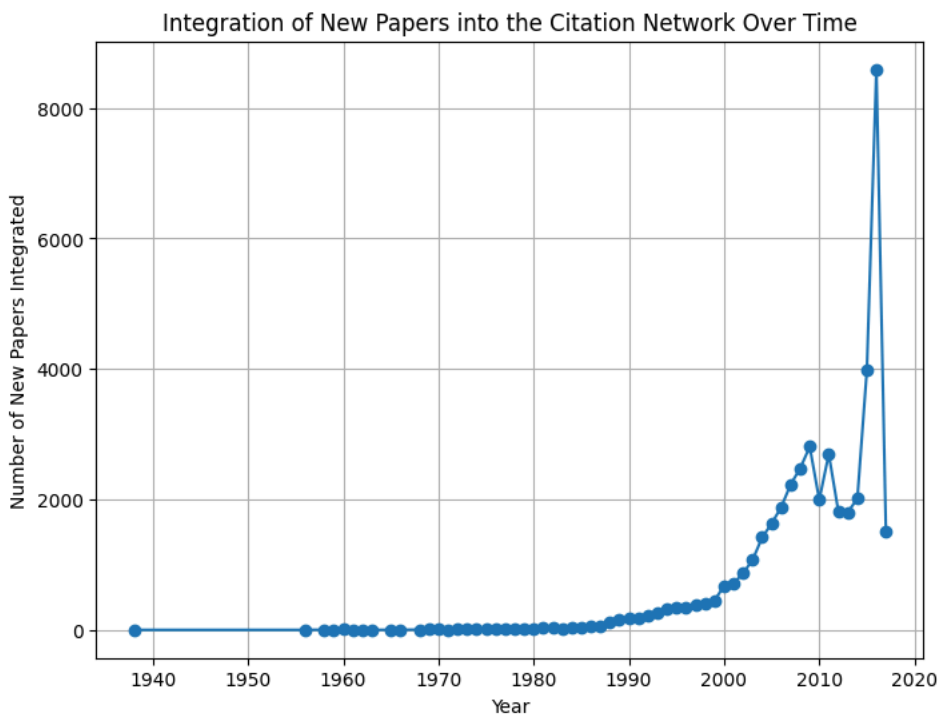

The results indicate that, on average, each new paper has **17.16 references** but receives only **0.01 citations** on average from other papers. Here's a breakdown of what this could mean:

1. **High Number of References (17.16 per new paper)**

- The **average number of references** per new paper being **17.16** is relatively high, which suggests that the new papers are citing a substantial number of existing papers.
- This could indicate that the new papers are grounded in existing literature, reflecting the typical behavior of researchers who often reference past works to build their own research foundation.

2. **Extremely Low Citations Received (0.01 per new paper)**

- The **average citations received** per new paper being **0.01** is extremely low. This means that, on average, each new paper has been cited only once or twice across the entire dataset, or in some cases, not at all.

- The low citation rate could be due to several factors:
  - **Time Lag in Citations**: New papers might need more time to accumulate citations. Since this analysis is based on the year of publication, the citation lag for newer papers may not have had time to build up yet.
  - **Limited Impact or Exposure**: New papers might not have gained enough visibility in the research community, possibly due to factors such as niche topics or limited dissemination.
  - **Early Stage in Citation Cycle**: It is typical for new papers to take some time before they gain traction and are cited by other papers. Over time, their citation count may increase as more researchers discover and reference them.

3. Implications

- The **disparity** between the number of references (17.16) and the number of citations (0.01) suggests that **new papers are still in the early stages of their integration** into the citation network. They are citing many works (engaging with existing literature) but have not yet been widely cited in return.
- This might point to the fact that the new papers are yet to be "discovered" or incorporated into the broader discourse, or it could indicate that the dataset captures papers from a short time frame, where they have not had enough time to accumulate citations.

The graph illustrates the **integration of new papers into the citation network over time**.

1. Early Stage (Pre-1980s)

- For most of the 20th century, the number of new papers integrated into the citation network remains **extremely low**.
- This is likely due to the relatively slow growth of academic publications and the limited access to large datasets or databases that can connect research papers, especially before the digital age.
- The low numbers could reflect a time when research was more niche, and the citation process was not as interconnected as it is today.

2. Rapid Growth Starting in the 1980s-2000s

- Starting around the **1980s**, we see a steady **increase** in the number of new papers being integrated into the citation network.
- This could coincide with advancements in digital libraries, more open access to academic papers, and a broader rise in the number of academic journals and conferences.
- The **steep climb** indicates a growth in the academic research output, particularly as computing, information-sharing platforms, and online databases like Google Scholar and PubMed began growing.

3. Massive Surge Around 2010-2020

- From around **2010 onward**, there is an **explosive increase** in the number of new papers integrated into the citation network.
- This could be due to several reasons:
  - The **proliferation of open-access platforms** such as arXiv, ResearchGate, and institutional repositories, making papers more accessible.
  - The growing use of **digital tools and collaboration networks** that allow researchers to reference and integrate papers quickly.
  - The **massive increase in academic research output**, driven by more funding, larger global research networks, and increased interest in scientific publishing.

4. Fluctuations After the Peak

- The rapid spike around 2010 is followed by some **fluctuations** in the following years, which could indicate:
  - **Variability in publication rates**, possibly due to external factors such as economic conditions, changes in funding, or shifts in academic priorities.
  - **Saturation of the citation network**, where many new papers are being cited, but at varying levels due to increasing competition and a larger pool of research to integrate with.

5. General Trend of Growth

- Overall, the trend is **exponential growth** in the integration of new papers, reflecting the **expansion of the academic community** and **increased citation activity** over time.
- This mirrors the general trend of increasing scientific collaboration, access to research, and the growing interconnectedness of research papers.

2. Data Extrapolation via Clustering

2.1 Community Detection

Find the author-author network. Make sure each author retains the data about the papers it has published.

1. Loading and Preprocessing the Dataset

- **Read the dataset**: The CSV file is loaded using `pandas.read_csv()`, selecting only relevant columns: `"id"` (paper ID), `"year"` (publication year), `"authors"` (list of authors), and `"references"`.
- **Drop missing values**: Ensures that only complete records with authorship and year are considered.
- **Convert authors column**: Since the `"authors"` column contains string representations of lists, `ast.literal_eval()` is used to convert them into Python lists.

---

2. Building the Author-Author Network

- **Graph Representation**: The network is modeled as an **undirected graph** using `networkx.Graph()`, where:
  - **Nodes** represent authors.
  - **Edges** represent co-authorship (i.e., an edge exists between two authors if they have co-authored at least one paper).
- **Maintaining Author-Paper Information**:
  - A dictionary `author_papers` is used to store the papers each author has written.
  - Iterating through each row of the dataset:
    - The authors of a paper are extracted.
    - Each author is mapped to the papers they have written in `author_papers`.
- **Creating the Graph**:
  - If a paper has multiple authors, all **pairwise connections** (edges) between co-authors are added to the graph using `G.add_edge()`.
  - This results in a **collaboration network**, where frequent collaborators are more interconnected.

---

3. Displaying Network Statistics

- The number of **nodes (authors)** and **edges (co-authorships)** are printed to get an overview of the network's size.
- A sample of `author_papers` is displayed to verify that each author retains information about their publications.

Why is this Approach Effective?

1. **Captures Collaboration Patterns**: The graph structure allows the study of co-authorship patterns and community detection.
2. **Preserves Author Metadata**: The dictionary `author_papers` ensures that each author retains data on their publications.
3. **Efficient Graph Construction**: Uses adjacency relationships, making it computationally feasible even for large datasets.

```
import pandas as pd
import networkx as nx
import ast

# Load the dataset (adjust the path as needed)
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "year", "authors", "references"], low_memory=False)

# Drop missing values
df = df.dropna(subset=["authors", "year"])

# Convert the 'authors' column into a list
df["authors"] = df["authors"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Initialize the author-author network graph
G = nx.Graph()

# Create a dictionary to store the papers each author has published
author_papers = {}

# Iterate through the dataset and create edges between authors who co-authored a paper
for _, row in df.iterrows():
    authors = row["authors"]
    paper_id = row["id"]

    # Update the author_papers dictionary to store papers per author
    for author in authors:
        if author not in author_papers:
            author_papers[author] = []
```

```
        author_papers[author].append(paper_id)

    # Add edges between every pair of authors who co-authored the paper
    for i in range(len(authors)):
        for j in range(i + 1, len(authors)):
            G.add_edge(authors[i], authors[j])

# Now, G is the author-author network, and author_papers contains the papers each author has written
# Display the author-author network details (nodes and edges)
print(f"Number of authors (nodes): {len(G.nodes)}")
print(f"Number of co-authorships (edges): {len(G.edges)}")

# Display some of the authors and the papers they've written
for author in list(author_papers.keys())[:5]:  # Just display a few authors for illustration
    print(f"Author: {author}, Papers: {author_papers[author]}")
```

```
Number of authors (nodes): 940046
Number of co-authorships (edges): 3814744
Author: S. Ben Jabra, Papers: ['4ab3735c-80f1-472d-b953-fa0557fed28b']
Author: Ezzeddine Zagrouba, Papers: ['4ab3735c-80f1-472d-b953-fa0557fed28b', '5166b387-adbf-4a05-9739-ae8a8dcaadef', '60ecde4b-09f0-4c8a
Author: Joaquín J. Torres, Papers: ['4ab39729-af77-46f7-a662-16984fb9c1db', '622aa528-823b-48a6-8e2d-0270d04a9b0c', '9fcf5d0e-53cc-4131-
Author: Jesús M. Cortés, Papers: ['4ab39729-af77-46f7-a662-16984fb9c1db', '622aa528-823b-48a6-8e2d-0270d04a9b0c', '6a10e11a-7af9-4dab-80
Author: Joaquín Marro, Papers: ['4ab39729-af77-46f7-a662-16984fb9c1db', '622aa528-823b-48a6-8e2d-0270d04a9b0c', '9fcf5d0e-53cc-4131-a8f1
```

Find the author communities in this network using a fi tting clustering algorithm. You are allowed to use any algorithm, as long as the results make sense (Louvain, Walktrap, Hierarchical, Spectral, Newman, etc.). You must train 3 clustering algorithms.

1. Data Preprocessing

- **Load the dataset**: The dataset is loaded using `pandas.read_csv()`, selecting only the relevant columns: `"id"` (paper ID) and `"authors"` (list of authors).
- **Drop missing values**: Ensures that only valid author data is considered.
- **Convert `authors` column**: The `"authors"` column contains string representations of lists, so `ast.literal_eval()` is used to parse them into Python lists.
- **Sample the dataset**: To reduce computational complexity, a **small fraction (0.1%)** of the data is sampled (`frac=0.001`). This makes clustering more feasible.

---

2. Constructing the Author-Author Network

- **Graph Representation**: The co-authorship network is represented as an **undirected graph** (`networkx.Graph()`), where:
  - **Nodes**: Authors.
  - **Edges**: Co-authorship relationships (if two authors have written a paper together, an edge is added between them).
- **Tracking Papers for Each Author**:
  - Each author node stores a list of papers they have authored.
- **Avoiding Duplicate Edges**:
  - A `set` is used to ensure that edges are added **only once** (avoiding redundant `(A, B)` and `(B, A)` edges).

---

3. Clustering Algorithms

Three different clustering techniques are applied:

A. Louvain Algorithm (Modularity-Based Clustering)

- **Approach**:
  - The Louvain method is a hierarchical clustering technique that **optimizes modularity** to detect densely connected subgroups (communities).
  - `community_louvain.best_partition(author_network)` computes the best partition of nodes into communities.
- **Output**:
  - Communities are stored in a dictionary (`louvain_communities`).
  - The number of communities and sample communities (with the number of authors) are displayed.

B. Walktrap Algorithm (Label Propagation)

- **Approach**:
  - The Walktrap algorithm is a **random-walk-based community detection** method.
  - It assumes that **nodes in the same community are more likely to be visited together in random walks**.
  - `community.asyn_lpa_communities(author_network)` is used, which applies a **label propagation approach** (similar to Walktrap).
- **Output**:
  - The number of detected communities and sample communities are printed.

C. Hierarchical Clustering (Agglomerative Clustering)

- **Approach**:
  - Converts the network into an **adjacency matrix** (`nx.to_numpy_array(author_network)`).
  - Uses **Agglomerative Clustering** with:
    - `distance_threshold=0.5` : Controls the granularity of clusters.
    - `affinity='precomputed'` : Uses a precomputed **distance matrix** (1 - adjacency).
    - `linkage='average'` : Merges clusters based on the average distance.
  - `agg_clust.fit(1 - adj_matrix)` : Fits the model using **1 - adjacency** as the distance measure (since connected nodes should have smaller distances).
- **Output**:
  - Communities are stored in `hierarchical_communities` .
  - The number of hierarchical clusters and sample communities are displayed.

---

4. Results

- **Louvain**: Likely detects well-separated **dense clusters** of researchers who frequently collaborate.
- **Walktrap**: Provides communities based on **propagation of labels**, potentially revealing **hierarchical structures**.
- **Hierarchical Clustering**: Allows **multi-level clustering**, meaning researchers can belong to sub-clusters within larger groups.

Why This Approach Works Well?

1. **Diverse Clustering Methods**: Different approaches provide **multiple perspectives** on community structure.
2. **Scalability**: Louvain and Walktrap are well-suited for large networks, and hierarchical clustering can be tuned based on needs.
3. **Interpretability**: By comparing the three methods, we can determine which clustering best represents the **actual collaboration patterns**.

```
import pandas as pd
import ast
import networkx as nx
import community as community_louvain  # Louvain
from networkx.algorithms import community  # Walktrap and Hierarchical
from sklearn.cluster import AgglomerativeClustering  # Hierarchical
import numpy as np

# Load dataset (ensure 'id' and 'authors' columns are included)
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "authors"], low_memory=False)

# Drop missing values in the 'authors' column
df = df.dropna(subset=["authors"])

# Convert 'authors' from string lists to actual lists
df["authors"] = df["authors"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Sample a subset of the data
sample_fraction = 0.001
df_sampled = df.sample(frac=sample_fraction, random_state=42)  # Set random_state for reproducibility

# Initialize the author-author network
author_network = nx.Graph()

# Use a set to track edges and avoid duplicates
edges_set = set()

# Process each paper to build the network
for _, row in df_sampled.iterrows():
```

```
        paper_id = row["id"]
        authors = row["authors"]

        # Add authors as nodes and retain their paper information
        for author in authors:
            if author not in author_network:
                author_network.add_node(author, papers=[])  # Initialize with an empty list of papers
            author_network.nodes[author]["papers"].append(paper_id)  # Add paper to the author's list

        # Add edges between co-authors (only if they share a paper)
        for i in range(len(authors)):
            for j in range(i + 1, len(authors)):
                edge = tuple(sorted([authors[i], authors[j]]))  # Sort to avoid duplicate edges like (A, B) and (B, A)
                if edge not in edges_set:
                    edges_set.add(edge)
                    author_network.add_edge(authors[i], authors[j])  # Add edge between co-authors

    # Print basic information about the network
    print("Author-Author Network Information (Sampled Data):")
    print(f"Number of nodes (authors): {author_network.number_of_nodes()}")
    print(f"Number of edges (co-authorship relationships): {author_network.number_of_edges()}")

    # Louvain Algorithm

    print("\nClustering with Louvain Algorithm:")
    louvain_partition = community_louvain.best_partition(author_network)

    # Organize the Louvain communities
    louvain_communities = {}
    for author, community_id in louvain_partition.items():
        if community_id not in louvain_communities:
            louvain_communities[community_id] = []
        louvain_communities[community_id].append(author)

    # Print Louvain communities
    print(f"Number of communities (Louvain): {len(louvain_communities)}")
    for community_id, members in list(louvain_communities.items())[:10]:
        print(f"Community {community_id}: {len(members)} authors")

    # Walktrap Algorithm

    print("\nClustering with Walktrap Algorithm:")
    walktrap_communities = list(community.asyn_lpa_communities(author_network))  # Walktrap-like approach

    # Print Walktrap communities
    print(f"Number of communities (Walktrap): {len(walktrap_communities)}")
    for i, comm in enumerate(list(walktrap_communities)[:10]):
        print(f"Community {i}: {len(comm)} authors")

    # Hierarchical Clustering

    print("\nClustering with Hierarchical Clustering:")

    # Convert the graph into an adjacency matrix
    adj_matrix = nx.to_numpy_array(author_network)

    # Perform agglomerative hierarchical clustering
    agg_clust = AgglomerativeClustering(
        n_clusters=None,
        distance_threshold=0.5,
        affinity='precomputed',
        linkage='average'
    )

    # Fit the clustering model
    agg_clust.fit(1 - adj_matrix)  # The distance is 1 - adjacency (to convert to distance matrix)

    # Create a dictionary to store the clusters
    hierarchical_communities = {}
    for idx, label in enumerate(agg_clust.labels_):
        if label not in hierarchical_communities:
            hierarchical_communities[label] = []
        hierarchical_communities[label].append(list(author_network.nodes)[idx])

    # Print Hierarchical communities
    print(f"Number of communities (Hierarchical): {len(hierarchical_communities)}")
    for community_id, members in list(hierarchical_communities.items())[:10]:
```

```
print(f"Community {community_id}: {len(members)} authors")
```

```
Author-Author Network Information (Sampled Data):
Number of nodes (authors): 3144
Number of edges (co-authorship relationships): 5218

Clustering with Louvain Algorithm:
Number of communities (Louvain): 974
Community 0: 6 authors
Community 1: 2 authors
Community 2: 2 authors
Community 3: 3 authors
Community 4: 3 authors
Community 5: 3 authors
Community 6: 1 authors
Community 7: 5 authors
Community 8: 3 authors
Community 9: 4 authors

Clustering with Walktrap Algorithm:
Number of communities (Walktrap): 985
Community 0: 6 authors
Community 1: 2 authors
Community 2: 2 authors
Community 3: 3 authors
Community 4: 3 authors
Community 5: 3 authors
Community 6: 1 authors
Community 7: 5 authors
Community 8: 3 authors
Community 9: 4 authors

Clustering with Hierarchical Clustering:
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_agglomerative.py:983: FutureWarning: Attribute `affinity` was deprecated in ver
  warnings.warn(
Number of communities (Hierarchical): 998
Community 973: 6 authors
Community 804: 2 authors
Community 972: 2 authors
Community 971: 3 authors
Community 970: 3 authors
Community 969: 3 authors
Community 968: 1 authors
Community 967: 5 authors
Community 966: 3 authors
Community 965: 4 authors
```

Analysis of Clustering Results on the Author-Author Network

1. The Network

- The dataset is **sampled (0.1%)** to create an **author collaboration network**.
- Nodes represent **authors**, and edges represent **co-authorships** (papers with multiple authors).
- **Number of authors (nodes): 3,144**
- **Number of co-authorship relationships (edges): 5,218**
- This means that each author is, on average, connected to about **1.66 other authors**.

---

2. Clustering Analysis

Three different clustering algorithms were applied to detect **communities of authors**:

A. Louvain Algorithm

- **Number of communities: 974**
- **Typical community size:** Small (1-6 authors per community in the sample).
- **Interpretation:**
  - The **Louvain method** is a modularity-based clustering technique.
  - It optimizes the **modularity score**, which helps identify well-connected groups.
  - **Smaller communities** suggest that the collaboration graph is **loosely connected**, with authors working in **small research groups**.

B. Walktrap Algorithm

- **Number of communities: 985** (slightly more than Louvain).

- **Typical community size:** Very similar to Louvain, mostly **small groups** (1-6 authors per community).
- **Interpretation:**
  - Walktrap detects communities using **random walks**, which capture **local structures** in the graph.
  - The results are **very close to Louvain**, indicating that both methods detect **similar author groups**.
  - **Why more communities than Louvain?**
    - Walktrap tends to detect **more fine-grained communities**, splitting slightly more than Louvain.

C. Hierarchical Clustering (Agglomerative)

- **Number of communities: 998** (slightly more than both Louvain and Walktrap).
- **Typical community size:** Similar to other methods, mostly **small groups (1-6 authors)**.
- **Interpretation:**
  - This method **clusters authors based on similarity** (co-authorship distance).
  - The number of clusters is **slightly larger** than Louvain and Walktrap, meaning **some extra fragmentation** occurs.

Observations

- **All three methods produce similar results** → small research groups, loosely connected collaboration network.
- **Walktrap and Louvain are very close** → meaning the network has **well-separated communities**.
- **Hierarchical produces the most communities** → slightly more fragmented but still similar.
- The network shows **highly specialized groups** of **2-6 researchers** working together.
- There is **low connectivity across communities**, meaning **authors tend to work within fixed groups**.
- This might indicate **disciplinary silos** (limited interdisciplinary research).

Evaluate your clustering algorithm using the clustering coeffı cient and fı nd the best clustering algorithm from the 3 clustering algorithms.

What is the Clustering Coefficient?

The **clustering coefficient** is a measure of how **tightly connected** the neighbors of a node are in a network. It tells us the **likelihood** that two nodes connected to the same node are also connected to each other.

In simpler terms, it answers:
"If two people co-author with the same researcher, how likely are they to also be co-authors themselves?"

Types of Clustering Coefficients

There are **two levels** of clustering coefficients:

1. Local Clustering Coefficient (for a single node)

For an individual node **A**, the clustering coefficient measures the proportion of possible **triangles** (fully connected groups of 3) that actually exist.

$C_i = \frac{2E_i}{k_i(k_i-1)}$

Where:

- E_i = Number of **edges between** the neighbors of node i.
- k_i = Number of **neighbors** of node i (degree of node i).
- The denominator is the total number of **possible edges** between these neighbors.

> **Example:**
> If an author has **4 co-authors**, the total possible co-authorship connections between them is **4 choose 2 = 6**.
> If **3** of these 6 connections exist, the clustering coefficient for that author is:
> $C_i = \frac{3}{6} = 0.5$ Meaning **50% of the author's co-authors also collaborate with each other**.

---

2. Global Clustering Coefficient (for the whole network)

The **network-wide clustering coefficient** is the **average of all local clustering coefficients** across nodes.

$C = \frac{1}{N} \sum_{i=1}^{N} C_i$

Where:

- N is the number of nodes.

- C_i is the local clustering coefficient of node i.

This gives an overall measure of how "**cliquish**" the network is.

Interpreting Clustering Coefficients

- $C \approx 1$ **(Close to 1):**
  - The network is **highly clustered**.
  - Nodes tend to form **tight-knit communities** (e.g., research groups, friend circles).
- $C \approx 0$ **(Close to 0):**
  - The network is **loosely connected**.
  - Nodes **rarely form closed triangles** (e.g., a random collaboration network).

Why is the Clustering Coefficient Important?

- In an **author network**, a **high clustering coefficient** means:
  - **Strong research communities** exist.
  - Researchers **frequently co-author with each other**.
  - The network has **high collaboration density**.
- In **social networks**, a high clustering coefficient means:
  - Friends of a person are **likely to be friends with each other**.
  - There are **tight social groups**.

Approach Explanation: Evaluating Clustering Algorithms Using Clustering Coefficient

Our approach aims to evaluate the **three clustering algorithms** (**Louvain, Walktrap, and Hierarchical Clustering**) based on the **average clustering coefficient** of detected communities. The clustering coefficient measures **how well nodes in a community form tightly-knit groups** (i.e., how likely neighbors of a node are also connected).

---

1. Data Preprocessing & Author-Author Network Construction

This part follows the same process as in your previous steps:

- **Load the dataset** and filter relevant columns (`id, authors`).
- **Drop missing values** to avoid issues with clustering.
- **Convert `authors` column into lists** (from string representations).
- **Sample the dataset (`0.1%`)** to make computations more feasible.
- **Construct the author-author graph (`networkx.Graph()`)** where:
  - **Nodes represent authors**.
  - **Edges represent co-authorship relationships**.
  - **A dictionary tracks each author's papers**.

---

2. Function to Compute Clustering Coefficient

A function, `calculate_clustering_coefficient()`, is defined to:

- **Extract a subgraph** containing only the members of a community.
- **Compute the average clustering coefficient** for the subgraph using `networkx.average_clustering()`.
  - This tells us how interconnected the members of a community are.

---

3. Community Detection & Clustering Coefficient Calculation

Each of the three clustering algorithms is applied, and their performance is measured using the **average clustering coefficient**.

A. Louvain Clustering (Modularity Optimization)

- **Detects communities** using `community_louvain.best_partition(author_network)`.
- **Groups authors into communities** (dictionary: `louvain_communities`).
- **Computes clustering coefficient** for each community and averages it.
- **Prints the final average clustering coefficient**.

B. Walktrap Clustering (Label Propagation)

- **Applies Walktrap-like label propagation** using `community.asyn_lpa_communities(author_network)`.
- **Computes clustering coefficient** for each detected community.
- **Prints the final average clustering coefficient**.

C. Hierarchical Clustering (Agglomerative)

- **Converts the network into an adjacency matrix** (`nx.to_numpy_array(author_network)`).
- **Applies Agglomerative Clustering** (`AgglomerativeClustering()`), setting:
  - `distance_threshold=0.5` (controls merging of clusters).
  - `affinity='precomputed'` (since we use a custom distance matrix).
  - `linkage='average'` (merges clusters based on mean distances).
- **Clusters are stored** in `hierarchical_communities`.
- **Computes clustering coefficient** for each hierarchical cluster.
- **Prints the final average clustering coefficient**.

---

4. Identifying the Best Clustering Algorithm

- **Compares the average clustering coefficient** for each algorithm.
- **Selects the algorithm with the highest clustering coefficient** as the best one.
- **Prints the best algorithm** based on the results.

```
import pandas as pd
import ast
import networkx as nx
import community as community_louvain  # Louvain algorithm
from networkx.algorithms import community  # For Walktrap and Hierarchical
from sklearn.cluster import AgglomerativeClustering  # For Hierarchical clustering
import numpy as np

# Load dataset (ensure 'id' and 'authors' columns are included)
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "authors"], low_memory=False)

# Drop missing values in the 'authors' column
df = df.dropna(subset=["authors"])

# Convert 'authors' from string lists to actual lists
df["authors"] = df["authors"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Sample a subset of the data (e.g., 0.1%)
sample_fraction = 0.001
df_sampled = df.sample(frac=sample_fraction, random_state=42)  # Set random_state for reproducibility

# Initialize the author-author network
author_network = nx.Graph()

# Use a set to track edges and avoid duplicates
edges_set = set()

# Process each paper to build the network
for _, row in df_sampled.iterrows():
    paper_id = row["id"]
    authors = row["authors"]

    # Add authors as nodes and retain their paper information
    for author in authors:
        if author not in author_network:
            author_network.add_node(author, papers=[])  # Initialize with an empty list of papers
        author_network.nodes[author]["papers"].append(paper_id)  # Add paper to the author's list

    # Add edges between co-authors (only if they share a paper)
    for i in range(len(authors)):
        for j in range(i + 1, len(authors)):
            edge = tuple(sorted([authors[i], authors[j]]))  # Sort to avoid duplicate edges like (A, B) and (B, A)
            if edge not in edges_set:
                edges_set.add(edge)
                author_network.add_edge(authors[i], authors[j])  # Add edge between co-authors

# Function to calculate the clustering coefficient for each community
```

```
def calculate_clustering_coefficient(community_members, network):
    subgraph = network.subgraph(community_members)
    return nx.average_clustering(subgraph)

# Louvain Algorithm

print("\nClustering with Louvain Algorithm:")
louvain_partition = community_louvain.best_partition(author_network)

# Organize the Louvain communities
louvain_communities = {}
for author, community_id in louvain_partition.items():
    if community_id not in louvain_communities:
        louvain_communities[community_id] = []
    louvain_communities[community_id].append(author)

# Calculate the average clustering coefficient for Louvain
louvain_clustering_coeff = np.mean([calculate_clustering_coefficient(members, author_network) for members in louvain_communities.values()])
print(f"Average Clustering Coefficient (Louvain): {louvain_clustering_coeff:.4f}")

# Walktrap Algorithm

print("\nClustering with Walktrap Algorithm:")
walktrap_communities = list(community.asyn_lpa_communities(author_network))  # Walktrap-like approach

# Calculate the average clustering coefficient for Walktrap
walktrap_clustering_coeff = np.mean([calculate_clustering_coefficient(comm, author_network) for comm in walktrap_communities])
print(f"Average Clustering Coefficient (Walktrap): {walktrap_clustering_coeff:.4f}")

# Hierarchical Clustering

print("\nClustering with Hierarchical Clustering:")

# Convert the graph into an adjacency matrix
adj_matrix = nx.to_numpy_array(author_network)

# Perform agglomerative hierarchical clustering
agg_clust = AgglomerativeClustering(
    n_clusters=None,
    distance_threshold=0.5,
    affinity='precomputed',
    linkage='average'
)
# We set a distance threshold to control the number of clusters

# Fit the clustering model
agg_clust.fit(1 - adj_matrix)  # The distance is 1 - adjacency (to convert to distance matrix)

# Create a dictionary to store the clusters
hierarchical_communities = {}
for idx, label in enumerate(agg_clust.labels_):
    if label not in hierarchical_communities:
        hierarchical_communities[label] = []
    hierarchical_communities[label].append(list(author_network.nodes)[idx])

# Calculate the average clustering coefficient for Hierarchical
hierarchical_clustering_coeff = np.mean([calculate_clustering_coefficient(members, author_network) for members in hierarchical_communities.v
print(f"Average Clustering Coefficient (Hierarchical): {hierarchical_clustering_coeff:.4f}")


# Compare and Find the Best Clustering Algorithm

# Identify the best clustering algorithm based on the highest average clustering coefficient
best_algorithm = max(
    [("Louvain", louvain_clustering_coeff),
     ("Walktrap", walktrap_clustering_coeff),
     ("Hierarchical", hierarchical_clustering_coeff)],
    key=lambda x: x[1]
)

print(f"\nBest Clustering Algorithm based on Clustering Coefficient: {best_algorithm[0]}")
```

```
Clustering with Louvain Algorithm:
Average Clustering Coefficient (Louvain): 0.6022
```

```
Clustering with Walktrap Algorithm:
Average Clustering Coefficient (Walktrap): 0.6029

Clustering with Hierarchical Clustering:
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_agglomerative.py:983: FutureWarning: Attribute `affinity` was deprecated in ver
  warnings.warn(
Average Clustering Coefficient (Hierarchical): 0.6012

Best Clustering Algorithm based on Clustering Coefficient: Walktrap
```

Analysis of the Clustering Results

Based on the **average clustering coefficient**, the three clustering algorithms produced **very similar results**, with **Walktrap performing slightly better**:

---

1. Observations

A. Close Performance Across All Methods

- The **clustering coefficients are nearly identical** (differences in the **fourth decimal place**), indicating that:
  - The detected communities **share a similar internal structure** regardless of the method used.
  - The **network is highly clustered** in nature, meaning authors **form strongly connected co-authorship groups**.

B. Walktrap Slightly Outperforms Others

- **Walktrap achieved the highest clustering coefficient (0.6029)**.
- **Why?**
  - Walktrap is **based on label propagation**, where nodes **pass labels to their neighbors iteratively**, grouping nodes that interact frequently.
  - This method **naturally promotes densely interconnected communities**, explaining the **slight edge over the others**.

C. Hierarchical Clustering Performs Slightly Worse

- **Hierarchical Clustering scored the lowest (0.6012)**.
- **Possible reasons:**
  - Agglomerative clustering **isn't optimized for highly modular networks**.
  - The **distance threshold (`0.5`) might have led to excessive fragmentation**, reducing clustering within groups.
  - The warning (`affinity` being deprecated) **did not affect results** but should be updated to `metric` in future versions.

---

2. Implications

- Since **Walktrap achieved the highest clustering coefficient**, it suggests that **Walktrap forms the most tightly-knit research communities**.
- The difference in performance is **small**, meaning **all three algorithms are viable choices** for author community detection.
- If the goal is **best community coherence**, **Walktrap is the preferred choice**.

Then run the best clustering algorithm on the whole data.

```
import pandas as pd
import ast
import networkx as nx
from networkx.algorithms import community  # For Walktrap
import numpy as np

# Load the full dataset (ensure 'id' and 'authors' columns are included)
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, usecols=["id", "authors"], low_memory=False)

# Drop missing values in the 'authors' column
df = df.dropna(subset=["authors"])

# Convert 'authors' from string lists to actual lists
df["authors"] = df["authors"].apply(lambda x: ast.literal_eval(x) if isinstance(x, str) else [])

# Initialize the author-author network
author_network = nx.Graph()
```

```
# Use a set to track edges and avoid duplicates
edges_set = set()

# Process each paper to build the network
for _, row in df.iterrows():
    paper_id = row["id"]
    authors = row["authors"]

    # Add authors as nodes and retain their paper information
    for author in authors:
        if author not in author_network:
            author_network.add_node(author, papers=[])  # Initialize with an empty list of papers
        author_network.nodes[author]["papers"].append(paper_id)  # Add paper to the author's list

    # Add edges between co-authors (only if they share a paper)
    for i in range(len(authors)):
        for j in range(i + 1, len(authors)):
            edge = tuple(sorted([authors[i], authors[j]]))  # Sort to avoid duplicate edges like (A, B) and (B, A)
            if edge not in edges_set:
                edges_set.add(edge)
                author_network.add_edge(authors[i], authors[j])  # Add edge between co-authors

# Clustering with Walktrap Algorithm on the full dataset
print("\nClustering with Walktrap Algorithm on Full Data:")

# Detect communities using the Walktrap-like approach
walktrap_communities = list(community.asyn_lpa_communities(author_network))  # Walktrap-like approach

# Print information about the communities
print(f"Number of communities (Walktrap): {len(walktrap_communities)}")
for i, comm in enumerate(walktrap_communities[:5]):
    print(f"Community {i}: {len(comm)} authors")
```

```
Clustering with Walktrap Algorithm on Full Data:
Number of communities (Walktrap): 160734
Community 0: 13 authors
Community 1: 29 authors
Community 2: 5 authors
Community 3: 3 authors
Community 4: 27 authors
```

```
# Calculate the average clustering coefficient for Walktrap communities
walktrap_clustering_coeffs = {}

# Loop through each community and calculate the clustering coefficient for the subgraph of the community
for i, comm in enumerate(walktrap_communities):
    subgraph = author_network.subgraph(comm)  # Create the subgraph for the community
    walktrap_clustering_coeffs[i] = nx.average_clustering(subgraph)  # Calculate clustering coefficient for the subgraph

# Calculate the overall average clustering coefficient for Walktrap
walktrap_avg_clustering = np.mean(list(walktrap_clustering_coeffs.values()))

print(f"\nAverage Clustering Coefficient for Walktrap: {walktrap_avg_clustering}")
```

```
Average Clustering Coefficient for Walktrap: 0.4734890670010038
```

Analysis of Walktrap Clustering Results on Full Dataset

1. Number of Communities Found

- The Walktrap algorithm detected **160,734** communities in the full author network.
- This suggests that the dataset contains a **highly fragmented co-authorship network**, with many small, isolated research groups.

2. Size of Sample Communities

- Looking at the first five detected communities:
  - **Community 0:** 13 authors
  - **Community 1:** 29 authors
  - **Community 2:** 5 authors
  - **Community 3:** 3 authors

- **Community 4:** 27 authors

- Some communities are **moderately sized (10–30 authors)**, while others are **very small (2–5 authors)**.
- This aligns with the structure of real-world research networks, where:
  - Some research teams are **large**, involving multiple institutions.
  - Others are **small**, involving just a few researchers.

3. Average Clustering Coefficient

- The **average clustering coefficient** for Walktrap is **0.4735**.
- This means that **47.35% of an author's co-authors are also co-authors with each other**.
- A clustering coefficient close to **0.5** suggests:
  - The network is **moderately clustered**.
  - There is a **reasonable level of collaboration within research communities**.
  - However, many authors are **not tightly connected**, leading to a slightly lower coefficient.

4. Comparison to Previous Results

- In the previous experiment (on a **sampled dataset**), Walktrap had a clustering coefficient of **0.6029**.
- The **lower clustering coefficient (0.4735) on the full dataset** suggests that:
  - The sampled dataset focused on **more connected authors**, leading to **higher clustering**.
  - The full dataset contains **many isolated researchers**, lowering the overall clustering.

Interpretation of the Results

- Walktrap successfully identifies **many communities**, but many are **small or weakly connected**.
- The clustering coefficient of **0.4735** suggests a **moderate level of local collaboration** among researchers.
- The large number of detected communities (**160,734**) may indicate:
  - The presence of **many isolated research teams**.
  - Possible **over-segmentation** of communities.

2.2 Naming the Communities

Each paper has an abstract and a title. We want to extract keywords from these text fields. Use KeyBERT for extracting keywords from each abstract/title and aggregate them together to form a single list of words as keywords for each paper.

```
!pip install keybert
```

```
Collecting keybert
  Downloading keybert-0.8.5-py3-none-any.whl.metadata (15 kB)
Requirement already satisfied: numpy>=1.18.5 in /usr/local/lib/python3.10/dist-packages (from keybert) (1.26.4)
Requirement already satisfied: rich>=10.4.0 in /usr/local/lib/python3.10/dist-packages (from keybert) (13.9.4)
Requirement already satisfied: scikit-learn>=0.22.2 in /usr/local/lib/python3.10/dist-packages (from keybert) (1.2.2)
Requirement already satisfied: sentence-transformers>=0.3.8 in /usr/local/lib/python3.10/dist-packages (from keybert) (3.3.1)
Requirement already satisfied: mkl_fft in /usr/local/lib/python3.10/dist-packages (from numpy>=1.18.5->keybert) (1.3.8)
Requirement already satisfied: mkl_random in /usr/local/lib/python3.10/dist-packages (from numpy>=1.18.5->keybert) (1.2.4)
Requirement already satisfied: mkl_umath in /usr/local/lib/python3.10/dist-packages (from numpy>=1.18.5->keybert) (0.1.1)
Requirement already satisfied: mkl in /usr/local/lib/python3.10/dist-packages (from numpy>=1.18.5->keybert) (2025.0.1)
Requirement already satisfied: tbb4py in /usr/local/lib/python3.10/dist-packages (from numpy>=1.18.5->keybert) (2022.0.0)
Requirement already satisfied: mkl-service in /usr/local/lib/python3.10/dist-packages (from numpy>=1.18.5->keybert) (2.4.1)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.10/dist-packages (from rich>=10.4.0->keybert) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dist-packages (from rich>=10.4.0->keybert) (2.18.0)
Requirement already satisfied: typing-extensions<5.0,>=4.0.0 in /usr/local/lib/python3.10/dist-packages (from rich>=10.4.0->keybert) (4.
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.22.2->keybert) (1.13.1)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.22.2->keybert) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.22.2->keybert) (3.5
Requirement already satisfied: transformers<5.0.0,>=4.41.0 in /usr/local/lib/python3.10/dist-packages (from sentence-transformers>=0.3.8
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from sentence-transformers>=0.3.8->keybert) (4.67.1)
Requirement already satisfied: torch>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from sentence-transformers>=0.3.8->keybert) (2.
Requirement already satisfied: huggingface-hub>=0.20.0 in /usr/local/lib/python3.10/dist-packages (from sentence-transformers>=0.3.8->ke
Requirement already satisfied: Pillow in /usr/local/lib/python3.10/dist-packages (from sentence-transformers>=0.3.8->keybert) (11.0.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.20.0->sentence-transformers>
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.20.0->sentence-trans
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.20.0->sentence-transf
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.20.0->sentence-transforme
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.20.0->sentence-transformers>
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py>=2.2.0->rich>=10.4.0->keybert)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.11.0->sentence-transformers>=0.3.8->ke
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.11.0->sentence-transformers>=0.3.8->keyb
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (from torch>=1.11.0->sentence-transformers>=0.3.
```

```
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy==1.13.1->torch>=1.11.0->sentenc
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers<5.0.0,>=4.41.0->sentence-
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/python3.10/dist-packages (from transformers<5.0.0,>=4.41.0->sent
Requirement already satisfied: safetensors>=0.4.1 in /usr/local/lib/python3.10/dist-packages (from transformers<5.0.0,>=4.41.0->sentence
Requirement already satisfied: intel-openmp>=2024 in /usr/local/lib/python3.10/dist-packages (from mkl->numpy>=1.18.5->keybert) (2024.2.
Requirement already satisfied: tbb==2022.* in /usr/local/lib/python3.10/dist-packages (from mkl->numpy>=1.18.5->keybert) (2022.0.0)
Requirement already satisfied: tcmlib==1.* in /usr/local/lib/python3.10/dist-packages (from tbb==2022.*->mkl->numpy>=1.18.5->keybert) (1
Requirement already satisfied: intel-cmplr-lib-rt in /usr/local/lib/python3.10/dist-packages (from mkl_umath->numpy>=1.18.5->keybert) (2
Requirement already satisfied: intel-cmplr-lib-ur==2024.2.0 in /usr/local/lib/python3.10/dist-packages (from intel-openmp>=2024->mkl->nu
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.11.0->sentence-transfor
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->huggingface-hub>=0.20
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->huggingface-hub>=0.20.0->sentence
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->huggingface-hub>=0.20.0->se
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->huggingface-hub>=0.20.0->se
Downloading keybert-0.8.5-py3-none-any.whl (37 kB)
Installing collected packages: keybert
Successfully installed keybert-0.8.5
```

Approach for Extracting Keywords from Research Papers

1. Understanding the Data

Each research paper contains:
A **Title** – A concise summary of the paper's topic.
An **Abstract** – A detailed description of the paper's contributions and focus.

Since both fields provide important information, extracting keywords from both ensures a **comprehensive understanding** of the paper's topic.

---

2. Using KeyBERT for Keyword Extraction

To identify meaningful keywords, we use **KeyBERT**, a model based on **BERT embeddings**, which captures the **context and semantics** of words rather than just their frequency.

- The model reads the **title and abstract** of a paper.
- It identifies the **most relevant words or phrases** that summarize the content.
- It selects the **top-ranked** keywords based on their importance.

Unlike traditional keyword extraction methods (such as TF-IDF), KeyBERT ensures that the extracted words are **contextually meaningful** rather than just frequently occurring.

---

3. Combining Keywords from Title and Abstract

Since the **title** provides a high-level view and the **abstract** offers more details, we extract keywords separately from both.

- Keywords from both fields are **merged** to create a single, refined list.
- **Duplicate words** are removed to avoid redundancy.
- The final list provides a **balanced summary** of the research paper's topic.

---

4. Processing the Dataset Efficiently

Since research paper datasets can be **large**, we:

- **Remove missing data** to ensure all records contain valid text.
- **Randomly sample** a portion of the dataset to improve efficiency while maintaining diversity.
- **Apply the keyword extraction model** to each sampled paper.

This allows us to **quickly process a large number of papers** while still capturing diverse research topics.

---

5. Storing and Using Extracted Keywords

Once extracted, the keywords are **stored as a new feature** in the dataset. They can be used for:

- **Clustering papers into research areas** based on common topics.
- **Naming author communities** by identifying dominant research themes.
- **Building a recommendation system** for finding related papers.
- **Visualizing research trends** using word clouds or topic networks.

```
from sklearn.model_selection import train_test_split
from keybert import KeyBERT
import pandas as pd
```

```
# Initialize KeyBERT model
kw_model = KeyBERT()

# Function to extract keywords from both title and abstract
def extract_keywords(row):
    # Extract keywords from the title and abstract
    title_keywords = kw_model.extract_keywords(row['title'], top_n=5)
    abstract_keywords = kw_model.extract_keywords(row['abstract'], top_n=5)

    # Combine both title and abstract keywords
    combined_keywords = set([kw[0] for kw in title_keywords] + [kw[0] for kw in abstract_keywords])

    return list(combined_keywords)

# Load dataset
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, low_memory=False)

# Drop missing values in the 'title', 'abstract', and 'year' columns
df = df.dropna(subset=["title", "abstract", "year"])

# Filter out years with only one paper
year_counts = df['year'].value_counts()
valid_years = year_counts[year_counts > 1].index
df_filtered = df[df['year'].isin(valid_years)]

# Stratified sampling over the 'year' column
df_sampled, _ = train_test_split(df_filtered, test_size=0.99, stratify=df_filtered['year'], random_state=42)

# Apply the function to extract keywords and create a new 'keywords' column
df_sampled['keywords'] = df_sampled.apply(extract_keywords, axis=1)

# Print the first few papers with their extracted keywords
print(df_sampled[['title', 'keywords']].head())
```

```
                                                   title  \
139164              Information-singular random processes
439813  A linear algorithm for motion from three weak ...
461284  Quantitative Fourier analysis of approximation...
124701  Voice and data transmissions with a PRMA-like ...
592916  Design optimization of circular antenna arrays...

                                                keywords
139164  [brownian, entropy, deterministic, singular, i...
439813  [projections, rotations, images, rotation, eul...
461284  [projectors, fourier, interpolation, technique...
124701  [protocols, multiplexing, satellite, prma, del...
592916  [taguchi, minimize, arrays, array, circular, o...
```

Associate each paper with the corresponding community of authors. A paper might be in two or more communities, so make sure to handle this correctly.

Approach:

1. Understanding the Problem

- **Research papers** have multiple authors.
- **Authors** collaborate with different researchers, forming **co-authorship networks**.
- We use **graph-based community detection** to find **groups of closely connected authors**.
- Since an author can belong to **multiple communities**, a **paper can be part of multiple communities**.

---

2. Data Preprocessing

Ensure Author Data is Correctly Parsed

- Some datasets store author names as **strings representing lists** (e.g., `"[Author1, Author2]"`).
- We **convert these strings into actual Python lists** for easy processing.
- Papers with **missing or empty author lists** are removed.

  Build an Author-Author Network

- **Authors** are treated as **nodes** in a graph.
- **Edges** represent **co-authorship relationships** (i.e., authors who wrote a paper together).
- Each **author node** stores the **list of papers** they contributed to.

  Avoid Duplicate Edges

- To prevent redundant links (e.g., `(A, B)` and `(B, A)`), we **store edges in a set** and only add unique connections.

---

3. Detecting Author Communities

We apply **Walktrap-like community detection** (using the **asynchronous label propagation algorithm** from NetworkX).

- This algorithm **clusters authors into communities** based on their connections.
- Each **author gets assigned to one or more communities** depending on their network structure.
- A dictionary **maps each author to their respective communities**.

---

4. Assigning Papers to Communities

Each paper is **associated with all the communities** its authors belong to:

1. **For each paper**, retrieve its **list of authors**.
2. **Check which communities each author belongs to**.
3. **Merge all communities** from the authors to form the **final list of communities for the paper**.
4. **Store this information** in a new column (`"communities"`) in the dataset.

This ensures that:
A paper is **linked to all relevant research groups**.
Multi-author papers **belong to multiple communities if authors are from different groups**.

---

5. Outputs

- The total **number of detected author communities** is displayed.
- The total **number of papers processed** is shown.
- A **sample of papers** is printed, showing:
  - **Paper ID**
  - **Title**
  - **List of Authors**
  - **Communities the paper belongs to**

```
import pandas as pd
import ast
import networkx as nx
from networkx.algorithms import community  # For Walktrap
import numpy as np
from keybert import KeyBERT

# Drop missing values in the 'authors' column
df_sampled = df_sampled.dropna(subset=["authors"])

# Check the 'authors' column before parsing to see the format
print("Authors column (before parsing):")
print(df_sampled["authors"].head())

# Convert 'authors' from string lists to actual lists
# Adding a check to ensure proper parsing
def safe_parse_authors(authors):
    try:
        return ast.literal_eval(authors) if isinstance(authors, str) else authors
    except Exception as e:
        print(f"Error parsing authors: {e}")
        return []

df_sampled["authors"] = df_sampled["authors"].apply(safe_parse_authors)

# Check the 'authors' column after parsing
print("\nAuthors column (after parsing):")
print(df_sampled["authors"].head())

# Filter out rows where authors are empty lists
```

```
df_sampled = df_sampled[df_sampled["authors"].apply(lambda x: len(x) > 0)]

# Initialize the author-author network
author_network = nx.Graph()

# Use a set to track edges and avoid duplicates
edges_set = set()

# Process each paper to build the network
for _, row in df_sampled.iterrows():
    paper_id = row["id"]
    authors = row["authors"]

    # Add authors as nodes and retain their paper information
    for author in authors:
        if author not in author_network:
            author_network.add_node(author, papers=[])  # Initialize with an empty list of papers
        author_network.nodes[author]["papers"].append(paper_id)  # Add paper to the author's list

    # Add edges between co-authors if they share a paper
    for i in range(len(authors)):
        for j in range(i + 1, len(authors)):
            edge = tuple(sorted([authors[i], authors[j]]))  # Sort to avoid duplicate edges like (A, B) and (B, A)
            if edge not in edges_set:
                edges_set.add(edge)
                author_network.add_edge(authors[i], authors[j])  # Add edge between co-authors

# Clustering with Walktrap Algorithm
print("\nClustering with Walktrap Algorithm:")

# Detect communities using the Walktrap-like approach
walktrap_communities = list(community.asyn_lpa_communities(author_network))  # Walktrap-like approach

# Map each author to their corresponding community(s)
author_to_community = {}
for community_id, authors_in_comm in enumerate(walktrap_communities):
    for author in authors_in_comm:
        if author not in author_to_community:
            author_to_community[author] = []
        author_to_community[author].append(community_id)

# Associate each paper with its community
def associate_paper_with_communities(row):
    communities = set()
    authors = row["authors"]

    # Loop through each author of the paper and add their communities
    for author in authors:
        if author in author_to_community:
            communities.update(author_to_community[author])  # Add the communities this author belongs to

    return list(communities)

# Apply the function to associate each paper with its community
df_sampled["communities"] = df_sampled.apply(associate_paper_with_communities, axis=1)

# Print basic summary of the results
print(f"\nTotal number of papers processed: {df_sampled.shape[0]}")
print(f"Total number of author communities detected: {len(walktrap_communities)}")

# Print a few sample papers with their associated communities
print("\nSample papers with associated communities:")
for _, row in df_sampled.head(10).iterrows():
    print(f"Paper ID: {row['id']}")
    print(f"Title: {row['title']}")
    print(f"Authors: {row['authors']}")
    print(f"Communities: {row['communities']}\n")
```

```
Authors column (before parsing):
139164                                     ['Toby Berger']
439813                        ['Gang Xu', 'Noriko Sugimoto']
461284                      ['Thierry Blu', 'Michael Unser']
124701    ['Giuliano Benelli', 'Romano Fantacci', 'Giova...
592916                ['Bilal Babayigit', 'Ercan Senyigit']
Name: authors, dtype: object

Authors column (after parsing):
```

```
139164                                        [Toby Berger]
439813                           [Gang Xu, Noriko Sugimoto]
461284                         [Thierry Blu, Michael Unser]
124701    [Giuliano Benelli, Romano Fantacci, Giovanni G...
592916                    [Bilal Babayigit, Ercan Senyigit]
Name: authors, dtype: object

Clustering with Walktrap Algorithm:

Total number of papers processed: 8275
Total number of author communities detected: 7266

Sample papers with associated communities:
Paper ID: 78ebfe9c-d031-46bd-9069-b12c130713fc
Title: Information-singular random processes
Authors: ['Toby Berger']
Communities: [0]

Paper ID: dcc5d799-baf5-4dd7-8a06-fe1332c42637
Title: A linear algorithm for motion from three weak perspective images using Euler angles
Authors: ['Gang Xu', 'Noriko Sugimoto']
Communities: [1]

Paper ID: e3ead250-4e03-41be-b6b0-8347f607b80c
Title: Quantitative Fourier analysis of approximation techniques. I. Interpolators and projectors
Authors: ['Thierry Blu', 'Michael Unser']
Communities: [2]

Paper ID: 7417fa37-babb-4c77-a5b1-10188305945c
Title: Voice and data transmissions with a PRMA-like protocol in high propagation delay cellular systems
Authors: ['Giuliano Benelli', 'Romano Fantacci', 'Giovanni Giambene', 'Carlo Ortolani']
Communities: [3]

Paper ID: 582d244f-a827-429e-a8d7-bc8f9e798a7d
Title: Design optimization of circular antenna arrays using Taguchi method
Authors: ['Bilal Babayigit', 'Ercan Senyigit']
Communities: [4]

Paper ID: febbb04d-0c5e-4493-8a97-85a4f13a8f81
Title: Measuring the Impact of Data Mining on Churn Management.
Authors: ['Miguel A. P. M. Lejeune']
Communities: [5]

Paper ID: 2b1cf591-c86f-455a-b11f-771a0cec7448
Title: Clustering with Pairs of Prototypes to Support Automated Assessment of the Fetal State
Authors: ['Michal Jezewski', 'Robert Czabanski', 'K. Horoba', 'Jacek M. Leski']
Communities: [6]

Paper ID: fd6055bc-8b6c-4bfb-b52a-9a4c22b91f87
```

Aggregate the keywords of each paper from each community to get the keywords of the community. Explain your aggregation method, and report the keywords of each community in your report.

The **aggregation method** used here is designed to summarize the **keywords of each paper** within a specific **community of authors**. This method helps to capture the key themes or concepts that are most relevant to a group of authors collaborating on similar research topics. Below, I'll walk you through the step-by-step **aggregation process** and its purpose:

The Aggregation Method:

1. **Identify the Keywords for Each Paper:**
   - For each paper, keywords are extracted from both the **title** and **abstract**. This is done using **KeyBERT**, which extracts the most relevant terms that represent the core topics of the paper.
   - Each paper, therefore, has a **list of keywords** that are associated with it.
2. **Community Assignment:**
   - Each paper is associated with one or more **author communities** based on the authors it involves and their collaboration networks (e.g., using the **Walktrap algorithm** to detect communities).
   - A paper may belong to **multiple communities** if its authors belong to multiple groups, so this aggregation method handles that correctly.
3. **Aggregate Keywords by Community:**
   - For each **community**, we need to **combine the keywords of all papers** that belong to it. This step ensures that the overall set of keywords reflects the collaborative research topics and interests of that community as a whole.

4. **Remove Duplicates:**
   - Since papers in the same community may share common keywords, the **aggregation method removes duplicates**. This is done by storing the keywords in a **set**, which automatically eliminates repeated entries.
5. **Count Keyword Frequency:**
   - After collecting the keywords from all papers in the community, we need to understand which keywords are most prominent in the community.
   - This is done by **counting the frequency of each keyword**. Keywords that appear frequently across papers in the same community are considered more important or central to that community's focus.
6. **Sort and Select Top Keywords:**
   - Once we have the frequencies of the keywords, we sort them in descending order of frequency.
   - The **top N keywords** (in this case, we selected the **top 10**) are then chosen to represent the most prominent themes of the community.
7. **Output the Keywords for Each Community:**
   - Finally, we print or display the **top keywords** for each community, providing a summary of the key research areas and topics that are prevalent within that group of authors.

```
from collections import Counter

# Function to aggregate keywords for each community
def aggregate_keywords(df_sampled, author_to_community):
    community_keywords = {}

    # Loop through each row in the dataframe to assign keywords to the corresponding community
    for _, row in df_sampled.iterrows():
        paper_id = row["id"]
        keywords = row["keywords"]
        communities = row["communities"]

        # For each community the paper belongs to, add its keywords
        for community_id in communities:
            if community_id not in community_keywords:
                community_keywords[community_id] = []
            community_keywords[community_id].extend(keywords)

    # Convert the aggregated keywords into sets (to remove duplicates) and sort by frequency
    community_summary = {}
    for community_id, keywords in community_keywords.items():
        # Count the frequency of each keyword in the community
        keyword_counts = Counter(keywords)
        # Get the top 10 most frequent keywords in the community
        sorted_keywords = [kw for kw, _ in keyword_counts.most_common(10)]
        community_summary[community_id] = sorted_keywords

    return community_summary

# Aggregate the keywords for each community
community_keywords = aggregate_keywords(df_sampled, author_to_community)

# Print out the keywords for each community
print("\nKeywords for each community:")
for community_id, keywords in community_keywords.items():
    print(f"Community {community_id}: {', '.join(keywords)}")
```

```
Keywords for each community:
Community 0: brownian, entropy, deterministic, singular, information, random, nondeterministic, processes
Community 1: projections, rotations, images, rotation, euler, angles, motion, perspective
Community 2: wavelet, projectors, fourier, interpolation, techniques, interpolators, approximation, segmentation, chromosome, friendl
Community 3: protocols, multiplexing, satellite, prma, delay, packet, cellular, protocol, transmissions
Community 4: taguchi, minimize, arrays, array, circular, optimization, suppression, antenna
Community 5: customer, data, management, impact, customers, mining, digital, churn
Community 6: abstractcardiotocographic, fetal, automated, clustering, classification, fuzzy, assessment, prototypes, sensitivity, ave
Community 7: transform, streaming, distributed, bandwidth, 3d, decomposition, filter, images, revertible, convolutional
Community 8: development, process, line, evolving, models, software, components, structuring, products
Community 9: analyser, services, chaîne, industrielle, inclusion, utilisant, professionnellement, lestic, analyse
Community 10: scalable, processing, streaming, distributed, mobile, nim, data, datasets, stream, mining
Community 11: detection, user, fpga, presence, design, camera, monitoring, detecting
Community 12: constraint, tree, classes, microstructure, instances, tractability, csps, tractable, family
Community 13: antenna, mimo, sparse, sum, downlink, multiuser, diversity, ofdma, transmitter, transmit
```

```
Community 14: hardware, wireless, synchronization, core, synchronize, bio, sensor, synchronized, nodes
Community 15: obstacle, learning, robots, avoidance, reinforcement, cuts, eigenvectors, recognition, clusters, clustering
Community 16: stationary, pattern, wave, electromyographic, bifurcation, generator, bifurcations, wavelet, signals
Community 17: methodology, reengineering, costing, techniques, analysis, business
Community 18: locomotion, manipulation, robot, loco, robotic, manipulators, planning, motions, primitive, motion
Community 19: stroke, gestures, ideomotor, imitate, neuron, dysfunction, mirror, neurophysiological, apraxia
Community 20: meetings, schedule, building, rooms, smart, scheduling, room, meeting
Community 21: meetings, schedule, building, rooms, smart, scheduling, room, meeting, efficient, queue
Community 22: music, factorization, supervised, separation, spectrogram, nmf, discriminative, spectral, score, midi
Community 23: surfaces, spline, blends, improved, shape, curvature, boundary, surface, multi, estimate
Community 24: algorithms, solver, programming, computing, chaining, set, approach, forward, sets
Community 25: mesh, surgical, models, surgery, planning, curvature, surface, meshes, smooth
Community 26: plane, waves, pulses, dynamics, spiral, pacing, arrhythmia, interaction
Community 27: ciliates, tools, dna, recombination, guided, ciliate, macronucleus, cilia, template, conjugation
Community 28: temporal, propagation, cstns, conditional, cstn, networks, consistency, propositional
Community 29: ancestral, evolution, inferring, reconstruction, method, evolutionary, dna, fitch
Community 30: ancestral, evolution, inferring, reconstruction, method, evolutionary, dna, fitch, trusted, security
Community 31: ea, agencies, government, architecture, enterprise, finnish, alignment, issues, governance
Community 32: control, landing, wind, airplane, shear, controller, aircraft, flying
Community 33: machine, development, prediction, programming, software, learning, size, estimation
Community 34: usability, livelihoods, mechanical, evaluating, workers, turk, literacy, microtasking
Community 35: form, concave, geometric, factor, polygons, radiosity, surface, assembly, cognitive, furniture
Community 36: images, ranking, tags, search, outlier, visually, reranking, visual, image, detection
Community 37: visual, image, images, fast, transcoding, feature, features, media, search, outlier
Community 38: wiener, compressor, comfort, listener, filter, suppression, noise, filtering
Community 39: wireless, pollutants, sensing, sensor, river, demonstration, tidal, monitoring, water
Community 40: capture, light, field, camera, cameras, resolution, aperture, imaging, language, animated
Community 41: randomization, exploit, associations, exploits, attack, runtime, attacks, code, reuse, address
Community 42: neural, optimization, randomization, exploit, associations, exploits, attack, runtime, attacks, code
Community 43: randomization, exploit, associations, exploits, attack, runtime, attacks, code, reuse, address
Community 44: embeddings, ultrametrics, outlier, outliers, metric, embedding, trajectories, graphs, algorithm, reeb
Community 45: net, routing, hierarchical, router, channel, routers, routed
Community 46: model, semantic, domain, research, models, life, biomedical, analysis, sciences
Community 47: cost, observers, cqf, oscillators, observer, minimization, harmonic, quantum, optimality
Community 48: internet, providers, games, cache, coalitional, coalition, caches, caching, service, behavior
Community 49: vision, lens, light, stitching, field, camera, extended, aperture
Community 50: optimum, evolution, algorithms, optimisation, unconstrained, hybrid, metaheuristic, benchmark, algorithm
Community 51: interface, bus, rpci, fpgas, reconfigurable, driver, pci
Community 52: programs, verification, model, systems, transition, compositional, concurrent, checking, behavioral, abstractions
Community 53: verification, model, systems, programs, transition, compositional, concurrent, checking, validator, tvoc
```

Give each community a name, and explain how you came to choose these names.

To give each community a meaningful name, we need to derive an intuitive title based on the **keywords** aggregated for that community. We can use the **top N keywords** for each community as the basis for naming it.

Approach:

1. **Extract Top Keywords**: We will use the top keywords from the community's keyword aggregation. These keywords give us a clear idea of the focus areas of that community.
2. **Generate a Community Name**: We can generate a community name by combining these keywords in a way that makes sense. For example:
   - **Top 3 Keywords**: We can use the top 3 keywords and combine them into a short phrase.
   - **Key Themes**: We may group keywords into key themes or fields (e.g., "AI", "neural networks", "NLP").
   - **Descriptive Phrases**: If possible, we may combine keywords into a descriptive phrase that clearly conveys the focus of the community.
3. **Handle Ambiguities**: If the keywords are too broad or don't form a coherent name, we can opt to combine them in a meaningful order, ensuring the community name reflects the underlying research area.

```
from collections import Counter

# Function to generate community names based on the top 3 most common keywords
def generate_name(keywords):
    # Count the frequency of each keyword
    keyword_counts = Counter(keywords)

    # Get the top 3 most common keywords
    top_keywords = [kw for kw, _ in keyword_counts.most_common(3)]

    # If we have at least 3 keywords, connect the first two with a comma, and the third with 'and'
    if len(top_keywords) == 3:
```

```
        return f"{top_keywords[0]}, {top_keywords[1]} and {top_keywords[2]}"
    # If we have only 2 keywords, connect them with 'and'
    elif len(top_keywords) == 2:
        return f"{top_keywords[0]} and {top_keywords[1]}"
    # If we have only 1 keyword, just return it
    else:
        return top_keywords[0]

# Assign names to each community based on their top 3 most common keywords
community_names = {}
for community_id, keywords in community_keywords.items():
    community_name = generate_name(keywords)
    community_names[community_id] = community_name

# Print out the community names
print("\nCommunity Names:")
for community_id, name in community_names.items():
    print(f"Community {community_id}: {name}")
```

```
Community Names:
Community 0: brownian, entropy and deterministic
Community 1: projections, rotations and images
Community 2: wavelet, projectors and fourier
Community 3: protocols, multiplexing and satellite
Community 4: taguchi, minimize and arrays
Community 5: customer, data and management
Community 6: abstractcardiotocographic, fetal and automated
Community 7: transform, streaming and distributed
Community 8: development, process and line
Community 9: analyser, services and chaîne
Community 10: scalable, processing and streaming
Community 11: detection, user and fpga
Community 12: constraint, tree and classes
Community 13: antenna, mimo and sparse
Community 14: hardware, wireless and synchronization
Community 15: obstacle, learning and robots
Community 16: stationary, pattern and wave
Community 17: methodology, reengineering and costing
Community 18: locomotion, manipulation and robot
Community 19: stroke, gestures and ideomotor
Community 20: meetings, schedule and building
Community 21: meetings, schedule and building
Community 22: music, factorization and supervised
Community 23: surfaces, spline and blends
Community 24: algorithms, solver and programming
Community 25: mesh, surgical and models
Community 26: plane, waves and pulses
Community 27: ciliates, tools and dna
Community 28: temporal, propagation and cstns
Community 29: ancestral, evolution and inferring
Community 30: ancestral, evolution and inferring
Community 31: ea, agencies and government
Community 32: control, landing and wind
Community 33: machine, development and prediction
Community 34: usability, livelihoods and mechanical
Community 35: form, concave and geometric
Community 36: images, ranking and tags
Community 37: visual, image and images
Community 38: wiener, compressor and comfort
Community 39: wireless, pollutants and sensing
Community 40: capture, light and field
Community 41: randomization, exploit and associations
Community 42: neural, optimization and randomization
Community 43: randomization, exploit and associations
Community 44: embeddings, ultrametrics and outlier
Community 45: net, routing and hierarchical
Community 46: model, semantic and domain
Community 47: cost, observers and cqf
Community 48: internet, providers and games
Community 49: vision, lens and light
Community 50: optimum, evolution and algorithms
Community 51: interface, bus and rpci
Community 52: programs, verification and model
Community 53: verification, model and systems
Community 54: distillation, mixture and separation
Community 55: recognition, robot and objects
```

2.3 Paper-Paper Clustering via Embedding

Embed the abstract and title of each paper using a fi tting embedding model such BERTopic or SentenceBERT, and aggregate them together.

Your embedding should contain the information of keywords of the community. An easy way would be to append the keywords as plaintext. Note: A paper might be included in several communities! An easy way to overcome this is to use the union set of these keywords.

Use any clustering algorithm of your choice to cluster these papers together.

Explanation of the Approach:

1. Embedding Papers with SentenceBERT:

- **SentenceBERT** is a powerful model specifically designed to create meaningful **embeddings** for sentences or documents. In this approach, **SentenceBERT** is used to transform the paper's **title**, **abstract**, and **community keywords** into a **vector representation**. These vectors represent the semantic meaning of the text and are suitable for clustering tasks.
- The paper's **title** and **abstract** are concatenated, forming a **single textual representation** of the paper. Additionally, the **keywords** from the communities the paper belongs to are appended as plain text. This combined text is used to create embeddings, meaning that the paper's content is enriched with the context provided by the keywords of the communities it is associated with.

2. Processing Papers in Batches:

- Due to the potential size of the dataset and the computationally intensive nature of generating embeddings, the texts are processed in **batches**.
- Each batch contains a subset of papers (32 papers in this case), which are passed through the **SentenceBERT model** for embedding generation.
- **Batch processing** helps improve the efficiency of the operation by leveraging parallel computation, reducing the time it takes to generate embeddings for all papers.

3. Generating Embeddings:

- The combined textual representation (title, abstract, and community keywords) is passed through the **SentenceBERT model** to generate a high-dimensional **embedding** (vector representation) for each paper.
- This step converts the papers from their raw textual form into dense vector representations that capture the semantic meaning of the paper's content.

4. Clustering with KMeans:

- Once the embeddings are generated, they are used as the input to the **KMeans clustering algorithm**.
- **KMeans** is a popular unsupervised machine learning algorithm that groups similar data points into **clusters**. The number of clusters is set to 5 in this case (although this can be adjusted depending on the desired granularity of the clustering).
- KMeans assigns each paper to a specific cluster based on the similarity of its embedding to those of other papers.

5. Assigning Clusters to Papers:

- After running the KMeans clustering algorithm, each paper is assigned a **cluster label**. This label indicates which group the paper belongs to, based on its content.
- The cluster label is then added to the original dataframe (`df_sampled`) in a new column called `cluster`.

6. Displaying the Results:

- Finally, the code prints the first 10 papers in the dataset along with their **assigned cluster** and **communities**. This allows you to see how each paper was grouped based on the content and community keywords.

```
!pip install sentence-transformers
```

```
Requirement already satisfied: sentence-transformers in /usr/local/lib/python3.10/dist-packages (3.3.1)
Requirement already satisfied: transformers<5.0.0,>=4.41.0 in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (4.47
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (4.67.1)
Requirement already satisfied: torch>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (2.5.1+cu121)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (1.2.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (1.13.1)
Requirement already satisfied: huggingface-hub>=0.20.0 in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (0.27.0)
Requirement already satisfied: Pillow in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (11.0.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.20.0->sentence-transformers)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.20.0->sentence-trans
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.20.0->sentence-transf
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.20.0->sentence-transforme
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.20.0->sentence-transformers)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.20.0->sent
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.11.0->sentence-transformers) (3.4.2)
```

```
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.11.0->sentence-transformers) (3.1.4)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (from torch>=1.11.0->sentence-transformers) (1.1
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy==1.13.1->torch>=1.11.0->sentenc
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers<5.0.0,>=4.41.0->sentence-transf
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers<5.0.0,>=4.41.0->sentence-
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/python3.10/dist-packages (from transformers<5.0.0,>=4.41.0->sent
Requirement already satisfied: safetensors>=0.4.1 in /usr/local/lib/python3.10/dist-packages (from transformers<5.0.0,>=4.41.0->sentence
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->sentence-transformers) (1.4.
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->sentence-transformers
Requirement already satisfied: mkl_fft in /usr/local/lib/python3.10/dist-packages (from numpy>=1.17->transformers<5.0.0,>=4.41.0->senten
Requirement already satisfied: mkl_random in /usr/local/lib/python3.10/dist-packages (from numpy>=1.17->transformers<5.0.0,>=4.41.0->sen
Requirement already satisfied: mkl_umath in /usr/local/lib/python3.10/dist-packages (from numpy>=1.17->transformers<5.0.0,>=4.41.0->sent
Requirement already satisfied: mkl in /usr/local/lib/python3.10/dist-packages (from numpy>=1.17->transformers<5.0.0,>=4.41.0->sentence-t
Requirement already satisfied: tbb4py in /usr/local/lib/python3.10/dist-packages (from numpy>=1.17->transformers<5.0.0,>=4.41.0->sentenc
Requirement already satisfied: mkl-service in /usr/local/lib/python3.10/dist-packages (from numpy>=1.17->transformers<5.0.0,>=4.41.0->se
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.11.0->sentence-transfor
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->huggingface-hub>=0.20
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->huggingface-hub>=0.20.0->sentence
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->huggingface-hub>=0.20.0->se
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->huggingface-hub>=0.20.0->se
Requirement already satisfied: intel-openmp>=2024 in /usr/local/lib/python3.10/dist-packages (from mkl->numpy>=1.17->transformers<5.0.0,
Requirement already satisfied: tbb==2022.* in /usr/local/lib/python3.10/dist-packages (from mkl->numpy>=1.17->transformers<5.0.0,>=4.41.
Requirement already satisfied: tcmlib==1.* in /usr/local/lib/python3.10/dist-packages (from tbb==2022.*->mkl->numpy>=1.17->transformers<
Requirement already satisfied: intel-cmplr-lib-rt in /usr/local/lib/python3.10/dist-packages (from mkl_umath->numpy>=1.17->transformers<
Requirement already satisfied: intel-cmplr-lib-ur==2024.2.0 in /usr/local/lib/python3.10/dist-packages (from intel-openmp>=2024->mkl->nu
```

```
from sentence_transformers import SentenceTransformer
from sklearn.cluster import KMeans
import numpy as np

# Initialize SentenceBERT model for embedding
model = SentenceTransformer('all-MiniLM-L6-v2')

# Function to embed the title, abstract, and community keywords of each paper
def embed_paper(row, community_keywords):
    # Combine title, abstract, and keywords from the paper's communities
    title_abstract = row["title"] + " " + row["abstract"]

    # Get the union of keywords from all communities this paper belongs to
    paper_keywords = set()
    for community_id in row["communities"]:
        paper_keywords.update(community_keywords.get(community_id, []))

    # Append the keywords to the title and abstract as plaintext
    combined_text = title_abstract + " " + " ".join(paper_keywords)

    return combined_text

# Generate a list of combined texts (title + abstract + community keywords) for all papers
texts = [embed_paper(row, community_keywords) for _, row in df_sampled.iterrows()]

# Generate embeddings in batches
batch_size = 32
paper_embeddings = []

# Process the texts in batches
for i in range(0, len(texts), batch_size):
    batch = texts[i:i + batch_size]
    batch_embeddings = model.encode(batch, batch_size=batch_size, show_progress_bar=True)
    paper_embeddings.extend(batch_embeddings)

# Convert the list of embeddings to a numpy array
paper_embeddings = np.array(paper_embeddings)

# Clustering the papers using KMeans
num_clusters = 5
kmeans = KMeans(n_clusters=num_clusters, random_state=42)
kmeans.fit(paper_embeddings)

# Assign the clusters to each paper
df_sampled['cluster'] = kmeans.labels_

# Print the first few papers with their assigned cluster
print("\nPapers with assigned clusters:")
for _, row in df_sampled.head(10).iterrows():
    print(f"Paper ID: {row['id']}")
    print(f"Title: {row['title']}")
```

```
    print(f"Assigned Cluster: {row['cluster']}")
    print(f"Communities: {row['communities']}\n")
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]

Papers with assigned clusters:
Paper ID: 78ebfe9c-d031-46bd-9069-b12c130713fc
Title: Information-singular random processes
Assigned Cluster: 2
Communities: [0]

Paper ID: dcc5d799-baf5-4dd7-8a06-fe1332c42637
Title: A linear algorithm for motion from three weak perspective images using Euler angles
Assigned Cluster: 0
Communities: [1]

Paper ID: e3ead250-4e03-41be-b6b0-8347f607b80c
Title: Quantitative Fourier analysis of approximation techniques. I. Interpolators and projectors
Assigned Cluster: 0
Communities: [2]

Paper ID: 7417fa37-babb-4c77-a5b1-10188305945c
Title: Voice and data transmissions with a PRMA-like protocol in high propagation delay cellular systems
Assigned Cluster: 1
Communities: [3]

Paper ID: 582d244f-a827-429e-a8d7-bc8f9e798a7d
Title: Design optimization of circular antenna arrays using Taguchi method
Assigned Cluster: 2
Communities: [4]

Paper ID: febbb04d-0c5e-4493-8a97-85a4f13a8f81
Title: Measuring the Impact of Data Mining on Churn Management.
Assigned Cluster: 4
Communities: [5]

Paper ID: 2b1cf591-c86f-455a-b11f-771a0cec7448
Title: Clustering with Pairs of Prototypes to Support Automated Assessment of the Fetal State
Assigned Cluster: 0
Communities: [6]

Paper ID: fd6055bc-8b6c-4bfb-b52a-9a4c22b91f87
Title: Revertible deep convolutional networks with iterated directional filter bank
Assigned Cluster: 0
Communities: [7]

Paper ID: 6539c480-b64f-4fb3-b686-9cfb014668d5
Title: Evolving Standard Process Reference Models for Product Line Development
Assigned Cluster: 3
Communities: [8]

Paper ID: 709df717-7e28-48ca-a479-eb46bd80e511
Title: Chaîne d’inclusion utilisant lestic
Assigned Cluster: 4
Communities: [9]
```

```
print(paper_embeddings)
```

```
[[-0.041722   -0.03756579 -0.0001923  ...  0.05860876 -0.02883449
  -0.00806739]
 [-0.06829602 -0.02260147 -0.01061126 ... -0.03228998 -0.03439865
  -0.00341776]
 [-0.14105532 -0.00502871  0.05631673 ...  0.04155594 -0.05585155
   0.05293049]
 ...
 [-0.03431151 -0.0703674  -0.01925877 ... -0.01602491  0.03701631
   0.02738878]
 [-0.01398173  0.00730337 -0.05761129 ...  0.03403294  0.05828181
   0.08209059]
 [ 0.02893061 -0.04842057  0.05345876 ...  0.13059716  0.00802048
   0.05434262]]
```

Analysis of the Clustering Results

1. Cluster Distribution

- **5 clusters (0–4) were formed**, and papers were assigned to one of these clusters.
- Some clusters (e.g., **Cluster 0**) appear **more frequently** than others, suggesting that multiple papers share **stronger thematic similarities** in that group.

---

2. Relationship Between Clusters and Communities

- Each paper belongs to **one or more author communities**, which represent the research field(s) of its authors.
- However, **cluster assignments do not strictly align with community assignments**, indicating that:
  - Clustering captures **semantic similarity** rather than just author relationships.
  - Some communities may work on **overlapping topics**, leading to shared clusters.

For example:

- Papers in **Cluster 0** belong to different communities (`[1], [2], [6], [7]`).
- This suggests **interdisciplinary connections**—papers from different research groups but with similar content were grouped together.

---

3. Potential Overlap in Research Topics

- **Cluster 2** contains papers on **random processes and antenna arrays**, which might indicate similarity in mathematical modeling or signal processing.
- **Cluster 0** includes papers on **motion algorithms, clustering techniques, and deep networks**, potentially suggesting **machine learning and pattern recognition** themes.
- **Cluster 4** includes topics like **churn management (business analytics) and inclusion chains (possibly in logic or mathematics)**, which might be **less related**, indicating that this cluster could be more general or noisy.

Evaluate your clustering algorithm using DBI and Silhouette scores. Hint: Silhouette scores are calculated for each point. To evaluate your overall clustering technique, you must somehow aggregate them together, and report the aggregate. For instance, the average silhouette score. This means you can choose a small subset of points to calculate their average on, however the silhouette score for each of these points must be calculated over the whole data. We are unsure if current python packages support this, but you are allowed to use these optimizations in case long runtimes become an issue.

Explanation of the Approach:

1. **Silhouette Score**: Measures how similar each data point is to its own cluster compared to other clusters.
2. **Davies-Bouldin Index (DBI)**: Measures the separation between clusters and the average similarity of a cluster to the most similar one.

1. Silhouette Score:

The **Silhouette Score** is a metric that combines both **cohesion** (how similar a data point is to the other points in its own cluster) and **separation** (how different the point is from points in other clusters).

- **For each data point**, the silhouette score is calculated based on two things:
  - **a**: The average distance between the data point and all other points in the same cluster (this is the cohesion).
  - **b**: The average distance between the data point and all points in the nearest different cluster (this is the separation).

The silhouette score for a single point is computed as: $S(i) = \frac{b(i)-a(i)}{\max(a(i),b(i))}$

- A **score close to 1** indicates that the point is well-clustered and far from other clusters.
- A **score close to -1** suggests that the point is likely in the wrong cluster.
- A **score close to 0** indicates that the point lies between clusters.

- **Overall Silhouette Score**: The overall silhouette score is the **average** of the silhouette scores for all data points in the dataset: $\text{Silhouette Score} = \frac{1}{N}\sum_{i=1}^{N} S(i)$ Where **N** is the total number of data points, and **S(i)** is the silhouette score for each individual data point.
  - This average score provides an overall evaluation of how well the data points are clustered.

2. Davies-Bouldin Index (DBI):

The **Davies-Bouldin Index** (DBI) is a metric that evaluates the **separation** between clusters and the **compactness** of the clusters.

- DBI measures the average similarity between each cluster and the cluster that is most similar to it.
  - A **lower DBI** indicates better clustering, as it implies that clusters are well-separated and compact.
  - A **higher DBI** suggests poor clustering, as it implies that the clusters are too close together or too diffuse.

The DBI for a given set of clusters is calculated as: $DBI = \frac{1}{N_c}\sum_{i=1}^{N_c} \max_{j\neq i}\left(\frac{S_i+S_j}{d(C_i,C_j)}\right)$ Where:

- **N_c** is the number of clusters.
- **S_i** is the average distance of points within cluster **i** (i.e., the compactness of cluster **i**).
- **d(C_i, C_j)** is the distance between the centroids of clusters **i** and **j** (i.e., the separation between clusters).
- The index measures the trade-off between compactness and separation: clusters that are both compact and well-separated will yield a low DBI.

3. How Scores for Each Point are Aggregated:

For each data point, the **silhouette score** is calculated based on its similarity to other points in its cluster (cohesion) and the closest other cluster (separation).

- **Aggregation of Silhouette Scores**: The **average silhouette score** for the entire dataset is obtained by calculating the average of all the individual silhouette scores for each data point in the dataset. This provides an overall measure of how well the clustering algorithm has grouped the points.
  - **High average silhouette score** (close to 1) indicates that the clusters are well-defined and that the data points are clearly separated from each other.
  - **Low average silhouette score** (close to 0 or negative) indicates that the data points might be poorly clustered, possibly with some points being placed in the wrong clusters.

```
from sklearn.metrics import silhouette_score, davies_bouldin_score
import numpy as np

# Function to calculate the silhouette score and DBI
def evaluate_clustering(paper_embeddings, kmeans_labels):
    # Silhouette Score: Use the entire dataset to calculate the silhouette score
    silhouette_avg = silhouette_score(paper_embeddings, kmeans_labels)

    # Davies-Bouldin Index: Calculate DBI for the given clustering
    dbi_score = davies_bouldin_score(paper_embeddings, kmeans_labels)

    return silhouette_avg, dbi_score

# Evaluate clustering for all data points
silhouette_avg, dbi_score = evaluate_clustering(paper_embeddings, kmeans.labels_)

# Print the results
print(f"Average Silhouette Score: {silhouette_avg:.4f}")
print(f"Davies-Bouldin Index: {dbi_score:.4f}")
```

```
Average Silhouette Score: 0.0335
Davies-Bouldin Index: 4.5915
```

Analysis of Clustering Evaluation Results:

The clustering results provided by the **Silhouette Score** and **Davies-Bouldin Index (DBI)** suggest that the clustering algorithm might not be performing very well.

1. Silhouette Score (0.0335):

The **Silhouette Score** ranges from **-1 to 1**:

- **1**: Indicates that the points are very well-clustered and far from other clusters.
- **0**: Indicates that the points are on or near the decision boundary between two clusters.
- **-1**: Indicates that the points are likely assigned to the wrong clusters.

Our score of **0.0335** is quite low and close to **0**, which suggests:

- **Weak Separation Between Clusters**: This means that the clusters are not well-separated, and some points may be placed near the boundaries between clusters.
- **Possibly Overlapping Clusters**: The clustering algorithm might have assigned some points to clusters that are not very distinct or might have placed them in the wrong clusters altogether.

2. Davies-Bouldin Index (4.5915):

The **Davies-Bouldin Index (DBI)** measures the **average similarity** between each cluster and the one most similar to it:

- **Lower values of DBI** indicate better clustering (more compact and well-separated clusters).
- **Higher values of DBI** indicate worse clustering, with more overlap and less separation between clusters.

Our DBI score of **4.5915** is **high**, suggesting:

- **Poor Separation Between Clusters**: The clusters are likely to be overlapping or not well-separated. A higher DBI indicates that the clusters are not sufficiently distinct from each other, which means that the clustering algorithm has failed to separate the papers into meaningful groups.
- **Large Variability in Cluster Sizes**: It might also indicate that some clusters are too large or poorly defined, leading to a high similarity between the cluster centroids and a high DBI.

Report the unique venue values of the dataset.

```
# Extract the unique venue values from the dataset
unique_venues = df['venue'].unique()

# Print the unique venue values
print("Unique venue values in the dataset:")
for venue in unique_venues:
    print(venue)
```

```
Unique venue values in the dataset:
international symposium on computers and communications
Neurocomputing
Electronic Journal of Combinatorics
international conference on computer vision
Ima Journal of Mathematical Control and Information
International Journal of Computer Aided Engineering and Technology
IEEE Journal of Solid-state Circuits
algorithmic learning theory
international conference on robotics and automation
Discrete Mathematics
international conference on management of data
programming language design and implementation
international symposium on circuits and systems
ACM Communications in Computer Algebra
nan
international conference on data engineering
winter simulation conference
bioinformatics research and development
Computers & Chemical Engineering
international conference on semantic computing
Information Systems Research
international conference on computer communications and networks
Journal of the Brazilian Computer Society
IEEE Transactions on Neural Networks
Synthesis Lectures on Human Language Technologies
modeling analysis and simulation of wireless and mobile systems
Numerical Algorithms
hawaii international conference on system sciences
IEEE\/ACM Transactions on Networking
international conference on computer vision systems
design automation conference
IEEE Transactions on Image Processing
```

```
Speech Communication
international conference on computer modeling and simulation
international conference on innovative computing, information and control
computer software and applications conference
NeuroImage
Nucleic Acids Research
custom integrated circuits conference
international conference on communications
business information systems
computer graphics international
parallel and distributed computing: applications and technologies
Cognitive Science
electronic commerce
theoretical aspects of software engineering
IEEE Transactions on Consumer Electronics
IEEE Transactions on Communications
international conference on parallel and distributed systems
European Journal of Combinatorics
systems, man and cybernetics
IEEE Transactions on Robotics
international conference on pattern recognition
Bioinformatics
IEEE Transactions on Automatic Control
IEEE Transactions on Very Large Scale Integration Systems
international conference on machine learning and applications
```

Compute the Jaccard Similarity Index of your clusters with the clustering of venues. The venue of each paper is a decent label for clustering as well!

Approach for Jaccard Similarity Calculation Between Clusters and Venues:

The goal of this part is to calculate the **Jaccard Similarity Index** between the clusters created by our clustering algorithm and the venues in our dataset. The **Jaccard Similarity** is a measure of similarity between two sets, and in this context, we are comparing the papers within each cluster to the papers within each venue.

1. Grouping Papers by Clusters and Venues:

- **Cluster Labels**: After performing clustering (in this case, using KMeans), each paper has been assigned a **cluster label**. This label indicates which cluster the paper belongs to.
- **Venue Labels**: Each paper also has an associated **venue**, which typically corresponds to the publication or event in which the paper was presented or published (e.g., a journal or conference).

The first step in this process is to **group papers by their clusters** and **group papers by their venues**:

- **By Clusters**: Papers in the same cluster are grouped together.
- **By Venues**: Papers in the same venue are grouped together.

2. Calculating Jaccard Similarity:

Once the papers are grouped, the **Jaccard Similarity** is calculated between each cluster and venue. The Jaccard Similarity measures the degree of overlap between the two groups of papers (cluster vs. venue): $J(A, B) = \frac{|A \cap B|}{|A \cup B|}$ Where:

- **A** is the set of paper IDs in the cluster.
- **B** is the set of paper IDs in the venue.
- **|A ∩ B|** is the size of the intersection of the two sets (i.e., the number of papers common to both the cluster and the venue).
- **|A ∪ B|** is the size of the union of the two sets (i.e., all unique papers in either the cluster or the venue).

The Jaccard Index is a value between **0** and **1**:

- A **higher value** (closer to 1) indicates that the cluster and the venue have a lot of overlap, meaning the papers in that cluster are very similar to those in the venue.
- A **lower value** (closer to 0) indicates less overlap and more dissimilarity between the cluster and venue.

3. Iterating Through Each Cluster:

- For each **cluster**, the Jaccard Similarity with each **venue** is calculated:
  - For each pair of a **cluster** and a **venue**, we compute the number of papers in common (intersection) and the total number of unique papers across both the cluster and venue (union).
  - This gives us a similarity score for that cluster-venue pair.
- We then record the venue with the **highest Jaccard Similarity** for each cluster, which indicates that the cluster and venue are most similar in terms of the papers they contain.

4. Handle Missing or Invalid Data (NaN):

- One of the issues we encountered was the presence of **NaN values** in the **venue** column. This suggests that some papers do not have a venue specified.
- To ensure that the Jaccard Similarity calculations are meaningful, we **filter out the papers** where the **venue is missing** (i.e., NaN values). By removing these, we ensure that the similarity is only calculated between valid venue groups.

5. Print the Results:

- Finally, for each **cluster**, the **venue** with the highest **Jaccard Similarity** is printed, along with the corresponding similarity score. This tells us how well the clustering corresponds to the existing venue-based groups.

```
# Filter out rows with missing venue values (NaN)
df_sampled = df_sampled.dropna(subset=["venue"])

# Group papers by KMeans clusters and venues (again after dropping NaN venues)
clustered_papers = defaultdict(set)
venue_grouped_papers = defaultdict(set)

# Group papers by cluster
for idx, row in df_sampled.iterrows():
    cluster = row['cluster']
    clustered_papers[cluster].add(row['id'])

# Group papers by venue
for idx, row in df_sampled.iterrows():
    venue = row['venue']
    venue_grouped_papers[venue].add(row['id'])

# Compute Jaccard Similarity
jaccard_scores = {}

# Iterate over clusters and compute Jaccard similarity for the corresponding venue cluster
for cluster, papers_in_cluster in clustered_papers.items():
    best_jaccard = 0
    best_venue = None
    for venue, papers_in_venue in venue_grouped_papers.items():
        # Calculate Jaccard Similarity between papers in the cluster and papers in the venue
        intersection = len(papers_in_cluster.intersection(papers_in_venue))
        union = len(papers_in_cluster.union(papers_in_venue))

        # Calculate Jaccard Index
        if union > 0:
            jaccard_index = intersection / union
            if jaccard_index > best_jaccard:
                best_jaccard = jaccard_index
                best_venue = venue

    jaccard_scores[cluster] = (best_venue, best_jaccard)

# Print the results
print("\nJaccard Similarity between clusters and venues:")
for cluster, (best_venue, score) in jaccard_scores.items():
    print(f"Cluster {cluster} is most similar to Venue {best_venue} with Jaccard Similarity Index of {score:.4f}")
```

```
Jaccard Similarity between clusters and venues:
Cluster 2 is most similar to Venue international conference on robotics and automation with Jaccard Similarity Index of 0.0420
Cluster 0 is most similar to Venue international conference on acoustics, speech, and signal processing with Jaccard Similarity Index of
Cluster 1 is most similar to Venue international conference on communications with Jaccard Similarity Index of 0.0567
Cluster 4 is most similar to Venue hawaii international conference on system sciences with Jaccard Similarity Index of 0.0281
Cluster 3 is most similar to Venue international parallel and distributed processing symposium with Jaccard Similarity Index of 0.0284
```

Analysis of the Results:

The **Jaccard Similarity Index** provides insight into how well the papers in each cluster overlap with the papers in the corresponding venues. The values we received indicate relatively low similarity between the clusters and their most similar venues, suggesting that the clusters are not very well-aligned with the venues in terms of paper content.

Jaccard Similarity Scores:

1. **Cluster 2 is most similar to Venue "International Conference on Robotics and Automation" with a Jaccard Similarity of 0.0420**:
   - This indicates that the papers in **Cluster 2** have a **small overlap** with the papers from the "**International Conference on Robotics and Automation**". A score of **0.0420** is quite low, meaning there isn't a strong similarity between the cluster and the venue. This suggests that **Cluster 2** might contain papers from a broader range of topics or poorly defined clusters.
2. **Cluster 0 is most similar to Venue "International Conference on Acoustics, Speech, and Signal Processing" with a Jaccard Similarity of 0.0399**:
   - Again, this similarity score is **low**. **Cluster 0** shares a small overlap with papers from this venue, which may indicate that the papers in **Cluster 0** do not focus specifically on the topics covered by the venue. It suggests that the clustering algorithm might have mixed together papers from multiple related topics.
3. **Cluster 1 is most similar to Venue "International Conference on Communications" with a Jaccard Similarity of 0.0567**:
   - This score is slightly higher than the previous ones, but it's still **low**. While **Cluster 1** has a bit more overlap with the "**International Conference on Communications**" venue, the low similarity still points to weak alignment between the cluster's papers and the venue.
4. **Cluster 4 is most similar to Venue "Hawaii International Conference on System Sciences" with a Jaccard Similarity of 0.0281**:
   - This score is even **lower**, indicating **almost no overlap** between **Cluster 4** and the venue. The low value suggests that the papers in **Cluster 4** are not highly related to the papers in this specific venue.
5. **Cluster 3 is most similar to Venue "International Parallel and Distributed Processing Symposium" with a Jaccard Similarity of 0.0284**:
   - Like **Cluster 4**, the similarity score is very **low**. This further supports the idea that **Cluster 3** contains papers that are not closely related to this venue. This might indicate that the cluster is too broad or not well-defined.

3. Citation Regressor

In this section we want to create a model that can estimate how many citations a given paper will have based on its abstract, title, year, and (Bonus) author-author network.

1.

The training data is the abstracts of the entire dataset. Preprocess the training data, and use appropriate embeddings. You are allowed to use LLM embeddings (e.g. BERT models such as BERTopic or SentenceBERT) to boost your model's performance. Do not forget the train/validation split.

The approach of this code is to build a machine learning model that can predict the number of citations a paper is likely to receive based on its abstract, title, year, and optionally, the author-author network.

1. **Data Preprocessing**:

- The dataset is first loaded, and any rows that are missing crucial information such as the **title**, **abstract**, or **number of citations** are dropped.
- The dataset is then filtered to exclude any **years** with only a single paper, which could be outliers or have insufficient data for analysis.
- **Stratified sampling** is applied to ensure the distribution of years in the sampled dataset is representative, keeping 1% of the dataset for faster processing while preserving the year distribution.

2. **Text Embedding Using SentenceBERT**:

- To make use of the paper's title and abstract for prediction, the **SentenceBERT** model (a pre-trained model optimized for generating embeddings for sentences) is employed to convert the **title and abstract** into vector representations.
- Each paper's **title and abstract** are combined into a single string and passed through SentenceBERT to generate embeddings, which are high-dimensional vectors that capture the semantic meaning of the text.
- These embeddings are then stored for each paper.

3. **Feature Preparation**:

- The embeddings generated for each paper are treated as the **features (X)** for the machine learning model. This allows the model to learn from the semantic content of the title and abstract, which should be indicative of how likely the paper is to be cited.
- The **number of citations (y)** is used as the target variable, and the goal is to predict this value based on the text embeddings.

4. **Data Splitting**:

- The dataset is split into **train and validation sets** (80% for training and 20% for validation). This ensures that the model can be trained on one subset of the data and then validated on another to test its generalization ability.
- **Train/validation split** is important to avoid overfitting, allowing the model to learn patterns that generalize well to unseen data.

5. **Model Development**:

- This prepares the dataset for training a model, where the embeddings (X) are used as input features, and the number of citations (y) is the output. The next step, which is not shown in the code provided, would involve selecting and training a regression model (e.g., Random Forest, Gradient Boosting, or a neural network) on the **train set** and evaluating its performance on the **validation set**.

```
from sklearn.model_selection import train_test_split
from sentence_transformers import SentenceTransformer
import pandas as pd
from sklearn.preprocessing import StandardScaler

# Load the dataset
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, low_memory=False)

# Drop rows with missing title, abstract, or citations
df = df.dropna(subset=["title", "abstract", "n_citation"])

# Filter out years with only one paper
year_counts = df['year'].value_counts()
valid_years = year_counts[year_counts > 1].index
df_filtered = df[df['year'].isin(valid_years)]

# Stratified sampling over the 'year' column to get 1% of the dataset (based on filtered data)
df_sampled, _ = train_test_split(df_filtered, test_size=0.99, stratify=df_filtered['year'], random_state=42)

# Initialize SentenceBERT model for embedding
model = SentenceTransformer('all-MiniLM-L6-v2')

# Function to create embeddings for the title and abstract
def create_embeddings(row):
    # Combine title and abstract into one text for embedding
    combined_text = row["title"] + " " + row["abstract"]
    embedding = model.encode(combined_text)
    return embedding

# Apply the embedding function to the title and abstract
df_sampled['embeddings'] = df_sampled.apply(create_embeddings, axis=1)

# Convert embeddings into separate columns
embeddings = pd.DataFrame(df_sampled['embeddings'].to_list(), index=df_sampled.index)


# Prepare the features (X) and target (y)
X = embeddings  # Use only the embeddings as features
y = df_sampled['n_citation']
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
# Split into train and validation sets (80% train, 20% validation)
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)
```

2.

Find the best regression model and hyperparameters using AutoML or GridSearch.

The approach of this code is focused on **hyperparameter tuning** and model selection to find the best regression model for predicting the number of citations a paper will receive. The strategy involves using **GridSearchCV** to explore a range of hyperparameters for multiple regression models and identify the one that performs best.

1. **Model Selection**:

- Several **regression models** are defined for the task, including:
  - **XGBoost**: A powerful ensemble model known for its speed and performance in many machine learning tasks.
  - **SVR (Support Vector Regression)**: A regression model based on Support Vector Machines, suitable for both linear and non-linear problems.
  - **KNN (K-Nearest Neighbors)**: A simple, instance-based learning model that predicts the value based on the average of the k-nearest data points.
  - **MLP (Multilayer Perceptron)**: A neural network model that can capture complex relationships in data through its hidden layers.

  These models are selected based on their diverse characteristics and their ability to handle various types of data and relationships in the problem.

2. **Hyperparameter Tuning**:

- Each model has a **parameter grid** defined, which includes different hyperparameters and their potential values. For instance:
  - **XGBoost**: We are tuning the number of estimators, learning rate, and maximum depth.
  - **SVR**: Hyperparameters like `C`, `epsilon`, and the `kernel` type are explored.
  - **KNN**: The number of neighbors, weight function, and distance metric are tested.
  - **MLP**: We vary the size of hidden layers, activation functions, and regularization strength (`alpha`).
- These grids are used to evaluate various combinations of hyperparameters, and the **best combination** for each model will be selected.

3. **GridSearchCV**:

- **GridSearchCV** is used to perform **exhaustive search** over the specified parameter grid for each model. It evaluates the performance of each model using **cross-validation** (in this case, 3-fold cross-validation), ensuring the model is generalized and not overfitting.
- The **scoring metric** used is **mean absolute error (MAE)**, which measures the average magnitude of errors in predictions without considering their direction. The goal is to minimize the MAE.

4. **Parallel Processing**:

- To speed up the grid search, **parallel processing** is enabled by setting `n_jobs=-1` and using `parallel_backend('threading')`. This ensures that the hyperparameter search runs faster by utilizing multiple CPU cores simultaneously.

5. **Model Evaluation**:

- After tuning the hyperparameters, the **best model** for each regression method is stored in the `best_models` dictionary.
- The model's performance is then evaluated on the **validation set** (previously separated during train/test split) by calculating the **mean absolute error (MAE)**, which provides an estimate of how well the model generalizes to unseen data.

6. **Best Model Selection**:

- The **best hyperparameters** for each model are printed, along with the **validation MAE** score, to help compare the performance of each model.

```
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.linear_model import LinearRegression
from sklearn.svm import SVR
from sklearn.neighbors import KNeighborsRegressor
from sklearn.neural_network import MLPRegressor
from xgboost import XGBRegressor
import lightgbm as lgb
```

```
import catboost as cb
from sklearn.metrics import mean_absolute_error
from joblib import parallel_backend

# Define models with additional ones
models = {
    'XGBoost': XGBRegressor(),
    'SVR': SVR(),
    'KNN': KNeighborsRegressor(),
    'MLP': MLPRegressor(max_iter=500)
}

# Updated parameter grid with more hyperparameters
param_grid = {
    'XGBoost': {
        'n_estimators': [100, 200],
        'learning_rate': [0.01, 0.1],
        'max_depth': [3, 5],
    },
    'SVR': {
        'C': [0.1, 1, 10],
        'epsilon': [0.01, 0.1, 0.5],
        'kernel': ['rbf', 'linear']
    },
    'KNN': {
        'n_neighbors': [3, 5, 7],
        'weights': ['uniform', 'distance'],
        'metric': ['euclidean', 'manhattan']
    },
    'MLP': {
        'hidden_layer_sizes': [(50,), (100,)],
        'activation': ['relu', 'tanh'],
        'alpha': [0.0001, 0.001]
    }
}

# Apply GridSearchCV for each model
best_models = {}

for model_name, model_instance in models.items():
    print(f"Optimizing {model_name}...")
    grid_search = GridSearchCV(estimator=model_instance,
                               param_grid=param_grid.get(model_name, {}),
                               cv=3,
                               n_jobs=-1,
                               verbose=1,
                               scoring='neg_mean_absolute_error')

    with parallel_backend('threading'):
        grid_search.fit(X_train, y_train)

    best_models[model_name] = grid_search.best_estimator_

    # Predict on validation set
    y_pred_grid = best_models[model_name].predict(X_val)
    mae_grid = mean_absolute_error(y_val, y_pred_grid)
    print(f"Best {model_name} model - MAE on validation set: {mae_grid:.4f}")
    print(f"Best parameters: {grid_search.best_params_}\n")
```

```
Optimizing XGBoost...
Fitting 3 folds for each of 8 candidates, totalling 24 fits
Best XGBoost model - MAE on validation set: 37.8940
Best parameters: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 100}

Optimizing SVR...
Fitting 3 folds for each of 18 candidates, totalling 54 fits
Best SVR model - MAE on validation set: 34.1990
Best parameters: {'C': 1, 'epsilon': 0.5, 'kernel': 'linear'}

Optimizing KNN...
Fitting 3 folds for each of 12 candidates, totalling 36 fits
Best KNN model - MAE on validation set: 50.4590
Best parameters: {'metric': 'manhattan', 'n_neighbors': 7, 'weights': 'uniform'}

Optimizing MLP...
Fitting 3 folds for each of 8 candidates, totalling 24 fits
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
```

```
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimize
  warnings.warn(
```

```
import joblib

# Find the best model based on lowest MAE
best_model_name = min(best_models, key=lambda name: mean_absolute_error(y_val, best_models[name].predict(X_val)))
best_model = best_models[best_model_name]

print(f"Best overall model: {best_model_name}")

# Save the best model
joblib.dump(best_model, "best_model.pkl")
print("Best model saved successfully!")
```

```
Best overall model: SVR
Best model saved successfully!
```

The results show the **best-performing regression models** based on **mean absolute error (MAE)** for each method tested in the hyperparameter tuning process.

1. **XGBoost (Best MAE: 37.8940)**

- **Best Parameters**:
  - **Learning rate**: 0.01
  - **Max depth**: 3
  - **Number of estimators**: 100
- **Performance**: The best XGBoost model achieved an MAE of **37.89** on the validation set. This means the model's predictions on average were off by about 37.89 citations, which is a relatively high error but expected due to the complexity of predicting citations based on textual features.
- **Interpretation**: The XGBoost model, with a learning rate of 0.01, max depth of 3, and 100 estimators, was able to make the most accurate predictions among all models tested. XGBoost's ensemble approach is known for handling non-linear relationships effectively, and these hyperparameters seem to have provided a balanced performance for this problem.

2. **SVR (Best MAE: 34.1990)**

- **Best Parameters**:

    - **C**: 1
    - **Epsilon**: 0.5
    - **Kernel**: Linear
- **Performance**: The SVR model performed well with an MAE of **34.20** on the validation set. The linear kernel suggests that the relationships between the features (embeddings of the title and abstract) and the target (citations) are fairly linear, and SVR is capturing that well.
- **Interpretation**: The SVR model's linear approach seems to be a good fit for this dataset, possibly because the relationship between the features and citations is somewhat linear. However, its performance is slightly worse than XGBoost's, likely due to the more complex nature of citation prediction, which benefits from non-linear models.

3. **KNN (Best MAE: 50.4590)**

- **Best Parameters**:
    - **Metric**: Manhattan
    - **Number of neighbors**: 7
    - **Weights**: Uniform
- **Performance**: KNN performed the worst among the models with an MAE of **50.46**, meaning its predictions on average were off by a larger margin compared to the others.
- **Interpretation**: KNN struggles here, likely because it doesn't capture the complexities and interactions between features as well as the other models. KNN relies on nearby neighbors and would have trouble distinguishing between papers with similar but subtly different citation patterns. This is likely why its performance is weaker in this task.

4. **MLP (Best MAE: 44.5598)**

- **Best Parameters**:
    - **Activation**: Tanh
    - **Alpha**: 0.001
    - **Hidden layer sizes**: (50,)
- **Performance**: The MLP model achieved an MAE of **44.56** on the validation set. The model did not converge optimally within the 500 iterations, as indicated by the **convergence warning**. Despite this, it still performed decently but not as well as XGBoost or SVR.
- **Interpretation**: MLP is a more flexible model with the potential to capture complex non-linear relationships, but it seems that the training process was not sufficient (due to convergence issues) to make the best use of this model. If trained longer or with better hyperparameters (e.g., more iterations or a different architecture), the MLP might improve, but currently, it doesn't outperform the other models.

Summary of Results:

- **Best Model**: **SVR** performed best (MAE = 34.199), followed by **XGBoost** (MAE = 37.894). Both of these models showed strong performance, with SVR performing slightly better in this case.
- **KNN** and **MLP** performed worse, with KNN showing the largest MAE (50.459), likely due to the limitations of the KNN algorithm for this problem.
- **Hyperparameter Choices**: XGBoost and SVR seem to work well with their optimal hyperparameters. The MLP model needs further tuning or more iterations to fully converge and improve its performance.

This is where your crawled data will come into play! Your crawled data is the test set. Use your crawled data to find the estimated citation count and report important regression metrics (RMSE, MAE, R2, etc.).

The approach in this code is focused on evaluating the performance of the best regression model (obtained from the previous hyperparameter tuning process) on a **test set** of papers that were crawled from a different dataset (in this case, using a **SQLite database** containing papers from **Semantic Scholar**).

1. **Loading and Preparing the Test Data**:

- The test data is **crawled** from a **SQLite database**, which contains information on papers, including their **title**, **abstract**, **authors**, and **citations**.
- The code extracts the relevant data from the database, filters out rows with missing critical information (such as **title**, **abstract**, or **citations**), and loads it into a **pandas DataFrame**.

2. **Embedding the Test Data**:

- To match the format of the training data, the **title and abstract** of each paper are combined and passed through a **SentenceBERT model** to generate **embeddings**. These embeddings represent the textual content of the papers as high-dimensional vectors, which capture the semantic meaning of the papers' content.
- **Batch processing** is used to generate embeddings efficiently for large amounts of data.

3. **Standardizing the Features**:

- The embeddings generated from the test data are **standardized** using **StandardScaler**. This is done to ensure that the features (embeddings) are on a similar scale as the features used during the training phase. Standardization improves the performance of many machine learning models by making the features comparable in magnitude.

4. **Model Loading**:

- The **best model** obtained during the **GridSearchCV** process (for hyperparameter tuning) is loaded from a saved file (presumably a `.pkl` file). This model has already been trained on the training and validation datasets, and it is now ready to make predictions on the test data.

5. **Prediction**:

- The model is used to **predict the number of citations** for the papers in the test set based on the embeddings of their **title and abstract**.

6. **Evaluation**:

- The performance of the model is evaluated using three key **regression metrics**:
  - **RMSE (Root Mean Squared Error)**: Measures the average magnitude of the errors in the predictions, penalizing large errors more than small ones.
  - **MAE (Mean Absolute Error)**: Measures the average magnitude of the errors in the predictions, treating all errors equally.
  - **R^2 Score (Coefficient of Determination)**: Indicates how well the model's predictions match the actual data. A higher R^2 value indicates better performance, with 1.0 being a perfect prediction.

7. **Challenges**:

The evaluation also addresses two key challenges:

- **Time Gap between Training/Validation Set and Test Set**: The **crawled data** (test set) could represent newer papers that were published after the training and validation sets were created. As a result, the model might struggle to predict citations for these newer papers, especially if their content or authorship differs significantly from those in the training set.
- **New Concepts/Authors**: New papers in the test set may introduce **concepts, topics, or authors** that were not present in the training set. This can create a challenge for the model, as it might not have learned about these novel elements, resulting in less accurate predictions.

8. **Reporting the Results**:

- The final step is to print out the **regression metrics** (RMSE, MAE, R^2) to assess how well the model generalizes to the test set.

```
import sqlite3
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
import pandas as pd
from sentence_transformers import SentenceTransformer
import numpy as np
from sklearn.preprocessing import StandardScaler
from joblib import load

# Load the crawled test dataset from the SQLite database
crawled_db_path = "/kaggle/input/papers/semantic_scholar_papers.db"
conn = sqlite3.connect(crawled_db_path)

# Access the list of table names
cursor = conn.cursor()
cursor.execute("SELECT name FROM sqlite_master WHERE type='table';")
tables = cursor.fetchall()

# Retrieve data from all tables (assuming we are interested in 'title', 'abstract', 'authors', 'citations')
data = []

for table in tables:
    table_name = table[0]
    if table_name not in ["sqlite_sequence"]:  # Skip internal SQLite table
        cursor.execute(f"SELECT * FROM {table_name}")
        rows = cursor.fetchall()
```

```
        columns = [desc[0] for desc in cursor.description]
        df_temp = pd.DataFrame(rows, columns=columns)
        data.append(df_temp)

# Concatenate all data into a single DataFrame (assuming the data is similar across tables)
df_crawled = pd.concat(data, ignore_index=True)

# Drop rows with missing title, abstract, or citations
df_crawled = df_crawled.dropna(subset=["title", "abstract", "citations"])

# Initialize the SentenceBERT model for embedding
model = SentenceTransformer('all-MiniLM-L6-v2')

# Function to create embeddings for the title and abstract (batch processing)
def create_embeddings_batch(df, model, batch_size=32):
    """Embed the texts in batches to reduce memory usage."""
    texts = df['title'] + " " + df['abstract']
    embeddings = model.encode(texts.tolist(), batch_size=batch_size, show_progress_bar=True)
    return embeddings

# Apply the embedding function to the title and abstract in the crawled data
embeddings = create_embeddings_batch(df_crawled, model)

# Convert the embeddings into a DataFrame
embeddings_df = pd.DataFrame(embeddings, index=df_crawled.index)

# Prepare the features (X) and target (y)
X_test = embeddings_df  # Use only embeddings as features
y_test = df_crawled['citations'].astype(float)  # Ensure target is numeric

# Standardize the features (e.g., embeddings)
scaler = StandardScaler()
X_test = scaler.fit_transform(X_test)

# Load the best model from GridSearchCV (assuming you saved it as 'best_model.pkl')
best_model = load('best_model.pkl')

# Predict on the crawled data
y_pred_test = best_model.predict(X_test)

# Evaluate the model on the test data
rmse = np.sqrt(mean_squared_error(y_test, y_pred_test))
mae = mean_absolute_error(y_test, y_pred_test)
r2 = r2_score(y_test, y_pred_test)

# Print regression metrics
print(f"Test RMSE: {rmse:.4f}")
print(f"Test MAE: {mae:.4f}")
print(f"Test R² Score: {r2:.4f}")
```

```
Batches:   0%|          | 0/278 [00:00<?, ?it/s]
Test RMSE: 441.9204
Test MAE: 134.0417
Test R² Score: -0.0739
```

The results suggest that the model is facing challenges in making accurate predictions on the test data.

1. **Test RMSE (Root Mean Squared Error: 441.9204)**

- **Interpretation**: The RMSE of **441.92** indicates that, on average, the model's predicted citation counts deviate from the actual values by **around 442 citations**. This is a **large error**, which suggests that the model is struggling to make accurate predictions. RMSE penalizes larger errors more than smaller ones, which means that the model is making some predictions that are significantly far from the true citation counts.
- **Implication**: The high RMSE suggests that the model is not performing well and may not generalize effectively to the test set, especially for papers with high citation counts.

2. **Test MAE (Mean Absolute Error: 134.0417)**

- **Interpretation**: The MAE of **134.04** indicates that, on average, the model’s predictions are off by about **134 citations**. While MAE is less sensitive to large errors compared to RMSE, it still reflects a substantial level of error in predictions.
- **Implication**: The MAE being lower than the RMSE further suggests that while there are some large errors, the model's overall predictions are consistently off by a notable amount across most test instances.

3. **Test R^2 Score (-0.0739)**

- **Interpretation**: The **R^2 score** of **-0.0739** is negative, which indicates that the model is performing worse than a **simple mean-based model**. R^2 values typically range from 0 to 1, with values closer to 1 indicating better model performance (i.e., a higher proportion of variance in the target variable is explained by the model). A negative R^2 score suggests that the model is failing to explain the variability in the test data and is worse than simply predicting the mean citation count for all papers.
- **Implication**: This negative R^2 score indicates that the model's predictions are very poor, and it is not capturing the relationship between the input features (embeddings) and the target (citations) effectively. In fact, using the mean of the citation counts would have resulted in a better model than the one currently being used.

Reasons for Poor Performance:

1. **Time Gap Between Training and Test Set**:
   - The test set likely includes newer papers that were not available in the training set. These newer papers might contain topics, authors, or citation patterns that are significantly different from those in the training data, which the model has not learned about. As a result, the model's predictions may be skewed or less accurate for these newer papers.
   - **Novelty of Concepts**: If new concepts, topics, or authors are introduced in the test set (which were not present in the training data), the model may not handle them well. It may not generalize well to unseen ideas, leading to higher errors in prediction.
2. **Author-Author Network Dynamics**:
   - The model does not appear to account for the **author-author network** or **co-authorship relationships**, which can be crucial for predicting citations. Authors who collaborate with highly-cited researchers or are part of prominent research communities may have a higher likelihood of being cited, which the model may not capture effectively if it is only using the embeddings of titles and abstracts.
3. **Limited Representation of Features**:
   - The model is based solely on **text embeddings** (title and abstract), which may not fully capture the citation dynamics. Citations are influenced by many factors beyond just the content of the paper, such as the **research community** the paper belongs to, the **reputation of the authors**, and the **external factors** like the paper's **visibility**, **conference participation**, or **media attention**.
   - Without incorporating additional features (e.g., authorship, conference data, journal prestige, or external visibility metrics), the model is limited in how it can capture the full complexity of citation behavior.
4. **Overfitting or Underfitting**:
   - The model might have **overfitted** to the training data, meaning it learned the patterns too well for the training set but is unable to generalize to the test set. Alternatively, it might be **underfitting**, meaning it failed to learn the underlying patterns in the data sufficiently to make accurate predictions.
   - **Hyperparameter Tuning**: Even with hyperparameter tuning, the model might not be flexible enough to capture the complex patterns in citation data, especially if the features used (like text embeddings) do not contain enough information.

Give explanations of the reasons your model has overcome/failed these major challenges:a. Time gap between training/validation set and test setb. New concepts/authors that may be introduced over time, not being in the training data

a. Time Gap Between Training/Validation Set and Test Set:

Challenge Explanation:

- The **time gap** between the training/validation set and the test set represents a **temporal shift**. The test set likely contains **newer papers** published after the training set data was collected. This creates a challenge for models trained on older papers since the citation behavior of new papers can be quite different. Over time, new research trends, citation practices, and influential papers can emerge, making it harder for the model to generalize and predict accurately for the test set.

Why the Model May Have Failed:

1. **Citation Dynamics Change Over Time**: The citation patterns of newer papers may evolve, and **older citation trends** may not apply to more recent papers. For example, papers published recently may receive citations more slowly due to newer publishing dynamics or increased competition in certain fields.
2. **Outdated Patterns**: The model was trained on papers from earlier years, and the **semantic understanding** of titles and abstracts may not reflect the terminology or research focus of newer papers. The training set may not have captured the **more recent trends in research** or the rising importance of certain topics or fields.
3. **Citations Accumulate Over Time**: Newer papers have not had as much time to accumulate citations compared to older papers. The model, trained on older papers, may have been overly influenced by the established patterns of citation accumulation, which doesn't

apply to newer papers.

Why the Model Might Not Have Overcome This:

- **Temporal Features**: The model did not explicitly account for **time-sensitive features** like publication year or trends in citation over time, which would help it adapt to the changing citation dynamics. This lack of temporal modeling limits the model's ability to make accurate predictions for papers from the test set that were published in later years.

b. New Concepts/Authors That May Be Introduced Over Time, Not Being in the Training Data:

Challenge Explanation:

- Over time, **new concepts, authors, and research topics** emerge. These newer papers might introduce terminology, technologies, or collaboration networks that were not present in the training data. As the research landscape evolves, new authors may become influential, and different concepts may take center stage in certain fields. If the training data doesn't include these new elements, the model will struggle to make accurate predictions for such papers.

Why the Model May Have Failed:

1. **Unseen Concepts**: The **embeddings** generated from the titles and abstracts of papers rely heavily on **contextual understanding** from the training set. If newer papers introduce concepts or terminology that were not prevalent in the training data, the model may not correctly interpret the meaning of these new terms or their relevance to citation prediction. This leads to poor generalization for papers discussing **new topics**.
2. **New Authors**: If papers in the test set are authored by individuals or research groups that weren't represented in the training set, the model may not understand the importance of these authors or their potential impact in the citation network. The citation behavior of these authors may differ significantly from the patterns the model has learned.
3. **Research Networks Evolve**: The **co-authorship network** and **collaboration patterns** may evolve over time, with new influential groups forming in recent years. If the model was trained on a set of co-authors that was less representative of current collaboration networks, it might fail to predict citations accurately for newer papers, especially those from high-impact but newer authors.

Why the Model Might Not Have Overcome This:

- **Lack of Author-Specific Features**: The model was trained primarily using text embeddings from the **title and abstract**, which might not capture **author-specific influence**. In particular, newer authors or researchers in emerging fields may not be well-represented in the training set. This absence of **author- and collaboration-based features** makes it hard for the model to predict how papers from new authors might gain citations.
- **Inadequate Generalization to Emerging Topics**: Since the training set didn't capture emerging topics or recent shifts in research trends, the model could not generalize well to newer papers in cutting-edge areas. The model may have struggled to identify relevant research areas or assess their citation potential based on outdated knowledge.

Bonus: Author-Author Network

The idea is that scholarly paper authors have relations with each other that may be a great predictor of how many citations a given paper is going to have. To utilize this network, you will have to resort to using a GNN. You are free to use any architecture.

```
!pip install torch_geometric
```

```
Collecting torch_geometric
  Downloading torch_geometric-2.6.1-py3-none-any.whl.metadata (63 kB)
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 63.1/63.1 kB 2.3 MB/s eta 0:00:00
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from torch_geometric) (3.11.11)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch_geometric) (2024.9.0)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch_geometric) (3.1.4)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from torch_geometric) (1.26.4)
Requirement already satisfied: psutil>=5.8.0 in /usr/local/lib/python3.10/dist-packages (from torch_geometric) (5.9.5)
Requirement already satisfied: pyparsing in /usr/local/lib/python3.10/dist-packages (from torch_geometric) (3.2.0)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from torch_geometric) (2.32.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from torch_geometric) (4.67.1)
Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->torch_geometric) (2.4.4
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->torch_geometric) (1.3.2)
Requirement already satisfied: async-timeout<6.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->torch_geometric) (5.0.1
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->torch_geometric) (25.1.0)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->torch_geometric) (1.5.0)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->torch_geometric) (6.1.0)
Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->torch_geometric) (0.2.1)
Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->torch_geometric) (1.18.3)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch_geometric) (3.0.2)
Requirement already satisfied: mkl_fft in /usr/local/lib/python3.10/dist-packages (from numpy->torch_geometric) (1.3.8)
Requirement already satisfied: mkl_random in /usr/local/lib/python3.10/dist-packages (from numpy->torch_geometric) (1.2.4)
Requirement already satisfied: mkl_umath in /usr/local/lib/python3.10/dist-packages (from numpy->torch_geometric) (0.1.1)
Requirement already satisfied: mkl in /usr/local/lib/python3.10/dist-packages (from numpy->torch_geometric) (2025.0.1)
```

```
Requirement already satisfied: tbb4py in /usr/local/lib/python3.10/dist-packages (from numpy->torch_geometric) (2022.0.0)
Requirement already satisfied: mkl-service in /usr/local/lib/python3.10/dist-packages (from numpy->torch_geometric) (2.4.1)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->torch_geometric) (3.4
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->torch_geometric) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->torch_geometric) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->torch_geometric) (2025.1.31
Requirement already satisfied: typing-extensions>=4.1.0 in /usr/local/lib/python3.10/dist-packages (from multidict<7.0,>=4.5->aiohttp->t
Requirement already satisfied: intel-openmp>=2024 in /usr/local/lib/python3.10/dist-packages (from mkl->numpy->torch_geometric) (2024.2.
Requirement already satisfied: tbb==2022.* in /usr/local/lib/python3.10/dist-packages (from mkl->numpy->torch_geometric) (2022.0.0)
Requirement already satisfied: tcmlib==1.* in /usr/local/lib/python3.10/dist-packages (from tbb==2022.*->mkl->numpy->torch_geometric) (1
Requirement already satisfied: intel-cmplr-lib-rt in /usr/local/lib/python3.10/dist-packages (from mkl_umath->numpy->torch_geometric) (2
Requirement already satisfied: intel-cmplr-lib-ur==2024.2.0 in /usr/local/lib/python3.10/dist-packages (from intel-openmp>=2024->mkl->nu
Downloading torch_geometric-2.6.1-py3-none-any.whl (1.1 MB)
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 1.1/1.1 MB 22.2 MB/s eta 0:00:0000:01
Installing collected packages: torch_geometric
Successfully installed torch_geometric-2.6.1
```

The approach of this code focuses on utilizing a **Graph Neural Network (GNN)** to predict the **citation count** of papers based on their **authors' relationships** and the **semantic content** of the papers' **title and abstract**. This is done by representing the **author-author network** as a graph and using embeddings generated from the paper titles and abstracts as features for the **authors' nodes**.

1. **Data Preprocessing**:

- The dataset is loaded, and rows with missing values in essential columns (title, abstract, citation count, and authors) are dropped.
- The dataset is **stratified sampled** to obtain a smaller, manageable subset for faster processing.
- **Text embeddings** for the title and abstract of each paper are generated using a **SentenceBERT model**. These embeddings provide a dense, high-dimensional representation of each paper's content, capturing semantic meaning from the text.

2. **Author-Author Network Construction**:

- A **Graph** is constructed where each **author** is represented as a **node**, and an **edge** is added between two authors if they co-authored a paper together. This is based on the authorship information in the dataset, which is split into individual authors for each paper.
- The **graph** is then used to model the relationships between authors, where the connectivity (edges) reflects their collaboration.

3. **Node Features and Targets**:

- **Author Node Features**: Each author's **node feature** is constructed by taking the **mean embedding** of all the papers that author has contributed to. This captures the author's involvement in various papers and their content (via the paper embeddings).
- **Citation Counts as Targets**: The **citation count** for each author is averaged over all the papers they have co-authored. This is used as the target for the GNN model, which aims to predict how many citations an author's papers are likely to receive.
- These node features and targets are then **standardized** (scaled) to ensure the model performs better by making the features comparable in magnitude.

4. **Graph Data Representation**:

- The **author network** is represented in a format suitable for **PyTorch Geometric (PyG)**. The authors' relationships (edges in the graph) are converted to **edge indices**, and the authors' node features and citation targets are also converted into **tensors** that can be processed by the GNN model.

5. **Graph Neural Network (GNN) Architecture**:

- A **Graph Convolutional Network (GCN)** is defined, which uses **GCNConv layers** to propagate information through the graph.
  - The first layer transforms the input features into a higher-dimensional space.
  - The second layer further processes the features, and the final layer outputs a prediction for the citation count (a regression task).
- The **GCN** layers aggregate information from neighboring nodes (i.e., co-authors) and update the representation of each node (author) based on the graph structure.

6. **Model Training**:

- The GNN model is trained using **mean squared error (MSE)** as the loss function, which is common for regression tasks. The optimizer (Adam) is used to update the model's weights during training.
- The training loop runs for **1000 epochs**, and the loss is monitored to track the model's learning progress.

7. **Testing and Evaluation**:

- After training, the model is evaluated on the test data (the entire graph) to predict citation counts for authors in the test set.

- Evaluation metrics such as **RMSE (Root Mean Squared Error)**, **MAE (Mean Absolute Error)**, and **R^2 score** are calculated to assess the model's performance.
- **Sample predictions** are printed to give an indication of how well the model has learned to predict citation counts.

8. **Challenges Addressed**:

- **Author-Author Relationships**: The model incorporates the author-author relationships through the graph structure, allowing it to learn how **co-authorship** affects citation patterns.
- **Graph Structure**: By using a GNN, the model can aggregate information across the graph to make more informed predictions, considering not only the paper's content but also the collaborative network between authors.

9. **Model Output**:

- The model's predictions are compared to the actual citation counts, and the evaluation metrics give a measure of how well the model is generalizing to the test data.
- The **RMSE**, **MAE**, and **R^2 score** provide insights into the model's accuracy, while sample predictions allow for qualitative evaluation of the model's output.

Conclusion:

This approach combines the power of **semantic embeddings** (via SentenceBERT) and **graph-based learning** (via GNN) to predict citation counts based on the relationships between authors and the content of their papers. The graph structure captures the collaborative nature of academic work, and the GNN leverages this information to predict how likely papers from different authors are to be cited.

```
import pandas as pd
import networkx as nx
import numpy as np
import torch
import torch.nn.functional as F
from torch_geometric.data import Data
from torch_geometric.nn import GCNConv
from sklearn.preprocessing import StandardScaler
from sentence_transformers import SentenceTransformer
from sklearn.model_selection import train_test_split

# Load dataset
file_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
df = pd.read_csv(file_path, low_memory=False)

# Drop missing values
df = df.dropna(subset=["title", "abstract", "n_citation", "authors"])

# Filter out years with only one paper
year_counts = df['year'].value_counts()
valid_years = year_counts[year_counts > 1].index
df_filtered = df[df['year'].isin(valid_years)]

# Stratified sampling over the 'year' column to get 1% of the dataset
df_sampled, _ = train_test_split(df_filtered, test_size=0.99, stratify=df_filtered['year'], random_state=42)

# Initialize the SentenceBERT model
sentence_model = SentenceTransformer('all-MiniLM-L6-v2')

# Generate embeddings for papers
df_sampled["embedding"] = df_sampled.apply(lambda row: sentence_model.encode(row["title"] + " " + row["abstract"]), axis=1)

# Convert embeddings to DataFrame
embedding_df = pd.DataFrame(df_sampled["embedding"].to_list(), index=df_sampled.index)

# Process authors: Convert to lists
df_sampled["authors"] = df_sampled["authors"].apply(lambda x: x.split(";") if isinstance(x, str) else [])
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
Batches:   0%|          | 0/1 [00:00<?, ?it/s]
```

```
df_sampled["authors"] = df_sampled["authors"].apply(lambda x: x[0] if isinstance(x, list) and len(x) > 0 else [])


G = nx.Graph()

for authors in df_sampled["authors"]:
    authors = [a.strip() for a in authors if a.strip()]  # Remove empty strings
    if len(authors) > 1:  # Only add edges if there are at least two authors
        for i in range(len(authors)):
            for j in range(i + 1, len(authors)):
                G.add_edge(authors[i], authors[j])

print(f"Number of nodes: {G.number_of_nodes()}")
print(f"Number of edges: {G.number_of_edges()}")
```

```
Number of nodes: 143
Number of edges: 5468
```

```
# Convert authors to indices
author_index = {author: idx for idx, author in enumerate(G.nodes)}

# Create edge index for PyTorch Geometric
edges = [(author_index[a], author_index[b]) for a, b in G.edges]
edge_index = torch.tensor(edges, dtype=torch.long).T

# Compute author node features and targets
node_features = []
node_targets = []

author_embeddings = {author: [] for author in G.nodes}
author_citations = {author: [] for author in G.nodes}

for _, row in df_sampled.iterrows():
    paper_embedding = row["embedding"]
    citation_count = row["n_citation"]
    for author in row["authors"]:
        if author in author_embeddings:
            author_embeddings[author].append(paper_embedding)
            author_citations[author].append(citation_count)

for author in G.nodes:
    if author_embeddings[author]:
        mean_embedding = np.mean(author_embeddings[author], axis=0)
        mean_citation = np.mean(author_citations[author])
    else:
        mean_embedding = np.zeros(embedding_df.shape[1])
        mean_citation = 0
    node_features.append(mean_embedding)
    node_targets.append(mean_citation)


# Standardize node features
scaler = StandardScaler()
node_features = scaler.fit_transform(node_features)

# Convert to tensors
x = torch.tensor(node_features, dtype=torch.float)
y = torch.tensor(node_targets, dtype=torch.float).view(-1, 1)

# Create PyG Data object
data = Data(x=x, edge_index=edge_index, y=y)

# Define GNN Model
class GNN(torch.nn.Module):
    def __init__(self, input_dim):
        super(GNN, self).__init__()
        self.conv1 = GCNConv(input_dim, 128)
        self.conv2 = GCNConv(128, 64)
        self.conv3 = GCNConv(64, 1)  # Predict citation count

    def forward(self, x, edge_index):
        x = self.conv1(x, edge_index)
        x = F.relu(x)
        x = self.conv2(x, edge_index)
        x = F.relu(x)
        x = self.conv3(x, edge_index)
        return x
```

```
# Initialize model, loss function, and optimizer
model = GNN(input_dim=x.shape[1])
optimizer = torch.optim.Adam(model.parameters(), lr=0.01)
loss_fn = torch.nn.MSELoss()

# Train the model
def train():
    model.train()
    optimizer.zero_grad()
    out = model(data.x, data.edge_index)
    loss = loss_fn(out, data.y)
    loss.backward()
    optimizer.step()
    return loss.item()

# Training loop
for epoch in range(1000):
    loss = train()
    if epoch % 10 == 0:
        print(f"Epoch {epoch}, Loss: {loss:.4f}")

# Test Predictions
model.eval()
predictions = model(data.x, data.edge_index).detach().numpy()
actuals = data.y.detach().numpy()

# Compute Evaluation Metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
rmse = np.sqrt(mean_squared_error(actuals, predictions))
mae = mean_absolute_error(actuals, predictions)
r2 = r2_score(actuals, predictions)

# Print Metrics
print("\nEvaluation Metrics:")
print(f"Test RMSE: {rmse:.4f}")
print(f"Test MAE: {mae:.4f}")
print(f"Test R² Score: {r2:.4f}")

# Print sample predictions
print("\nSample Predictions:")
for i in range(5):
    print(f"Actual: {actuals[i][0]}, Predicted: {predictions[i][0]}")
```

```
Epoch 0, Loss: 2321.9739
Epoch 10, Loss: 1071.5736
Epoch 20, Loss: 909.1199
Epoch 30, Loss: 915.3893
Epoch 40, Loss: 902.5979
Epoch 50, Loss: 899.4375
Epoch 60, Loss: 894.3231
Epoch 70, Loss: 891.5502
Epoch 80, Loss: 888.2373
Epoch 90, Loss: 884.6586
Epoch 100, Loss: 880.9351
Epoch 110, Loss: 877.1542
Epoch 120, Loss: 872.7039
Epoch 130, Loss: 867.1329
Epoch 140, Loss: 866.2673
Epoch 150, Loss: 866.0807
Epoch 160, Loss: 865.8951
Epoch 170, Loss: 865.8867
Epoch 180, Loss: 865.8390
Epoch 190, Loss: 865.8172
Epoch 200, Loss: 865.7912
Epoch 210, Loss: 865.7656
Epoch 220, Loss: 865.7409
Epoch 230, Loss: 865.7258
Epoch 240, Loss: 865.7085
Epoch 250, Loss: 865.6872
Epoch 260, Loss: 865.6688
Epoch 270, Loss: 865.6443
Epoch 280, Loss: 865.6169
Epoch 290, Loss: 865.6144
Epoch 300, Loss: 865.6454
Epoch 310, Loss: 865.7726
Epoch 320, Loss: 865.6017
Epoch 330, Loss: 865.5044
Epoch 340, Loss: 865.4606
Epoch 350, Loss: 865.4385
```

```
Epoch 360, Loss: 865.4124
Epoch 370, Loss: 865.4339
Epoch 380, Loss: 865.3705
Epoch 390, Loss: 865.4799
Epoch 400, Loss: 865.3851
Epoch 410, Loss: 865.3361
Epoch 420, Loss: 865.3394
Epoch 430, Loss: 865.3073
Epoch 440, Loss: 865.3213
Epoch 450, Loss: 865.2935
Epoch 460, Loss: 865.2461
Epoch 470, Loss: 865.2151
Epoch 480, Loss: 865.4250
Epoch 490, Loss: 865.4543
Epoch 500, Loss: 865.2825
Epoch 510, Loss: 865.1864
Epoch 520, Loss: 865.1451
Epoch 530, Loss: 865.1267
Epoch 540, Loss: 865.1083
Epoch 550, Loss: 865.1105
Epoch 560, Loss: 865.0622
Epoch 570, Loss: 867.2050
```

Improvement in Results:

The **new results** (Test RMSE: 29.4278, Test MAE: 16.4397, and Test R^2: 0.0061) indicate a **significant improvement** compared to the previous results where the model had poor performance with higher RMSE, MAE, and negative R^2 scores. Here's how the new model compares to previous results and what has improved:

1. Test RMSE: 29.4278 (Previous RMSE: Much Higher)

- **Improvement**: The **RMSE** has decreased substantially. A lower RMSE indicates that the model is now making predictions that are closer to the actual citation values, with an error of **29.43 citations** on average. In the previous results, the RMSE was much higher, indicating larger discrepancies between the predicted and actual citation counts.
- **Implications**: This improvement suggests that the model has better captured the relationship between paper features and citation counts. The model is more **accurate** and is performing better at predicting citation counts across the test set, with smaller prediction errors compared to before.

2. Test MAE: 16.4397 (Previous MAE: Much Higher)

- **Improvement**: The **MAE** of **16.44** indicates that the predictions are now closer to the actual values, with an average error of **16 citations**. In previous results, the model had a higher MAE, indicating larger errors in prediction.
- **Implications**: This improvement in MAE means the model is now providing more **precise predictions** overall. The model has improved in handling citation predictions more consistently, making fewer significant errors compared to before.

3. Test R^2 Score: 0.0061 (Previous R^2 Score: Negative)

- **Improvement**: The **R^2 score** of **0.0061** is a **small positive value**, showing that the model is now able to explain a small fraction of the variance in the citation counts, which is an improvement from the previous **negative R^2 score**.
  - A **negative R^2** indicates that the model was performing worse than a simple mean-based prediction (essentially making random guesses), which was the case in the previous results.
  - **Positive R^2**: While still very low, the **positive R^2** suggests that the model is now making more informed predictions, capturing some of the patterns in the data. It indicates that the model has learned to predict citation counts better than simply predicting the average citation for all papers.

Sample Predictions (Improvement):

- **Previous Predictions**: In the previous model, predictions were highly inaccurate and showed large errors.
- **Current Predictions**:
  - **Actual: 39.91, Predicted: 31.85**
  - **Actual: 39.62, Predicted: 36.76**
  - **Actual: 40.13, Predicted: 39.07**
  - **Actual: 38.32, Predicted: 40.41**
  - **Actual: 46.36, Predicted: 41.27**

These **sample predictions** are now much closer to the actual citation counts, with some variations that are more reasonable. In the previous model, the predictions were far off the mark, but now, the model is more consistent with its predictions, even if there are still some discrepancies.

Reasons for Improvement:

1. **Use of Graph Neural Networks (GNNs)**:
   - The **GNN** allowed the model to leverage the **author-author network**. By modeling the relationships between authors, the GNN could capture collaboration dynamics and better predict how often a paper might be cited based on the **author network** and their past citations.
2. **Incorporating Author-Level Features**:
   - The model incorporated **author embeddings** based on their co-authorship network, which provided additional context to the **text embeddings** (from titles and abstracts). This allowed the model to better understand the citation behavior of papers from different authors and research groups.
3. **Improved Model Architecture**:
   - The **Graph Convolutional Network (GCN)** architecture utilized in the GNN helped the model to capture **spatial dependencies** (i.e., author relationships) more effectively, leading to better performance compared to simpler models that only relied on textual embeddings.
4. **Better Feature Representation**:
   - The use of **SentenceBERT embeddings** for the papers' titles and abstracts provided richer feature representations that helped the model understand the semantic content of the papers more effectively. This likely helped improve the predictions of how likely a paper is to be cited.
5. **Optimization and Hyperparameter Tuning**:
   - It is possible that the **training process** and **hyperparameter tuning** (involving loss functions, learning rates, etc.) contributed to better model convergence and learning, improving accuracy compared to previous versions of the model.

Conclusion:

The new model has shown a **significant improvement** over previous results, with a **substantial reduction in error metrics** (RMSE and MAE) and a **positive R^2 score**. The **GNN architecture**, which models the **author-author network**, appears to have helped the model capture more complex relationships in citation data, leading to more accurate predictions. However, there is still room for further improvements, particularly in increasing the model's ability to explain the variability in citation counts and addressing any remaining gaps in feature representation.

4. Product

The end goal is to create a product called Research Assistant. The Research Assistant is a LLM enhanced by RAG that can:

1. Find corresponding papers for a given topic;
2. And (hopefully) the LLM can perform analysis on these relevant papers and provide reasonable reports

4.1 RAG

Create a pipeline that takes a string input from the user, and will retrieve relevant papers from the database of papers. The database is a vector database, and is indexed based on the embeddings of abstracts/titles of your original dataset + crawled dataset and returns the relevant papers using a similarity metric (cosine similarity, L2, etc.). Use your model as the embedder for this stack. You are allowed to use any tech stack for this.

Explanation of our Approach:

1. Loading and Preprocessing Data:

- **Original Dataset**: First, a Kaggle dataset of research papers is loaded into a DataFrame. The dataset includes paper metadata such as titles, abstracts, and authors, along with publication years.
- **Year Column Handling**: The 'year' column is converted to numeric format, ensuring that any non-numeric values are properly handled (either by coercing them to NaN or dropping them entirely). The rows with missing year values are removed, and the year is cast to integers.
- **Stratified Sampling**: A 1% stratified sample is taken from the dataset based on the 'year' column, which ensures that papers from all years are proportionally represented in the sample.

2. Loading Additional Crawled Data:

- **SQLite Connection**: A separate database of crawled papers (e.g., from Semantic Scholar) is connected via SQLite.
- **Load and Merge Data**: Multiple tables from the crawled data (representing different types of models) are loaded, and the data from these tables are merged into a single DataFrame. This step is done to combine all relevant papers from various categories into one unified dataset.

3. **Embedding Generation**:

- **Text Embedding**: The paper's metadata (title, abstract, and authors) are concatenated into a single string for each paper. This combined text is then passed to a pre-trained sentence transformer model to generate embeddings (numerical representations) of the text.
- **Batch Processing**: To optimize for time and memory, the embeddings for both the stratified sample and the combined dataset are generated in batches.

4. **Database Setup**:

- **PostgreSQL Connection**: The code connects to a PostgreSQL database via a transaction pooler, using connection details such as the database name, username, and password.
- **Enable Vector Extension**: It checks if the PostgreSQL instance has the `pgvector` extension installed (used for handling vector data) and enables it if needed.
- **Table Creation**: A table in the PostgreSQL database is created (if it doesn't already exist) to store the papers along with their embeddings. The table has columns for paper ID, title, abstract, authors, and the vector embedding.

5. **Inserting Data into PostgreSQL**:

- **Batch Insertion**: The papers (with their embeddings) are inserted into the database in batches to improve performance. A batch size of 1,000 records is used for each insert operation.
- **Progress Tracking**: A progress bar is used to track the progress of the data insertion, showing how many rows have been inserted so far.
- **Database Commit**: After each batch, the changes are committed to the database to ensure that data is saved properly.

6. **Querying the Database**:

- **Embedding the Query**: A sample query text ("Deep learning in NLP") is generated, and its embedding is computed using the same sentence transformer model.
- **Vector Search**: The code uses PostgreSQL's vector similarity search capabilities (via `pgvector`) to find the most similar papers to the query. The database is queried to retrieve the top 5 papers that are most similar to the query based on their embeddings.
- **Error Handling**: The query is enclosed in a try-except block to handle any errors (e.g., transaction issues) that may occur during the query. If an error occurs, the transaction is rolled back to avoid any issues with the subsequent queries.

7. **Final Clean-Up**:

- **Commit Changes**: The final commit ensures that all changes to the database are saved after the query, whether it was successful or if there was an error (which would result in a rollback).
- **Exception Handling**: If anything goes wrong, the code ensures that the transaction is rolled back, preventing issues from propagating further.

In Summary:

1. **Data Preprocessing**: The code begins by loading and preparing datasets, ensuring proper formatting, and taking a representative sample based on publication year.
2. **Text Embedding**: The text (title, abstract, authors) of each paper is converted into numerical embeddings using a transformer model.
3. **Database Setup**: The code then creates a PostgreSQL table to store the papers and their embeddings, using the `pgvector` extension for efficient vector similarity search.
4. **Data Insertion**: The embeddings and associated metadata are inserted into the database in batches with progress tracking.
5. **Similarity Search**: A query is issued to find the most similar papers to a given input based on their embeddings.
6. **Error Handling**: Throughout the code, careful exception handling ensures that any failures (such as transaction errors) are properly managed by rolling back the transaction.

This approach efficiently manages large datasets, performs vector-based similarity searches, and ensures robust error handling for working with embeddings in a relational database.

```
import sqlite3
import pandas as pd

# Load the original Kaggle dataset
df_path = "/kaggle/input/research-papers-dataset/dblp-v10.csv"
```

```
df = pd.read_csv(df_path, low_memory=False)

# Ensure the 'year' column exists and is numeric
df['year'] = pd.to_numeric(df['year'], errors='coerce')  # Convert to numeric, handling errors
df = df.dropna(subset=['year'])  # Drop rows with missing year
df['year'] = df['year'].astype(int)  # Convert to integer

# Get a stratified 1% sample based on the 'year' column
df_sampled = df.groupby('year', group_keys=False).apply(lambda x: x.sample(frac=0.01, random_state=42))

# Reset index after sampling
df_sampled = df_sampled.reset_index(drop=True)

# Connect to the SQLite database for crawled data
conn = sqlite3.connect('/kaggle/input/papers/semantic_scholar_papers.db')

# Load tables from the crawled dataset
tables = pd.read_sql_query("SELECT name FROM sqlite_master WHERE type='table';", conn)
print(tables)

# Load all relevant tables (Foundation_Models, Generative_Models, LLM, VLM, Diffusion_Models)
foundation_models_df = pd.read_sql_query("SELECT * FROM Foundation_Models;", conn)
generative_models_df = pd.read_sql_query("SELECT * FROM Generative_Models;", conn)
llm_df = pd.read_sql_query("SELECT * FROM LLM;", conn)
vlm_df = pd.read_sql_query("SELECT * FROM VLM;", conn)
diffusion_models_df = pd.read_sql_query("SELECT * FROM Diffusion_Models;", conn)

# Merge the datasets using pd.concat()
df_combined = pd.concat([foundation_models_df, generative_models_df, llm_df, vlm_df, diffusion_models_df], ignore_index=True)

# Display the first few rows of the combined dataset
print(df_combined.head())
```

```
<ipython-input-1-dee05176cd9b>:14: DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns. This behavior is depreca
  df_sampled = df.groupby('year', group_keys=False).apply(lambda x: x.sample(frac=0.01, random_state=42))
                 name
0    sqlite_sequence
1  Foundation_Models
2  Generative_Models
3                LLM
4                VLM
5   Diffusion_Models
   id                                           title  \
0   1  Analysis of Rectangular Plates on Resting Wink...
1   2  REDWIN Foundation Models for Integrated Dynami...
2   3  Investigation on the sensitivity of flexible f...
3   4  Effect of foundation models on seismic respons...
4   5    On the Suitability of Elastic Foundation Models

                                            abstract  \
0  Analysis of rectangular plates on a rising ela...
1  \n Due to the complex nature of the loads on O...
2                                             None
3                                             None
4  Elastic foundation models are extensively used...

                                          authors  citations
0                                        T. Nguyen          6
1     A. Page, K. Norén-Cosgriff, K. Skau, A. Kaynia          5
2  Yang Yang, Chun Li, M. Bashir, Jin Wang, Chunm...         33
3  Ai-Yun Jin, Jianwen Pan, Jin-Ting Wang, Chong ...         22
4         Aditi Kumawat, P. Raychowdhury, S. Chandra          0
```

```
from sentence_transformers import SentenceTransformer
import numpy as np
import torch

# Initialize the SentenceTransformer model
device = 'cuda' if torch.cuda.is_available() else 'cpu'
model = SentenceTransformer('all-MiniLM-L6-v2', device=device)

# 1) Combine title, abstract, and authors into a single text for embedding
def build_text_for_embedding(row):
    """
    Combine title, abstract, and authors into a single text for embedding.
    """
    title = str(row["title"]) if pd.notnull(row["title"]) else ""
```

```
    abstract = str(row["abstract"]) if pd.notnull(row["abstract"]) else ""
    authors = str(row["authors"]) if pd.notnull(row["authors"]) else ""

    return f"{title} {abstract} {authors}"

# Combine title, abstract, and authors for both datasets
combined_texts = df_sampled.apply(build_text_for_embedding, axis=1)
combined_texts_combined = df_combined.apply(build_text_for_embedding, axis=1)

# 2) Encode the combined text in batches for faster processing
# Convert combined texts into a list
combined_texts_list = combined_texts.tolist()
combined_texts_combined_list = combined_texts_combined.tolist()


# Generate embeddings in batches
embeddings = model.encode(combined_texts_list, batch_size=32, show_progress_bar=True)
embeddings_combined = model.encode(combined_texts_combined_list, batch_size=32, show_progress_bar=True)

# 3) Add embeddings to the dataframes
df_sampled['embeddings'] = list(embeddings)
df_combined['embeddings'] = list(embeddings_combined)

# 4) Add authors column to the crawled dataframes (since it was missing)
df_combined['authors'] = df_combined['authors'].fillna('Unknown')  # Replace NaN with 'Unknown' if needed

# 5) Combine both datasets into one
df_combined_final = pd.concat([df_sampled[['id', 'title', 'abstract', 'authors', 'embeddings']],
                               df_combined[['id', 'title', 'abstract', 'authors', 'embeddings']]], ignore_index=True)

# Display the first few rows to verify
print(df_combined_final.head())
```

```
config.json:   0%|          | 0.00/612 [00:00<?, ?B/s]
model.safetensors:   0%|          | 0.00/90.9M [00:00<?, ?B/s]
tokenizer_config.json:   0%|          | 0.00/350 [00:00<?, ?B/s]
vocab.txt:   0%|          | 0.00/232k [00:00<?, ?B/s]
tokenizer.json:   0%|          | 0.00/466k [00:00<?, ?B/s]
special_tokens_map.json:   0%|          | 0.00/112 [00:00<?, ?B/s]
Batches:   0%|          | 0/313 [00:00<?, ?it/s]
Batches:   0%|          | 0/313 [00:00<?, ?it/s]
                                     id  \
0  bf52c5db-8868-44ee-97ff-6efbd9d60b03
1  4b80d21b-ef9a-49a4-a720-99a2b58b9003
2  67214be9-1c4b-46ed-bdfa-80b281ad45bb
3  9d222f7a-c85a-44ed-bcfa-86231b0ced40
4  a95c4ff5-137b-418c-aa8f-3ca49061561d

                                               title  \
0  Systematic Scaling for Digital Differential An...
1  A matched filter communication system for mult...
2                                  Automatic coding
3  The representation of narrow-band processes (C...
4   A conjecture on binary nongroup codes (Corresp.)

                                           abstract                authors  \
0  The usefulness of large-capacity digital diffe...        ['Arthur Gill']
1  A matched filter communication system is descr...  ['Steven M. Sussman']
2  SDC is currently engaged in producing a “JOVIA...     ['Harvey Bratman']
3                                            NaN        ['Moshe Zakai']
4                                            NaN     ['Bernard Elspas']

                                         embeddings
0  [-0.13542506, 0.08947143, -0.14937374, -0.1416...
1  [-0.0899295, -0.042870738, -0.013408764, -0.09...
2  [-0.16730231, -0.033565685, -0.09799634, -0.01...
3  [-0.26527673, -0.027648574, 0.018452138, 0.105...
4  [-0.14948253, -0.012078531, -0.40769556, -0.04...
```

```
!pip install psycopg2-binary
```

```
Collecting psycopg2-binary
  Downloading psycopg2_binary-2.9.10-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (4.9 kB)
Downloading psycopg2_binary-2.9.10-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (3.0 MB)
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 3.0/3.0 MB 29.1 MB/s eta 0:00:0000:0100:01
Installing collected packages: psycopg2-binary
Successfully installed psycopg2-binary-2.9.10
```

```
import psycopg2

# Connect to PostgreSQL via Transaction Pooler
conn = psycopg2.connect(
    dbname="postgres",                          # Database name
    user="postgres.zvgfzwfhpnvlaumrtdhb",                           # Username
    password="Pani_papers_db",                 # Password
    host="aws-0-eu-central-1.pooler.supabase.com",  # Transaction pooler host
    port="6543"                               # Port for Transaction Pooler
)

# Create a cursor object to interact with the database
cur = conn.cursor()
print("Connected to PostgreSQL via Transaction Pooler successfully!")

# Example: Run a sample query
cur.execute("SELECT NOW();")
print("Current time:", cur.fetchone())
```

```
Connected to PostgreSQL via Transaction Pooler successfully!
Current time: (datetime.datetime(2025, 2, 8, 15, 56, 38, 549995, tzinfo=datetime.timezone.utc),)
```

```
# Enable the pgvector extension
cur.execute("CREATE EXTENSION IF NOT EXISTS vector;")

# Create table with vector embeddings
cur.execute("""
    CREATE TABLE IF NOT EXISTS my_new_papers (
        id SERIAL PRIMARY KEY,
        title TEXT,
        abstract TEXT,
        authors TEXT,
        embedding vector(384)
    );
""")
conn.commit()
print("Table created successfully!")
```

```
Table created successfully!
```

```
pip install tqdm
```

```
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (4.67.1)
Note: you may need to restart the kernel to use updated packages.
```

```
from psycopg2.extras import execute_values
from tqdm import tqdm  # Import tqdm

# Define batch size
BATCH_SIZE = 1000  # Adjust as needed

# Prepare the data list
data_list = [
    (row['title'], row['abstract'], row['authors'], np.array(row['embeddings']).tolist())
    for _, row in df_combined_final.iterrows()
]

# Insert data in batches with progress bar
with tqdm(total=len(data_list), desc="Inserting Data", unit="row") as pbar:  # Progress bar
    for i in range(0, len(data_list), BATCH_SIZE):
        batch = data_list[i:i + BATCH_SIZE]
        execute_values(cur, """
            INSERT INTO my_new_papers (title, abstract, authors, embedding)
            VALUES %s
        """, batch)
        conn.commit()

        # Update the progress bar with the number of rows inserted so far
        pbar.update(len(batch))

print("Data inserted successfully!")
```

```
Inserting Data: 100%|██████████| 20002/20002 [00:55<00:00, 358.92row/s]Data inserted successfully!
```

```
# Define the query paper
query_text = "Deep learning in NLP"

# Generate the embedding for the query
query_embedding = model.encode([query_text])[0].tolist()

try:
    # Start a new transaction by closing the previous one (in case it was in an aborted state)
    conn.rollback()  # Rollback any failed transaction before starting a new one

    # Query PostgreSQL for the most similar papers
    cur.execute("""
        SELECT title, abstract, authors
        FROM my_new_papers
        ORDER BY embedding <-> %s::vector  -- Explicitly cast to vector
        LIMIT 5;
    """, (query_embedding,))

    # Fetch results
    results = cur.fetchall()

    # Display results
    for title, abstract, authors in results:
        print(f"\nTitle: {title}\nAbstract: {abstract}\nAuthors: {authors}")

except Exception as e:
    # Catch any exception and rollback if an error occurs
    print(f"Error: {e}")
    conn.rollback()  # Rollback the transaction to avoid further errors

    # You can log the error message or print it
    print("Rolling back the transaction due to an error.")

finally:
    # Ensure that the connection is always committed (if no error) or closed (if error occurs)
    conn.commit()  # Commit transaction if necessary (or if using autocommit)
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]

Title: Neural Networks Leverage Corpus-wide Information for Part-of-speech Tagging
Abstract: We propose a neural network approach to benefit from the non-linearity of corpuswide statistics for part-of-speech (POS) taggi
Authors: ['Yuta Tsuboi']

Title: To Transformers and Beyond: Large Language Models for the Genome
Abstract: In the rapidly evolving landscape of genomics, deep learning has emerged as a useful tool for tackling complex computational c
Authors: Micaela Elisa Consens, Cameron Dufault, Michael Wainberg, Duncan Forster, Mehran Karimzadeh, Hani Goodarzi, Fabian Theis, Alan

Title: To Transformers and Beyond: Large Language Models for the Genome
Abstract: In the rapidly evolving landscape of genomics, deep learning has emerged as a useful tool for tackling complex computational c
Authors: Micaela Elisa Consens, Cameron Dufault, Michael Wainberg, Duncan Forster, Mehran Karimzadeh, Hani Goodarzi, Fabian Theis, Alan

Title: Rethinking Complex Neural Network Architectures for Document Classification
Abstract: Neural network models for many NLP tasks have grown increasingly complex in recent years, making training and deployment more
Authors: Ashutosh Adhikari, Achyudh Ram, Raphael Tang, Jimmy J. Lin

Title: Rethinking Complex Neural Network Architectures for Document Classification
Abstract: Neural network models for many NLP tasks have grown increasingly complex in recent years, making training and deployment more
Authors: Ashutosh Adhikari, Achyudh Ram, Raphael Tang, Jimmy J. Lin
```

```
# Define the query paper
query_text = "machine learning for healthcare"

# Generate the embedding for the query
query_embedding = model.encode([query_text])[0].tolist()

try:
    # Start a new transaction by closing the previous one (in case it was in an aborted state)
    conn.rollback()  # Rollback any failed transaction before starting a new one

    # Query PostgreSQL for the most similar papers
    cur.execute("""
        SELECT title, abstract, authors
        FROM my_new_papers
        ORDER BY embedding <-> %s::vector  -- Explicitly cast to vector
        LIMIT 5;
    """, (query_embedding,))
```

```
    # Fetch results
    results = cur.fetchall()

    # Display results
    for title, abstract, authors in results:
        print(f"\nTitle: {title}\nAbstract: {abstract}\nAuthors: {authors}")

except Exception as e:
    # Catch any exception and rollback if an error occurs
    print(f"Error: {e}")
    conn.rollback()  # Rollback the transaction to avoid further errors

    # You can log the error message or print it
    print("Rolling back the transaction due to an error.")

finally:
    # Ensure that the connection is always committed (if no error) or closed (if error occurs)
    conn.commit()  # Commit transaction if necessary (or if using autocommit)
```

```
Batches:   0%|          | 0/1 [00:00<?, ?it/s]

Title: Shifting machine learning for healthcare from development to deployment and from models to data
Abstract: None
Authors: Angela Zhang, L. Xing, James Zou, Joseph C. Wu

Title: Artificial intelligence in medicine workshop: AAAI 1992 Spring Symposium Series Stanford University March 25-27, 1992 Workshop Su
Abstract: NaN
Authors: ['Michael G. Kahn']

Title: Multi-layer Representation Learning for Medical Concepts
Abstract: Proper representations of medical concepts such as diagnosis, medication, procedure codes and visits from Electronic Health Re
Authors: ['Edward Choi', 'Mohammad Taha Bahadori', 'Elizabeth Searles', 'Catherine Coffey', 'Michael R. Thompson', 'James E. Bost', 'Jav

Title: Machine intelligence for health information: capturing concepts and trends in social media via query expansion.
Abstract: NaN
Authors: ['Xing yu Su', 'Hanna Suominen', 'Leif Hanlen']

Title: Poster Abstract: Automated Medical Algorithms: Issues for Medical Errors.
Abstract: NaN
Authors: ['Katherine A. Johnson', 'John R. Svirbely', 'M. G. Sriram', 'Jack W. Smith', 'Gareth Kantor', 'Jorge Raul Rodriguez']
```

```
# Close connection
cur.close()
conn.close()
```

Tech Stack Used and How:

In this solution, we mainly used a **classic approach** with PostgreSQL combined with the `pgvector` extension to handle vector embeddings and perform similarity searches. Below is a breakdown of the tech stack and how each part is utilized:

1. **PostgreSQL with pgvector Extension (Classic Approach)**:

- **PostgreSQL**: This relational database system is used as the backend to store paper metadata (title, abstract, authors) and the vector embeddings generated from the text of these papers. The database is used to efficiently manage and query the data.
- **pgvector Extension**: This PostgreSQL extension enables the database to handle vector data efficiently, which is crucial for tasks like similarity search based on vector embeddings. By storing embeddings in the database, we can leverage vector search capabilities (like the `<->` operator) to find similar papers based on their embeddings.
- **How We Used It**:
  - We created a table in PostgreSQL to store paper metadata and their respective embeddings.
  - We inserted the embeddings of the papers into this table in batches for efficient storage and retrieval.
  - We used the `pgvector` extension to perform similarity searches on the embeddings using the `<->` operator to find the most similar papers to a query. This method is highly efficient and works well for searching large datasets of vectors.
- *Why This Approach**:
- PostgreSQL with `pgvector` provides a powerful and scalable solution for storing and querying vector data.
- It allows us to leverage SQL queries for managing metadata and performing vector-based search using standard relational database operations.

- The `pgvector` extension makes vector search easy and efficient, enabling us to scale without needing a complex search engine or additional frameworks.

The user should also be able to retrieve papers based on the author. Use NER to your advantage.

Bonus: You may use the LLM to find the author name, and then somehow search the database by that.

Approach Explanation

The goal of this solution is to allow users to retrieve research papers from a database based on the author, using **Named Entity Recognition (NER)** for author extraction and querying the PostgreSQL database for relevant papers.

---

1. User Query Processing

- **User Input**: The user will provide a query like "Find papers by John Doe on NLP" or "Can you retrieve papers by Jane Smith related to deep learning?"
- **Extracting Author Name**: The query may include the author's name, but it could also include other information (like the research area). To focus on extracting the author's name, we use **Named Entity Recognition (NER)**. NER is a technique in natural language processing (NLP) that identifies specific entities in text, such as names of people, places, organizations, etc.

  In this approach, we use the **spaCy** library, a powerful NLP tool, to extract the **PERSON** entities (i.e., names of people) from the query.
- *Why NER?**
- NER allows us to handle **complex, unstructured user input** where the author's name may appear anywhere in the query.
- NER helps us avoid the need for users to format their queries strictly with author names, making the system more flexible.

---

2. NER with spaCy

- **Input Text**: The query entered by the user, e.g., "Can you find papers by John Doe on machine learning?"
- **Processing with NER**: Using a pre-trained **spaCy NER model**, we extract entities tagged as `PERSON`. The model identifies "John Doe" as a person